



Ash Framework

Create Declarative
Elixir Web Apps

Rebecca Le and Zach Daniel

Series editor: Sophie DeBenedetto

Development editor: Kelly Lee



Under Construction: The book you're reading is still under development. As part of our Beta book program, we're releasing this copy well before a normal book would be released. That way you're able to get this content a couple of months before it's available in finished form, and we'll get feedback to make the book even better. The idea is that everyone wins!

Be warned: The book has not had a full technical edit, so it will contain errors. It has not been copyedited, so it will be full of typos, spelling mistakes, and the occasional creative piece of grammar. And there's been no effort spent doing layout, so you'll find bad page breaks, over-long code lines, incorrect hyphenation, and all the other ugly things that you wouldn't expect to see in a finished book. It also doesn't have an index. We can't be held liable if you use this book to try to create a spiffy application and you somehow end up with a strangely shaped farm implement instead. Despite all this, we think you'll enjoy it!

Download Updates: Throughout this process you'll be able to get updated ebooks from your account at pragprog.com/my_account. When the book is complete, you'll get the final version (and subsequent updates) from the same address.

Send us your feedback: In the meantime, we'd appreciate you sending us your feedback on this book at pragprog.com/titles/lash/errata, or by using the links at the bottom of each page.

Thank you for being part of the Pragmatic community!

The Pragmatic Bookshelf

Ash Framework

Create Declarative Elixir Web Apps

Rebecca Le
Zach Daniel

The Pragmatic Bookshelf

Dallas, Texas



See our complete catalog of hands-on, practical,
and Pragmatic content for software developers:

<https://pragprog.com>

Sales, volume licensing, and support:

support@pragprog.com

Derivative works, AI training and testing,
international translations, and other rights:

rights@pragprog.com

Copyright © 2025 The Pragmatic Programmers, LLC.

All rights reserved. No part of this publication may be reproduced by any means, nor may any derivative works be made from this publication, nor may this content be used to train or test an artificial intelligence system, without the prior consent of the publisher.

When we are aware that a term used in this book is claimed as a trademark, the designation is printed with an initial capital letter or in all capitals.

The Pragmatic Starter Kit, The Pragmatic Programmer, Pragmatic Programming, Pragmatic Bookshelf, PragProg, and the linking *g* device are trademarks of The Pragmatic Programmers, LLC.

Every precaution was taken in the preparation of this book. However, the publisher assumes no responsibility for errors or omissions or for damages that may result from the use of information (including program listings) contained herein.

ISBN-13: 979-8-88865-152-0

Encoded using recycled binary digits.

Book version: B1.0—January 14, 2025

Contents

Foreword	v
Change History	vii
Welcome!	ix
1. Building Our First Resource	1
Getting the Ball Rolling	1
Oh, CRUD! — Defining Basic Actions	11
Integrating Actions into LiveViews	22
2. Extending Resources with Business Logic	35
Resources and Relationships	35
Loading Related Resource Data	43
Structured Data with Validations and Identities	47
Deleting All of the Things	51
Changing Data Within Actions	54
3. Creating a Better Search UI	61
Custom Actions with Arguments	61
Dynamically Sorting Artists	68
Pagination of Search Results	73
No DB field? No Problems, with Calculations	78
Relationship Calculations as Aggregates	82

4.	Generating APIs Without Writing Code	89
	Model Your Domain, Derive The Rest	89
	Building a JSON REST Interface	90
	Building a GraphQL Interface	103
5.	Authentication: Who Are You?	113
	Introducing AshAuthentication	113
	Setting Up Password Authentication	115
	Automatic UIs With AshAuthenticationPhoenix	119
	Setting Up Magic Link Authentication	126
6.	Authorization: What Can You Do?	133
	Introducing Policies	133
	Authorizing API Access for Authentication	135
	Assigning Roles to Users	143
	Writing Policies for Artists	146
	Removing Forbidden Actions from the UI	153
	Writing Policies for Albums	161
7.	All About Testing	169
8.	Fun With Nested Forms	171
9.	PubSub and Real-Time Notifications	173
	Bibliography	175

Foreword

Content to be supplied later.

Change History

The book you're reading is in beta. This means that we update it frequently. Here is the list of the major changes that have been made at each beta release of the book, with the most recent change first.

B1.0: January 14th, 2025

- Initial beta release.

Welcome!

As software developers, we face new and interesting challenges daily. When one of these problems appears, our instincts are to start building a mental model of the solution. The model might contain high-level concepts, ideas, or *things* that we know we want to represent, and ways they might communicate with each other to carry out the desired task.

Your next job is to find a way to map this model onto the limitations of the language and frameworks available to you. But there's a mismatch: your internal model is a fairly abstract representation of the solution, but the tooling you use demands very specific constructs, often dictated by things such as database schemas and APIs.

These problems are hard, but they're not intractible — they *can* be solved, using a framework like Ash.

Ash lets you think and code at a higher level of abstraction, and your resulting code will be cleaner, easier to manage, and you'll be less frustrated.

This book will show you the power of Ash, and how to get the most out of it in your Elixir projects.

What is Ash?

Ash is a set of tools you can use to describe and build the *domain model* of your applications — the “things” that make up what your app is supposed to do, and the business logic of how they relate and interact with each other. If you’re building an e-commerce store, your domain model will have things like products, categories, suppliers, orders, customers, deliveries, and more; and you’ll already have a mental model to describe how they fit together. Ash is how you can translate that mental model into code, using standardized patterns and your own terminology.

Ash is a fantastic application framework, but it is *not* a web framework. This question comes up often, so we want to be really clear up front — Ash doesn’t

replace Phoenix, Plug, or any other web framework when building web apps in Elixir. It does, however, slide in very nicely alongside them and work with them, and when combined they can make the ultimate toolkit for building amazing apps.

What can Ash offer an experienced Elixir/Phoenix developer? You're already familiar with a great set of tools for building web applications today, and Ash *builds* on that foundation that you know and love. It leverages the rock-solid Ecto library for its database integrations, and its resource-oriented design helps bring structure and order to the wild west of Phoenix contexts. If this sounds interesting to you, keep reading!

And if you're only just starting on your web development journey, we'd love to introduce you to our battle-tested and highly-productive stack!

Why Ash?

Ash is built on three fundamental principles. These principles are rooted in the concept of *declarative design*, and have arisen from direct encounters with the good, bad and the ugly of software in the wild. They are:

- Data > Code
- Derive > Hand-write
- What > How

To paraphrase a famous manifesto, while there is value in the items on the right, we value the items on the left more.

No principle is absolute and each has its own tradeoffs, but together they can help us build rich, maintainable and scalable applications. The “why” of Ash is rooted in the “why” of each of these core principles.

Data > Code

With Ash, we model (describe) our application components with *resource* modules, using code that compiles into pre-defined data structures. These resources describe the interfaces to, and behavior of, the various components of our application.

Ash can take the data structures created by these descriptions, and use them to do wildly useful things with little-to-no effort. Furthermore, Ash contains tools that allow you to leverage your application-as-data to build and extend your application in fully custom ways. You can introspect and use the data structures in your own code, and you can even write transformers to extend the language that Ash uses and add new behaviour to existing data.

Taking advantage of these super powers, however, requires learning the language of Ash Framework, and this is what we'll teach you in this book.

Derive > Hand-write

We emphasize *deriving* application components from our descriptions, instead of hand-writing our various application layers. When building a JSON API, for example, you might end up hand-writing controllers, serializers, OpenAPI schemas, error handling, and the list goes on. If you want to add a GraphQL API as well, you have to do it all over again with queries, mutations, and resolvers. In Ash, this is all driven from your resource definitions, using them as the single source of truth for how your application should behave. Why should you need to restate your application logic in five different ways?

There *is* value in the separation of these concerns, but that value is radically overshadowed by all of the associated costs, such as:

- The cost of bugs via functionality drift in your various components
- The cost of the conceptual overhead required to implement changes to your application and each of its interfaces
- And the cost, especially, of every piece of your application being a special snowflake with its own design, idiosyncrasies and patterns.

When you see what Ash can derive automatically, without all of the costly spaghetti code necessary with other approaches, the value of this idea becomes very clear.

What > How

This is the core principle of declarative design, and you've almost certainly leveraged this principle already in your time as a developer without even realizing it.

Two behemoths in the world of declarative design are HTML and SQL. When writing code in either language, you don't describe *how* the target is to be achieved, only *what* the target is. For HTML, a renderer is in charge of turning your HTML descriptions into pixels on a screen; and for SQL, a query planner and engine are responsible for translating your queries into procedural code that reads data from storage.

An Ash resource behaves in the exact same way, as a description of the *what*. All of the code in Ash is geared towards looking at the descriptions of what you want to happen, and making it so. This is a crucial thing to keep in mind as you go through this book — when we write resources, we are only describing their behavior. Later, when we actually call the actions we describe,

or connect them to an API using an API extension for example, Ash looks at the description provided to determine what is to be done.

These principles, and the insights we derive from them, might take some time to really comprehend and come to terms with. As we go through the more concrete concepts presented in this book, revisit these principles. Ash is more than just a new tool, it's a new way of thinking about how we build applications in general.

We've seen time and time again, especially in our in-person workshops, that everyone has a moment when these concepts finally *click*. This is when Ash stops feeling like magic, and begins to look like what it really is: the principles of declarative design, taken to their natural conclusion.

Model your domain, derive the rest.

Is This Book For You?

If you've gotten this far, then yes, this book is for you!

If you have some experience with Elixir and Phoenix, have heard about this library called Ash, and are keen to find out more, then this book is *definitely* for you.

If you're a grizzled Elixir veteran wondering what all the Ash fuss is about, it's also for you!

If you've already been working with Ash, even professionally, you'll still learn new things from this book (but you can read it a bit faster).

If you haven't used Elixir before, this book is probably not for you *yet* — but it might be soon! To learn about this amazing functional programming language, we highly recommend working through [Elixir in Action \[Jur15\]](#). To get a feel for how modern web apps are built in Elixir with Phoenix and Phoenix LiveView, [Programming Phoenix LiveView \[TD24\]](#) will get you up to speed. And then you can come back here, and keep reading!

What's In This Book

This book is divided into nine chapters, each one building on top of the previous to flesh out the domain model for a music database. We'll provide the starter Phoenix LiveView application to get up and running, and then away we'll go!

In [Chapter 1, Building Our First Resource, on page 1](#), we'll set up the Tunez starter app, install and configure Ash, and get familiar with CRUD actions.

We'll build a full (simple) resource, complete with attributes, actions, and a database table; and integrate those actions into the web UI using forms and code interfaces.

In [Chapter 2, Extending Resources with Business Logic, on page 35](#), we'll create a second resource, and learn about linking resources together with relationships. We'll also cover more advanced features of resources, like preparations, validations, identities, and changes.

In [Chapter 3, Creating a Better Search UI, on page 61](#), we'll focus on features for searching, sorting, and pagination, to make our main catalog view much more dynamic. We'll also start to unlock some of the true power of Ash, by deriving new attributes with calculations and aggregates.

In [Chapter 4, Generating APIs Without Writing Code, on page 89](#), we'll see the principle of "model your domain, derive the rest" in action, when we learn how to create full REST JSON and GraphQL APIs from our existing resource and action definitions. It's not magic, we swear!

In [Chapter 5, Authentication: Who Are You?, on page 113](#), we'll set up authentication for Tunez, using the AshAuthentication library. We'll cover different strategies for authentication like username/password and logging in via magic link, as well as customizing the autogenerated liveviews to really make them seamless.

In [Chapter 6, Authorization: What Can You Do?, on page 133](#), we'll introduce authorization into the app, using policies and bypasses. We'll see how we can define a policy *once* and use it throughout the entire app, from securing our APIs to showing and hiding UI buttons and more.

In *the (as yet) unwritten Chapter 7, All About Testing*, we'll tackle the topic of testing — what should we test in an app built with Ash, and how should we do it? We'll go over some testing strategies, see what tools Ash provides to help with testing, and cover practical examples of testing Ash and LiveView apps.

In *the (as yet) unwritten Chapter 8, Fun With Nested Forms*, we'll dig a little deeper into Ash's integration with Phoenix, by expanding our domain model and building a nested form, including drag and drop re-ordering for nested records.

And lastly, in *the (as yet) unwritten Chapter 9, PubSub and Real-Time Notifications*, we'll use everything we've learned so far to build a user notification system. Using pubsub for broadcasting real-time updates, and AshOban for making sure batches of notifications are processed smoothly in the back-

ground, we'll create a simple yet robust system that allows for expansion as your apps grow.

Online Resources

All online resources for this book, such as errata and code samples, can be found on the Pragmatic Bookshelf product page:

>> [<<](http://pragprog.com/book/ldash)

We also invite you to join the greater Ash community, if you'd like to learn more, or contribute to the project and ecosystem: [<<](https://ash-hq.org/community)

And on that note, let's dig in! We've got a lot of exciting topics to cover, and can't wait to get started!

Building Our First Resource

Hello! You've arrived! Welcome!!

In this very first chapter, we'll start from scratch and work our way up. We'll set up the starter Tunez application, install Ash, and build our first resource (as the chapter title suggests!). We'll define attributes, set up actions, and connect to a database, all while seeing firsthand how Ash's declarative principles simplify the process. By the end, you'll have a working resource fully integrated with the Phoenix frontend — and the confidence to take the next step.

Getting the Ball Rolling

Throughout this book, we'll build Tunez, a music database app. Think of it as a lightweight Spotify, without actually playing music — users can browse a catalog of artists and albums, star their favorites, and receive notifications when new albums are released. On the management side, we'll implement a role-based access system with customizable permissions, and create APIs that allow users to integrate Tunez data into their own apps.

But Tunez is more than just an app — it's your gateway to mastering Ash's essential building blocks. By building Tunez step by step, you'll gain hands-on experience with resources, relationships, authentication, authorization, APIs, and more. Each feature we build will teach you foundational skills you can apply to any project, giving you the toolkit and know-how to tackle larger, more complex applications with the same techniques. Tunez may be small, but the lessons you'll learn here will have a big impact on your development workflow.

A demo version of the final Tunez app can be seen here:

<https://tunez.sevenseacat.net/>

Setting up your development environment

One of the (many) great things about the Elixir ecosystem is that we get a lot of great new functionality with every new version of Elixir, but nothing gets taken away (at worst, it gets deprecated). So while it would be awesome to always use the latest and greatest versions of everything, sometimes that's not possible, and that's okay! Our apps will still work with most recent versions of Elixir, Erlang, and PostgreSQL.

To work through this book, you'll need at least:

- Elixir 1.15
- Erlang 26.0
- PostgreSQL 14.0
- NodeJS 18.0

Any newer version will also be just fine!

To install these dependencies, we'd recommend a tool like `asdf`¹ or `mise`.²

We've built an initial version of the Tunez app, for you to use as a starting point. To follow along with this book, clone the app from the following repository:

<https://github.com/sevenseacat/tunez>

If you're using `asdf`, once you've cloned the app, you can run `asdf install` from the project folder to get all the language dependencies set up. The `.tool-versions` file in the app lists slightly newer versions than the dependencies listed above, but you can use any versions you prefer as long as they meet the minimum requirements.

Follow the setup instructions in the app `README`, including `npm install --prefix assets` and `mix setup`, to make sure everything is good to go. If you can run `mix phx.server` without errors and see a styled homepage with some sample artist data, you're ready to begin!



The code for each chapter can also be found in the Tunez repo on GitHub, in branches named after the chapter, e.g. the app at the end of chapter 1 can be found at <https://github.com/sevenseacat/tunez/tree/chapter-1>.

1. <https://asdf-vm.com/>
 2. <https://mise.jsx.dev/>

Welcome to Ash!

Before we can start using Ash in Tunez, we'll need to install it and configure it within the app. Tunez is a blank slate — it has a lot of the views and template logic, but no way of storing or reading data. This is where Ash comes in — it will be our main tool for building out the domain model layer of the app, the code responsible for reading and writing data from the database, and implementing our app's business logic.

To install Ash, we'll use the Igniter³ toolkit, which is already installed as a dependency in Tunez. Igniter gives library authors tools to write smarter code generators, including installers, and we'll see that here with the `igniter.install` Mix task.

Run `mix igniter.install ash` in the `tunez` folder, and it will patch the `mix.exs` file with the new package:

```
$ mix igniter.install ash
Fetch Required Dependencies:
Update: mix.exs

...
34 34  |  defp deps do
35 35  |    [
36 + |      {:ash, "~> 3.0"},
36 37  |      {:igniter, "~> 0.3"},
37 38  |      {:phoenix, "~> 1.7.14"},
...

```

These dependencies should be installed before continuing. Modify `mix.exs` and `install? [y/n]`

Confirm the change, and Igniter will install and compile the latest version of the `ash` package. This will trigger Ash's own installation Mix task, which will add Ash-specific formatting configuration in `.formatter.exs` and `config/config.exs`. The output is a little too long to print here, but we'll get consistent code formatting and section ordering across all of the Ash-related modules we'll write over the course of the project.

Starting a new app and want to use Igniter and Ash?



Much like the `phx_new` package is used to generate new Phoenix projects, Igniter has a companion `igniter_new` package for generating projects. You can install it with:

```
$ mix archive.install hex igniter_new
```

3. <https://hexdocs.pm/igniter/>

Starting a new app and want to use Igniter and Ash?

This gives access to the igniter.install⁴ Mix task, which is *very* powerful. It can also combine with phx.new so you can use Igniter to scaffold Phoenix apps that come pre-installed with any package you like (and will also pre-install Igniter). For example, a Phoenix app with Ash and ErrorTracker:

```
$ mix igniter.new your_new_app_name --with phx.new \
--install ash,error_tracker
```

You can also add Igniter and/or Ash to the mix.exs file of any existing app, to get access to all their goodies.

The Ash ecosystem is made up of many different packages for integrations with external libraries and services, allowing us to pick and choose only the dependencies we need. As we're building an app that will talk to a PostgreSQL database, we'll want the PostgreSQL Ash integration. We can use mix igniter.install to add it to Tunez as well:

```
$ mix igniter.install ash_postgres
```

You may see a warning at this point about uncommitted changes detected in your project. Igniter will warn you if your app has changes that haven't been committed into Git. This is to prevent accidental data loss if you've made code changes that Igniter might undo or overwrite. All we've done so far is install Ash so it's fine to continue, but in the future, if you see this message, you may want to back out and commit your code first, just in case.

Confirm the change to our mix.exs file, and the package will be downloaded and installed. After completion, this will:

- Add and fetch the ash_postgres Hex package (in mix.exs and mix.lock)
- Add code auto-formatting for the new dependency (in .formatter.exs and config/config.exs)
- Update the database Tunez.Repo module to use Ash, instead of Ecto (in lib/tunez/repo.ex). This also includes a list of PostgreSQL extensions to be installed and enabled by default
- Generate our first migration to set up the ash-functions pseudo-extension listed in the Tunez.Repo module (in priv/repo/migrations/<timestamp>_install_ash-functions_extension_<version>.exs)
- Generate an extension config file so Ash can keep track of which PostgreSQL extensions have been installed (in priv/resource_snapshots/repo/extensions.json)

4. https://hexdocs.pm/igniter_new/Mix.Tasks.Igniter.New.html

You'll also see a notice from AshPostgres. It has inferred the version of PostgreSQL you're running, and configured that in `Tunez.Repo.min_pg_version/0`.

And now we're good to go and can start building!

Resources and domains

In Ash, the central concept is the *resource*. Resources are domain model objects, the nouns that our app revolves around — they typically (but not always!) contain some kind of data and define some actions that can be taken on that data.

Related resources are grouped together into *domains*⁵ — context boundaries where we can define configuration and functionality that will be shared across all connected resources. This is also where we'll define the interfaces that the rest of the app uses to communicate with the domain model, much like a Phoenix context does.

Domains also provide an integration point for extensions. Some extensions, such as AshAdmin⁶ for automatic admin panels, can be enabled or disabled on a per-domain basis, instead of needing to be configured for every resource within the domain.

What does this mean for Tunez? Over the course of the book, we'll define several different domains for the distinct ideas within the app such as Music and Accounts; and each domain will have a collection of resources such as Album, Artist, Track, User, and Notification.

Each resource will define a set of *attributes* — data that maps to keys of the resource's struct. An Artist resource will read/modify records in the form of Artist structs, and each attribute of the resource will be a key in that struct. The resources will also define *relationships* — links to other resources — as well as actions, validations, pubsub configuration and more.

Do I need multiple domains in my app?

Technically you don't *need* multiple domains. For small apps, you can get away with defining a single domain and putting all of your resources in it, but we want to be really clear about keeping closely-related resources like Album and Artist away from other closely-related resources such as User and Notification.

-
- 5. <https://hexdocs.pm/ash/domains.html>
 - 6. https://github.com/ash-project/ash_admin

In some scenarios, we might even want to have resources with the same name, but different meanings. A Member or a Group resource could represent very different things in a music-related or a user-related context!

We've just thrown a lot of words and concepts at you — some may be familiar to you from other frameworks, others may not. We'll go over each of them as they become relevant to the app, including lots of other resources that can help you out, as well.

Generating the Artist resource

The first resource we'll create is for an Artist. It's the most important resource for anything music-related in Tunetz, that other resources such as albums will link back to. The resource will store information about an artist's name and biography — important so users can know who they're looking at!

To create our Artist resource, we'll use generators to get basic domain and resource modules. We'll generate the database migration to add a database table for storage for our resource, and then we can start fleshing out actions to be taken on our resource.

We'll start off with generating our modules. Ash comes with Igniter generators to create domains and resources for you. We'll use them to save some time, but you could just as well create the necessary files yourself. First we'll generate the Music domain:

```
$ mix ash.gen.domain Tunetz.Music
```

This will create a domain module in `lib/tunetz/music.ex`, as well as adding configuration to load the domain in `config/config.exs`. The domain module is pretty empty right now:

```
01/lib/tunetz/music.ex
defmodule Tunetz.Music do
  use Ash.Domain

  resources do
    end
end
```

We don't need to add anything to this domain yet — for now it's only a bucket to put resources in. As we create more resources relating to music, we'll be adding them to the resources block here. (If you use the resource generator, this will be done for you.)

Now that we have a domain, we can create the resource. The basic resource generator will create a nearly-empty Ash resource, so we can step through it and look through the parts. Run the following in your terminal:

```
$ mix ash.gen.resource Tunez.Music.Artist --extend postgres
```

This will generate a new resource module named `Tunez.Music.Artist` that extends PostgreSQL, and automatically add it as a resource in the `Tunez.Music` domain.

The code for the generated resource is in `lib/tunez/music/artist.ex`:

```
01/lib/tunez/music/artist.ex
defmodule Tunez.Music.Artist do
  use Ash.Resource, otp_app: :tunez, domain: Tunez.Music,
  data_layer: AshPostgres.DataLayer

  postgres do
    table "artists"
    repo Tunez.Repo
  end
end
```

Let's break down this generated code piece by piece, because this is our first introduction to Ash's domain-specific language (DSL).

Because we specified `--extend postgres` when calling the generator, the resource will be configured with PostgreSQL as its data store for reading from and writing to via `AshPostgres.DataLayer`. Each `Artist` struct will be persisted as a row in an artist-related database table.

This specific data layer is configured using the `postgres` code block. The minimum information we need is the `repo` and `table` name, but there is a lot of other behaviour that can be configured⁷ as well.



Ash has several different data layers built in using storage such as Mnesia⁸ and ETS.⁹ More can be added via external packages like we added for PostgreSQL, such as SQLite¹⁰ or CubDB.¹¹

To add attributes to our resource, add another block in the resource named `attributes`. Because we're using PostgreSQL, each attribute we define will be a column in the underlying database table. Ash provides macros we can call

-
7. https://hexdocs.pm/ash_postgres/dsl-ashpostgres-datalayer.html
 8. <https://hexdocs.pm/ash/dsl-ash-datalayer-mnesia.html>
 9. <https://hexdocs.pm/ash/dsl-ash-datalayer-ets.html>
 10. https://hexdocs.pm/ash_sqlite/
 11. https://hexdocs.pm/ash_cubdb/

to define different types of attributes,¹² so let's add some attributes to our resource.

A primary key will be critical to identify our artists, so we can call `uuid_primary_key` to create a autogenerated UUID primary key. Some timestamp fields would be useful, so we know when records are inserted and updated, and we can use `create_timestamp` and `update_timestamp` for those. Specifically for artists, we also know we want to store their `name` and a short `biography`, and they'll both be string values. They can be added to the attributes block using the attribute macro.

```
01/lib/tunez/music/artist.ex
defmodule Tunez.Music.Artist do
  # ...

  attributes do
    uuid_primary_key :id

    attribute :name, :string do
      allow_nil? false
    end

    attribute :biography, :string
    create_timestamp :inserted_at
    update_timestamp :updated_at
  end
end
```

And that's all the code we need to write to add attributes to our resource!



There's a rich set of configuration options for attributes. You can read more about them in the attribute DSL documentation.¹³ We've used one here, `allow_nil?`, but there are many more available.

You can also pass extra options like `--uuid-primary-key id` to the `ash.gen.resource` generator¹⁴ to generate attribute-related code (and more!) if you prefer.

Right now our resource is only a module. We've configured a database table for it, but that database table doesn't yet exist. To change that, we can use another generator. This one we'll get pretty familiar with over the course of the book.

12. <https://hexdocs.pm/ash/dsl-ash-resource.html#attributes>

13. <https://hexdocs.pm/ash/dsl-ash-resource.html#attributes-attribute>

14. <https://hexdocs.pm/ash/Mix.Tasks.Ash.Gen.Resource.html>

Auto-generating database migrations

If you've used Ecto for working with databases before, you'll be familiar with the pattern of creating or updating a schema module, then generating a blank migration and populating it with commands to mirror that schema. It can be a little bit repetitive, and has the possibility of your schema and your database getting out of sync. If someone updates the database structure but *doesn't* update the schema module, or vice versa, you can get some tricky and hard-to-debug issues.

Ash side-steps these kinds of issues by generating complete migrations for you based on your resource definitions. This is our first example of Ash's philosophy of "model your domain, derive the rest". Your resources are the source of truth for what your app should be and how it should behave, and everything else is derived from that.

What does this mean in practice? Every time you run the `ash.codegen mix` task, Ash (via AshPostgres) will:

- Create *snapshots* of your current resources
- Compare them with the previous snapshots (if they exist)
- And finally, generate deltas of the changes to go into the new migration.

This is data-layer agnostic: any data layer can provide its own implementation for what to do when `ash.codegen` is run. Because we're using AshPostgres, which is backed by Ecto, we get Ecto migrations.

Now we have an `Artist` resource with some attributes, so we can generate a migration for it using the mix task:

```
$ mix ash.codegen create_artists
```



The `create_artists` argument given here will become the name of the generated migration module, eg. `Tunez.Repo.Migrations.CreateArtists`. This can be anything, but it's a good idea to describe what the migration will actually do.

Running the `ash.codegen` task will create a few files:

- A snapshot file for our `Artist` resource, in `priv/resource_snapshots/repo/artists/[timestamp].json`. This is a JSON representation of our resource as it exists right now.
- A migration for our `Artist` resource, in `priv/repo/migrations/[timestamp]_create_artists.ex`. This contains the schema differences that Ash has detected between our current snapshot which was just created, and the previous snapshot (which in this case, is empty).

This migration contains the Ecto commands to set up the database table for our Artist resource, with the fields we added for a primary key, timestamps, name and biography:

```
01/priv/repo/migrations/[timestamp]_create_artists.exs
def up do
  create table(:artists, primary_key: false) do
    add :id, :uuid, null: false, default: fragment("gen_random_uuid()"),
        primary_key: true
    add :name, :text, null: false
    add :biography, :text

    add :inserted_at, :utc_datetime_usec,
        null: false,
        default: fragment("(now() AT TIME ZONE 'utc')")

    add :updated_at, :utc_datetime_usec,
        null: false,
        default: fragment("(now() AT TIME ZONE 'utc')")
  end
end
```

This looks a lot like what you would write if you were setting up a database table for a pure Ecto schema — but we didn't have to write it. We don't have to worry about keeping the database structure in sync manually. We can run mix ashcodegen every time we change anything database-related, and Ash will figure out what needs to be changed and create the migration for us.

This is the first time we've touched the database, but the database will already have been created when running mix setup earlier. To run the migration we generated, you can use Ash's ash.migrate Mix task:

```
$ mix ash.migrate
Getting extensions in current project...
Running migration for AshPostgres.DataLayer...

[timestamp] [info] == Running [timestamp] Tunez.Repo.Migrations
  .InstallAshFunctionsExtension[version][timestamp].up/0 forward
<<truncated SQL output>>
[info] == Migrated [timestamp] in 0.0s

[timestamp] [info] == Running [timestamp] Tunez.Repo.Migrations
  .CreateArtists.up/0 forward
[timestamp] [info] create table artists
[timestamp] [info] == Migrated [timestamp] in 0.0s
```

Now we have a database table, ready to store Artist data!



To roll back a migration, Ash also provides an `ash.rollback` Mix task, as well as `ash.setup`, `ash.reset`, and so on. These are more powerful than their Ecto equivalents — any Ash extension can set up their own functionality for each task. For example, `AshPostgres` provides an interactive UI to select how many migrations to roll back when running `ash.rollback`.

Note that if you roll back and then delete a migration to re-generate it, you'll also need to delete the snapshots that were created with the migration.

How do we actually *use* the resource to read or write data into our database, though? We'll need to define some *actions* on our resource.

Oh, CRUD! — Defining Basic Actions

An *action* describes an operation that can be performed for a given resource; it is the *verb* to a resource's *noun*. Actions can be loosely broken down into four types:

- Creating new persisted records (rows in the database table);
- Reading one or more existing records;
- Updating an existing record; and
- Destroying (deleting) an existing record.

These four types of actions are very common in web applications, and are often shortened to the acronym CRUD.



Ash also supports *generic actions* for any action that doesn't fit into any of those four categories. We won't be covering those in this book, but you can read the online documentation¹⁵ about them.

With a bit of creativity, we can use these four basic action types to describe almost any kind of action we might want to perform in an app.

Registering for an account? That's a type of create action on a User resource.

Searching for products to purchase? That sounds like a read action on a Product resource.

Publishing a blog post? It could be a create action if the user is writing the post from scratch, or an update action if they're publishing an existing saved draft.

15. <https://hexdocs.pm/ash/generic-actions.html>

In Tunez, we'll have functionality for users to list artists and view details of a specific artist (both read actions), create and update artist records (via forms), and also destroy artist records; so we'll want to use all four types of actions. This is a great time to learn how to define and run actions using Ash, with some practical examples.

In our Artist resource, we can add an empty block for actions, and then start filling it out with what we want to be able to do:

```
01/lib/tunez/music/artist.ex
defmodule Tunez.Music.Artist do
  # ...
  actions do
  end
end
```

Let's start with creating records with a create action, so we have some data to use when testing out other types of actions.

Defining a create action

Actions are defined by adding them to the actions block in a resource. At their most basic, they require a type (one of the four mentioned earlier — create, read, update and destroy), and a name. The name can be any atom you like, but should describe what the action is actually supposed to do. It's common to give the action the same name as the action type, until you know you need something different.

```
01/lib/tunez/music/artist.ex
actions do
  create :create do
  end
end
```

To create an Artist record, we need to provide the data to be stored — in this case, the name and biography attributes, in a map. (The other attributes, such as timestamps, will be automatically managed by Ash.) We call these the attributes that the action *accepts*, and can list them in the action with the accept macro.

```
01/lib/tunez/music/artist.ex
actions do
  create :create do
    accept [:name, :biography]
  end
end
```

And that's actually all we need to do to create the most basic create action. Ash knows that the core of what a create action should do is create a data layer record from provided data, so that's exactly what it will do when we run it.

Running actions

There are two basic ways we can run actions: the generic query/changeset method, and the more direct code interface method. We can test them both out in an iex session:

```
$ iex -S mix
```

Creating records via a changeset

If you've used Ecto before, this pattern may be familiar to you:

- Create a *changeset* (a set of data changes to apply to the resource)
- Pass that changeset to Ash for processing

In code, this might look like the following:

```
Tunez.Music.Artist
|> Ash.Changeset.for_create(:create, %{
  name: "Valkyrie's Fury",
  biography: "A power metal band hailing from Tallinn, Estonia"
})
|> Ash.create()
```

We specify the action that the changeset should be created for, with the data that we want to save. When we pipe that changeset into Ash, it will handle running all of the validations and creating the record in the database.

```
iex(1)> Tunez.Music.Artist |>
... (1)> Ash.Changeset.for_create(:create, %{
... (1)>   name: "Valkyrie's Fury",
... (1)>   biography: "A power metal band hailing from Tallinn, Estonia"
... (1)> }) |>
... (1)> Ash.create()
{:ok,
 #Tunez.Music.Artist<
  __meta__: #Ecto.Schema.Metadata<:loaded, "artists">,
  id: [uuid],
  name: "Valkyrie's Fury",
  biography: "A power metal band hailing from Tallinn, Estonia",
  ...
>}
```

The record is inserted into the database, and then returned as part of an :ok tuple. You can verify this in your database client of choice, for example, using `psql tunez_dev` in your terminal to connect using the inbuilt command-line client:

```
tunez_dev=# select * from artists;
-[ RECORD 1 ]-----
id          | [uuid]
name        | Valkyrie's Fury
biography   | A power metal band hailing from Tallinn, Estonia
inserted_at | [now]
updated_at  | [now]
```

What happens if we submit invalid data, such as an Artist without a name?

```
iex(2)> Tunez.Music.Artist |>
... (2)> Ash.Changeset.for_create(:create, %{name: ""}) |>
... (2)> Ash.create()
{:error,
%Ash.Error.Invalid{
  changeset: #Ash.Changeset<...>,
  errors: [
    %Ash.Error.Changes.Required{
      field: :name,
      type: :attribute,
      resource: Tunez.Music.Artist,
      ...
    }
  ]
}}
```

The record *isn't* inserted into the database, and we get an error record back telling us what the issue is: the name is required. Later on in this chapter, we'll see how these returned errors are used when we integrate the actions into our web interface.

Like a lot of other Elixir libraries, most Ash functions return data in :ok and :error tuples. This is handy because it lets you easily pattern match on the result, to handle the different scenarios. To raise an error instead of returning an error tuple, you can use the bang version of a function ending in an exclamation mark, ie. `Ash.create!` instead of `Ash.create`.

Creating records via a code interface

If you're familiar with Ruby on Rails or ActiveRecord, this pattern may be more familiar to you. It allows us to skip the step of manually creating a changeset, and lets us call the action directly as a function.

Code interfaces can be defined on either a domain module or on a resource directly. We'd generally recommend defining them on domains, similar to Phoenix contexts, because it lets the domain act as a solid boundary with the

rest of your application. With all your resources listed in your domain, it also gives a great overview of all your functionality in one place.

To enable this, we can use Ash's define macro when including the Artist resource in our Tunez.Music domain:

```
01/lib/tunez/music.ex
resources do
  resource Tunez.Music.Artist do
    define :create_artist, action: :create
  end
end
```

This will connect our domain function `create_artists`, to the `create` action of the resource. Once you've done this, if you recompile within iex the new function will now be available, complete with auto-generated documentation:

```
iex(2)> h Tunez.Music.create_artist
def create_artist(params_or_opts \\ %{}, opts \\ [])
Calls the create action on Tunez.Music.Artist.
```

You can call it like any other function, with the data to be inserted into the database:

```
iex(7)> Tunez.Music.create_artist(%{
... (7)>   name: "Valkyrie's Fury",
... (7)>   biography: "A power metal band hailing from Tallinn, Estonia"
... (7)> })
{:ok, #Tunez.Music.Artist<...>}
```

When would I use changesets instead of code interfaces, or vice versa?

Under the hood, the code interface is creating a changeset and passing it to the domain, but that repetitive logic is hidden away. So there's no real functional benefit, but the code interface is easier to use and more readable.

Where the changeset method shines is around forms on the page — we'll see shortly how AshPhoenix provides a thin layer over the top of changesets to allow all of Phoenix's existing form helpers to work seamlessly with Ash changesets instead of Ecto changesets.

We've provided some sample content for you to play around with — you can import it into your database by running the following on the command line:

```
$ mix run priv/repo/seeds/01-artists.exs
```

There are other seed files in there too, but we'll mention those when we get to them!

Now that we have some data in our database, we can look at other types of actions.

Defining a read action

In the same way we defined the create action on the Artist resource, we can define a read action by adding it to the actions block. We'll add just one extra option: we'll define it as a *primary action*.

```
01/lib/tunez/music/artist.ex
actions do
  # ...
  read :read do
    primary? true
  end
end
```

A resource can have one of each of the four action types (create, read, update and destroy) marked as the primary action of that type. These are used by Ash behind the scenes when actions aren't or *can't* be specified. We'll cover these in a little bit more detail later.

To be able to call the read action as a function, add it as a code interface in the domain just like we did with create.

```
01/lib/tunez/music.ex
resource Tunez.Music.Artist do
  # ...
  define :read_artists, action: :read
end
```

What does a read action *do*? As the name suggests, it will read data from our data layer based on any parameters we provide. We haven't defined any parameters in our action, so when we call the action, we should expect it to return all the records in the database.

```
iex(1)> Tunez.Music.read_artists()
{:ok,
 [
  #Tunez.Music.Artist<...>,
  #Tunez.Music.Artist<...>,
  ...
]}
```

While actions that modify data use changesets under the hood, read actions use *queries*. If we *do* want to provide some parameters to the action, such as filtering or sorting, we need to modify the query when we run the action. This can be done either as part of the action definition (we'll learn about that in [Designing a search action, on page 61!](#)), or inline when we call the action.

Manually reading records via a query

While this isn't something you'll do a lot of when building applications, it's still a good way of seeing how Ash builds up queries piece by piece.

There are a few steps to the process:

- Creating a basic query from the action we want to run
- Piping the query through other functions to add any extra parameters we want, and then
- Passing the final query to Ash for processing.

In iex you can test this out step by step, starting from the basic resource, and creating the query:

```
iex(2)> Tunez.Music.Artist
Tunez.Music.Artist
iex(3)> |> Ash.Query.for_read(:read)
#Ash.Query<resource: Tunez.Music.Artist>
```

Then you can pipe that query into Ash's query functions like `sort` and `limit`. The query keeps getting the extra conditions added to it, but it isn't yet being run in the database.

```
iex(4)> |> Ash.Query.sort(name: :asc)
#Ash.Query<resource: Tunez.Music.Artist, sort: [name: :asc]>
iex(5)> |> Ash.Query.limit(1)
#Ash.Query<resource: Tunez.Music.Artist, sort: [name: :asc], limit: 1>
```

Then when it's time to go, Ash can call it and return you the data you requested, with all conditions applied:

```
iex(6)> |> Ash.read()
[debug] QUERY OK source="artists" db=3.4ms queue=1.2ms idle=1863.7ms
SELECT a0."id", a0."name", a0."biography", a0."inserted_at", a0."updated_at"
FROM "artists" AS a0 ORDER BY a0."name" LIMIT $1 [1]
{:ok, [#Tunez.Music.Artist<...>]}
```

For a full list of the query functions Ash provides, you can check out the documentation.¹⁶ Note that to use any of the functions that use special syntax like `filter`, you'll need to require `Ash.Query` in your iex session first.

Reading a single record by identity

One very common requirement is to be able to read a single record by its primary key. We're building a music app, so we'll be building a page where

16. <https://hexdocs.pm/ash/Ash.Query.html>

we can view an artist's profile, and we'll want an easy way to fetch that single Artist record for display.

We have a basic read action already, and we *could* write another read action that applies a filter to only fetch the data by an ID we provide, but Ash provides a simpler way.

A neat feature of code interfaces is that we can automatically apply a filter for any *identity* of a resource. Ash defines an identity as any attribute (or combination of attributes) that can be used to uniquely identify a record — one is automatically generated for the primary key attribute, but we can also add others, such as for a user's email address.

To use this feature, we can add another code interface for the same read action, but adding the `get_by` option¹⁷ for the primary key, an attribute named `:id`:

```
01/lib/tunez/music.ex
resource Tunez.Music.Artist do
  # ...
  define :get_artist_by_id, action: :read, get_by: :id
end
```

Adding this code interface defines a new function on our domain:

```
iex(4)> h Tunez.Music.get_artist_by_id
           def get_artist_by_id(id, params_or_opts \\ %{}, opts \\ [])
Calls the read action on Tunez.Music.Artist.
```

Copy the ID from any of the records you loaded when testing the read action, and you'll see that this new function does exactly what we hoped — return the single record that has that ID.

```
iex(3)> Tunez.Music.get_artist_by_id("an-artist-id")
[debug] QUERY OK source="artists" db=0.3ms queue=0.4ms idle=331.4ms
SELECT a0."id", a0."name", a0."biography", a0."inserted_at", a0."updated_at"
FROM "artists" AS a0 WHERE (a0."id"::uuid = $1::uuid)
["an-artist-id"]
{:ok, #Tunez.Music.Artist<id: "an-artist-id", ...>}
```

Perfect! We'll be using that very soon.

Defining an update action

A basic update action is conceptually very similar to a create action. The main difference is that instead of building a new record with some provided data

17. https://hexdocs.pm/ash/dsl-ash-domain.html#resources-resource-define-get_by

and saving it into the database, we provide an existing record to be updated with the data and saved.

Let's add the basic action and code interface definition:

```
01/lib/tunez/music/artist.ex
actions do
  # ...
  update :update do
    accept [:name, :biography]
  end
end
```

```
01/lib/tunez/music.ex
resource Tunez.Music.Artist do
  # ...
  define :update_artist, action: :update
end
```

How would we call this new action? First we need a record to be updated. You can use the read action you defined earlier to find one, or if you've been testing the `get_artist_by_id` function we just wrote, you might have one right there.

```
iex(3)> artist = Tunez.Music.get_artist_by_id!("an-artist-id")
#Tunez.Music.Artist<id: "an-artist-id", ...>
```

Now we can either use the code interface we added or create a changeset and apply it, like we did for create.

```
# Via the code interface
iex(4)> Tunez.Music.update_artist(artist, %{name: "Hello"})
[debug] QUERY OK source="artists" db=1.4ms
UPDATE "artists" AS a0 SET "name" = $1, "updated_at" = $2 WHERE
(a0."id" = $3) RETURNING a0."name", a0."updated_at", a0."id" ["Hello",
[now], "an-artist-id"]
{:ok, #Tunez.Music.Artist<id: "an-artist-id", name: "Hello", ...>}

# Or via a changeset
iex(7)> artist |>
... (7)> Ash.Changeset.for_update(:update, %{name: "World"}) |>
... (7)> Ash.update()
«an almost-identical SQL statement»
{:ok, #Tunez.Music.Artist<id: "an-artist-id", name: "World", ...>}
```

Like create actions, we get either an `:ok` tuple with the updated record, or an `:error` tuple with an error record back.

Defining a destroy action

The last type of core action that Ash provides are destroy actions — when we want to get rid of data, delete it from our database and memories. Like update actions, they work on existing data records, so we need to provide a record when we call a destroy action... but that's the only thing we need to provide. Ash can do the rest!

You might be able to guess by now, how to implement a destroy action in our resource:

```
01/lib/tunez/music/artist.ex
actions do
  # ...
  destroy :destroy do
    end
end
```

```
01/lib/tunez/music.ex
resource Tunez.Music.Artist do
  # ...
  define :destroy_artist, action: :destroy
end
```

This will allow us to call the action, either by creating a changeset and submitting it or calling the action directly.

```
iex(3)> {:ok, artist} = Tunez.Music.get_artist_by_id("the-artist-id")
{:ok, #Tunez.Music.Artist<id: "an-artist-id", ...>}

# Via the code interface
iex(4)> Tunez.Music.destroy_artist(artist)
[debug] QUERY OK source="artists" db=0.3ms
DELETE FROM "artists" AS a0 WHERE (a0."id" = $1) ["the-artist-id"]
:ok

# Or via a changeset
iex(5)> artist |>
... (5)> Ash.Changeset.for_destroy(:destroy) |>
... (5)> Ash.destroy()
[debug] QUERY OK source="artists" db=1.1ms
DELETE FROM "artists" AS a0 WHERE (a0."id" = $1) ["the-artist-id"]
:ok
```

And with that, we have a solid explanation of the four main types of actions we can define in our resources.

Let's take a bit to let this really sink in. By creating a resource and adding a few lines of code to describe its attributes and what actions we can take on it, we now have:

- A database table to store records in
- Secure functions we can call to read and write data to the database (no storing any attributes that aren't explicitly allowed!)
- Database-level validations to ensure that data is present
- Automatic type-casting of attributes before they get stored

We didn't have to write any functions that query the database, or update our database schema when we added new attributes to the resource, or manually cast attributes. A lot of the boilerplate we would typically need to write has been taken care of for us because Ash handles translating *what* our resource should do into *how* it should be done. This is a pattern we'll see a lot!

Default actions

Now that you've learned about the different types of actions and how to define them in your resources, we can let you in on a little secret...

You don't actually *have* to define empty actions like this for CRUD actions.

You know how we said that Ash knows that the main purpose of a create action is to take the data and save it to the data layer? This is what we call the *default* implementation for a create action. Ash provides default implementations for all four action types, and if you want to use these implementations without any customization, you can use the `defaults` macro in your actions block:

```
01/lib/tunez/music/artist.ex
actions do
  defaults [:create, :read, :update, :destroy]
end
```

We still need the code interface definitions if we want to be able to call the actions as functions, but we can cut out the empty actions to save time and space. This also marks all four actions as primary? true, as a handy side-effect.

But what about the accept definitions that we added to the create and update actions, the list of attributes to save? We can define default values for that list, with the `default_accept`¹⁸ macro. This default list will then apply to all create and update actions unless specified otherwise (as part of the action definition).

So the actions for our whole resource as it stands right now could be written in a few short lines of code:

```
01/lib/tunez/music/artist.ex
actions do
  defaults [:create, :read, :update, :destroy]
```

18. https://hexdocs.pm/ash/dsl-ash-resource.html#actions-default_accept

```
  default_accept [:name, :biography]
end
```

That is a *lot* of functionality packed into those four lines!

Which version of the code you write is up to you — we would generally err on the side of explicitly defining actions with the attributes they accept, as you'll need to convert them whenever you need to add business logic to your actions anyway. For quick prototyping though, the shorthand can't be beat. Whichever way you go, it's critical to know what your code is doing for you, under the hood — generating a full CRUD interface to your resource, allowing you to manage your data.

Integrating Actions into LiveViews

We've talked a lot about Tunez the web app, but we haven't even *looked* at the app in an actual browser yet. Now that we have a fully-functioning resource, let's integrate it into the web interface so we can see the actions in, well, action!

Listing artists

In the app folder, you can start the Phoenix webserver with the following command in your terminal:

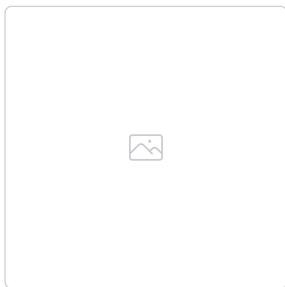
```
$ mix phx.server
[info] Running TunezWeb.Endpoint with Bandit 1.3.0 at 0.0.0.0:4000 (http)
[info] Access TunezWeb.Endpoint at http://localhost:4000
[watch] build finished, watching for changes...
Rebuilding...
Done in [time]ms.
```

Once you see that the build is ready to go, you can open a web browser at <http://localhost:4000> and see what we've got to work with.

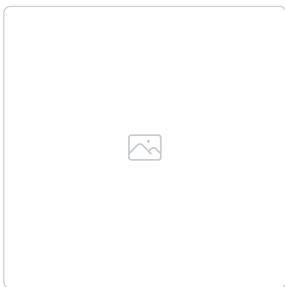


Artists

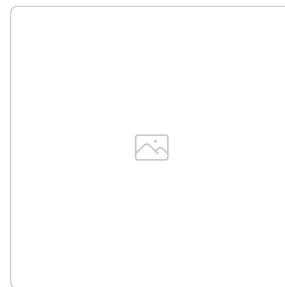
New Artist



Test Artist 1



Test Artist 2



Test Artist 3

The app homepage is the artist catalog, listing all of the Artists in the app. In the code, this catalog is rendered by the `TunezWeb.Artists.IndexLive` module, in `lib/tunez_web/live/artists/index_live.ex`.



We're not going to go into a lot of detail about Phoenix and Phoenix LiveView, apart from where we need to to ensure that our app is secure. If you need a refresher course (or want to learn them for the first time!), we can recommend reading through [Programming Phoenix 1.4 \[TV19\]](#) and [Programming Phoenix LiveView \[TD24\]](#).

If you're more of a video person, you can't go past Pragmatic Studio's Phoenix LiveView course.¹⁹

In the `IndexLive` module, we have some hardcoded maps of artist data defined in the `handle_params/3` function:

`01/lib/tunez_web/live/artists/index_live.ex`

```
def handle_params(_params, _url, socket) do
  artists = [
    %{id: "test-artist-1", name: "Test Artist 1"},
    %{id: "test-artist-2", name: "Test Artist 2"},
    %{id: "test-artist-3", name: "Test Artist 3"},
  ]
  socket =
    socket
    |> assign(:artists, artists)
```

19. <https://pragmaticstudio.com/courses/phoenix-liveview>

```
{:noreply, socket}
end
```

These are what are iterated over in the `render/1` function, using a function component to render an artist “card” for each artist — the image placeholders and names we can see in the browser.

```
01/lib/tunez_web/live/artists/index_live.ex
<li :for={artist <- @artists}>
  <.artist_card artist={artist} />
</li>
```

Earlier in [Defining a read action, on page 16](#), we defined a code interface function on the `Tunez.Music` domain for reading records from the database. It returns `Artist` structs that have a `name` key, just like the hardcoded data does. So to load real data from the database, we can replace the hardcoded data with a call to the `read` action.

```
01/lib/tunez_web/live/artists/index_live.ex
def handle_params(_params, _url, socket) do
  artists = Tunez.Music.read_artists!()
  # ...
```

And that’s it! The page should reload in your browser when you save the changes to the liveview, and you should be seeing the names of the test artists you created, rendered on the page.

Each of the placeholder images and names links to a separate profile page, where you can view details of a specific artist, and we’ll connect that up next.

Viewing an artist profile

Clicking on the name of one of the artists will bring you to their profile page. This liveview is defined in `TunezWeb.Artists.ShowLive`, which you can verify by checking the logs of the web server in your terminal:

```
[debug] MOUNT TunezWeb.Artists.ShowLive
Parameters: %{"id" => "[the artist UUID]"}
```

Inside that module, in `lib/tunez_web/live/artists/show_live.ex`, you’ll again see some hardcoded artist data defined in the `handle_params/3` function and added to the socket.

```
01/lib/tunez_web/live/artists/show_live.ex
def handle_params(_params, _url, socket) do
  artist = %{
    id: "test-artist-1",
    name: "Artist Name",
    biography: <<some sample biography content>>
```

```
}
```

```
# ...
```

Earlier in [Reading a single record by identity, on page 17](#), we defined a `get_artist_by_id` code interface function on the `Tunez.Music` domain, that reads a single Artist record from the database by its `id` attribute. The URL for the profile page contains the ID of the Artist to show on the page, and the terminal logs showed that the ID is available as part of the params. So we can replace the hardcoded data with a call to `get_by_id`, after first using pattern matching to get the ID from the params.

```
01/lib/tunez_web/live/artists/show_live.ex
```

```
def handle_params(%{"id" => artist_id}, _url, socket) do
  artist = Tunez.Music.get_artist_by_id!(artist_id)
  # ...
```

After saving the changes to the liveview, the page should refresh and you should see the correct data for the artist you're viewing the profile of.

Creating artists with AshPhoenix.Form

To create and edit Artist data, we'll have to learn how to handle forms and form data with Ash.

If we were building our app directly on Phoenix contexts using Ecto, we would have a schema module that would define the attributes for an Artist. The schema module would also define a changeset function to parse and validate data from forms before the context module would attempt to insert or update it in the database. If the data validation fails, the liveview (or a template) can take the resulting changeset with errors on it and use it to show the user what they need to fix.

In code, the changeset function might look something like this:

```
defmodule Tunez.Music.Artist do
  def changeset(artist, attrs) do
    artist
    |> cast(attrs, [:name, :biography])
    |> validate_required([:name])
  end
```

And the context module that uses it might look like this:

```
defmodule Tunez.Music do
  def create_artist(attrs \\ %{}) do
    %Artist{}
    |> Artist.changeset(attrs)
    |> Repo.insert()
  end
```

We can use a similar pattern with Ash, but we need a slightly different abstraction.

We have our Artist resource defined with attributes, similar to an Ecto schema. It has actions to create and update data, replacing the context part as well. What we're missing is the integration with our UI, a way to take the errors returned if the create action fails and show them to the user — and this is where AshPhoenix comes in.

Hello, AshPhoenix

As the name suggests, AshPhoenix is a core Ash library to make it much nicer to work with Ash in the context of a Phoenix application. We'll use it a few times over the course of building Tunez, but its main purpose is *form integration*.

Like AshPostgres and Ash itself, we can use `mix igniter.install` to install AshPhoenix in a terminal:

```
$ mix igniter.install ash_phoenix
```

Confirm the addition of AshPhoenix to your `mix.exs` file, and the package will be installed and we can start using it straight away.

A form for an action

Our Artist resource has a create action that accepts data for name and biography attributes. Our web interface will reflect this exactly — we'll have a form with a text field to enter a name, and a textarea to enter a biography.

We can tell AshPhoenix that what we want is a form to match the inputs for our create action, or more simply, a form *to* our create action. AshPhoenix will return an `AshPhoenix.Form` struct, and provides a set of intuitively-named functions for interacting with it. We can *validate* our form, we can *submit* our form, we can *add* and *remove* forms (for nested form data!) and more.

In an iex session, we can get familiar with using `AshPhoenix.Form`:

```
iex(1)> form = AshPhoenix.Form.for_create(Tunez.Music.Artist, :create)
#AshPhoenix.Form<
  resource: Tunez.Music.Artist,
  action: :create,
  type: :create,
  params: %{},
  source: #Ash.Changeset<
    domain: Tunez.Music,
    action_type: :create,
    action: :create,
    attributes: %{},
```

```
relationships: %{},  
...
```

An `AshPhoenix.Form` wraps an `Ash.Changeset`, which behaves similarly to an `Ecto.Changeset`. This allows the `AshPhoenix` form to be a drop-in replacement for an `Ecto.Changeset`, when calling the function components that Phoenix generates for dealing with forms. Let's keep testing.

```
iex(2)> AshPhoenix.Form.validate(form, %{name: "Best Band Ever"})  
#AshPhoenix.Form<  
  resource: Tunez.Music.Artist,  
  action: :create,  
  type: :create,  
  params: %{name: "Best Band Ever"},  
  source: #Ash.Changeset<  
    domain: Tunez.Music,  
    action_type: :create,  
    action: :create,  
    attributes: %{name: "Best Band Ever"},  
    relationships: %{},  
    errors: [],  
    data: #Tunez.Music.Artist<...>  
  ▶   valid?: true  
>,  
...
```

If we call `AshPhoenix.Form.validate` with valid data for an Artist, the changeset in the form is now valid. In a liveview, this is what we would call in a `phx-change` event handler, to make sure our form in-memory stays up-to-date with the latest data. Similarly, we can call `AshPhoenix.Form.submit` on the form in a `phx-submit` event handler.

```
iex(5)> AshPhoenix.Form.submit(form, params: %{name: "Best Band Ever"})  
...  
[debug] QUERY OK source="artists" db=6.3ms  
INSERT INTO "artists" ("id", "name", "inserted_at", "updated_at") VALUES  
($1, $2, $3, $4) RETURNING "updated_at", "inserted_at", "biography", "name", "id"  
[[uuid], "Best Band Ever", [timestamp], [timestamp]]  
↳ AshPostgres.DataLayer.bulk_create/3, at: lib/data_layer.ex:1244  
...  
{:ok,  
 #Tunez.Music.Artist<  
  __meta__: #Ecto.Schema.Metadata<:loaded, "artists">,  
  id: [uuid],  
  name: "Best Band Ever",  
  ...}
```

And it works! We get back a form-ready version of the return value of the action. If we had called submit with invalid data, we would get back an `{:error, %AshPhoenix.Form{}}` tuple back instead.

Using the AshPhoenix domain extension

We've defined code interface functions like `Tunez.Music.read_artists` for all of the actions in the Artist resource and used those code interfaces in our liveviews. It might feel a bit odd to now revert back to using action names directly when generating forms. And if the names of the code interface function and the action are really different, it could get a bit confusing!

AshPhoenix provides a solution for this, with a *domain extension*. If we add the AshPhoenix extension to the `Tunez.Music` domain module, it will define some new functions on the domain around form generation.

In the `Tunez.Music` module in `lib/tunez/music.ex`, add a new extensions option to the `use Ash.Domain` line:

```
01/lib/tunez/music.ex
defmodule Tunez.Music do
  >   use Ash.Domain, extensions: [AshPhoenix]
      # ...
```

Now, instead of calling `AshPhoenix.Form.for_create(Tunez.Music.Artist, :create)`, we can use a new function `Tunez.Music.form_to_create_artist`. This works for any code interface function, even for read actions, by prefixing `form_to_` to the function name.

```
iex(5)> AshPhoenix.Form.for_create(Tunez.Music.Artist, :create)
#AshPhoenix.Form<resource: Tunez.Music.Artist, action: :create, ...>
iex(6)> Tunez.Music.form_to_create_artist()
#AshPhoenix.Form<resource: Tunez.Music.Artist, action: :create, ...>
```

The result is the same — you get an `AshPhoenix.Form` struct to validate and submit, as before — but the way you get it is a lot more consistent with other function calls.

Integrating a form into a liveview

The liveview for creating an Artist is the `TunezWeb.Artists.FormLive` module, located in `lib/tunez_web/live/artists/form_live.ex`. In the browser, you can view it by clicking the New Artist button on the artist catalog, or visiting `/artists/new`.



New Artist

Name

Biography

Save

It looks good, but it's totally non-functional right now. We can use what we've learned so far about AshPhoenix.Form to make it work as we would expect.

It starts from the top — we want to build our initial form in the mount/3 function. Currently form is defined as an empty map, just to get the form to render. We can replace it with a function call to create the form, like we did in iex:

```
01/lib/tunez_web/live/artists/form_live.ex
def mount(_params, _session, socket) do
  >> form = Tunez.Music.form_to_create_artist()

  socket =
    socket
    |> assign(:form, to_form(form))
    # ...
```

The form has a phx-validate event handler attached that will fire after every pause in typing on the form. This will send the validate event to the liveview, handled by the handle_event/3 function head with the first argument "validate".

```
01/lib/tunez_web/live/artists/form_live.ex
def handle_event("validate", %{"form" => _form_data}, socket) do
  {:noreply, socket}
end
```

It doesn't currently do anything, but we know we need to update the form in the socket with the data from the form.

```
01/lib/tunex_web/live/artists/form_live.ex
def handle_event("validate", %{"form" => form_data}, socket) do
  socket =
    update(socket, :form, fn form =>
      AshPhoenix.Form.validate(form, form_data)
    end)

  {:noreply, socket}
end
```

Lastly, we need to deal with form submission. The form has a phx-submit event handler attached that will fire when the user presses the Save button (or presses Enter). This will send the save event to the liveview. The event handler currently doesn't do anything either (we told you the form was non-functional!) but we can add code to submit the form with the form data.

We also need to handle the response after submission, handling both the success and failure cases. If the user submits invalid data then we want to show errors, otherwise we can go back to the catalog and display a success message.

```
01/lib/tunex_web/live/artists/form_live.ex
def handle_event("save", %{"form" => form_data}, socket) do
  case AshPhoenix.Form.submit(socket.assigns.form, params: form_data) do
    {:ok, _artist} ->
      socket =
        socket
        |> put_flash(:info, "Artist saved successfully")
        |> push_navigate(to: ~p"/")

      {:noreply, socket}

    {:error, form} ->
      socket =
        socket
        |> put_flash(:error, "Could not save artist data")
        |> assign(:form, form)

      {:noreply, socket}
  end
end
```

Give it a try! Submit some invalid data, see the validation errors, correct the data, and submit the form again. It works great!

But what happens if you make a typo when entering data? No-one wants to read about Metlalica, do they? We need some way of editing Artist records, and updating any necessary information.

Updating artists with the same code

When we set up the update actions in our Artist resource in [Defining an update action, on page 18](#), we noted that it was pretty similar to the create action and that the only real difference for update is that we need to provide the record being updated. The rest of the flow — providing data to be saved, saving it to the database — is exactly the same.

In addition, the web interface for editing an artist should be exactly the same as for creating an artist. The only difference will be that the form for editing has the artist data pre-populated on it, so that it can be modified, and the form for creating will be totally blank.

We can actually use the same `TunezWeb.Artists.FormLive` liveview module for both creating and updating records. The routes are already set up for this: clicking the Edit Artist button on the profile page will take you to that liveview.

```
[debug] MOUNT TunezWeb.Artists.FormLive
Parameters: %{"id" => "[the artist UUID]"}
```



This won't be the case for all resources, all the time. You may need very different interfaces for creating and updating data. A lot of the time, though, this can be a neat way of building out functionality very quickly, and it can be changed later if your needs change.

The `FormLive` liveview will need to have different forms, depending on if an artist is being created or updated. Everything else can be the same because we still want to validate the data on keystroke, submit the form on form submission, and perform the same actions after submission.

As we build the form for create in the `mount/3` function, we can add another `mount/3` function head specifically for update, which sets `form` in the socket to a form built for a different action.

```
01/lib/tunez_web/live/artists/form_live.ex
def mount(%{"id" => artist_id}, _session, socket) do
  artist = Tunez.Music.get_artist_by_id!(artist_id)
  form = Tunez.Music.form_to_update_artist(artist)

  socket =
    socket
    |> assign(:form, to_form(form))
    |> assign(:page_title, "Update Artist")

  {:ok, socket}
end

def mount(_params, _session, socket) do
  form = Tunez.Music.form_to_create_artist()
```

```
# ...
```

This new function head (which has to come *before* the existing function head!) is differentiated by having an artist id in the params, just like the ShowLive module did when we viewed the artist profile. It sets up the form specifically for the update action of the resource, using the loaded Artist record as the first argument. It also sets a different page title, and... that's all that has to change! Everything else should keep behaving exactly the same.

Save the liveview and test it out in your browser. You should now be able to click Edit Artist, update the artist's details, save, and see the changes reflected back in the artist catalog.

Deleting artist data

The last action we need to integrate is the destroy_artist code interface function for removing records from the database. In the UI, this is done from a button at the top of the artist profile page, next to the edit button. The button, located in the template for TunezWeb.Artists.ShowLive, will send the destroy_artist event when pressed.

```
01/lib/tunez_web/live/artists/show_live.ex
<.button_link kind="error" text phx-click="destroy_artist"
  data-confirm={"Are you sure you want to delete #{@artist.name}?"}>
  Delete Artist
</.button_link>
```

We've already loaded the artist record from the database when rendering the page, and stored it in socket.assigns, so you can fetch it out again and attempt to delete it with the Music.destroy_artist function. The error return value would probably never be seen in practice, but just in case, we'll show the user a nice message anyway.

```
01/lib/tunez_web/live/artists/show_live.ex
def handle_event("destroy_artist", _params, socket) do
  case Tunez.Music.destroy_artist(socket.assigns.artist) do
    :ok ->
      socket =
        socket
        |> put_flash(:info, "Artist deleted successfully")
        |> push_navigate(to: ~p"/")
    {:noreply, socket}
    {:error, error} ->
      Logger.info("Could not delete artist '#{socket.assigns.artist.id}':
      #{inspect(error)})"
    socket =
```

```
    socket
      |> put_flash(:error, "Could not delete artist")
      {:noreply, socket}
  end
end
```

There are lots of different stylistic ways that this type of code could be written, but this is typically the way we would write it. If things go wrong, we want errors logged as to what happened, and we always want users to get feedback about what's going on.

And that's it! We've set up Ash in the Tunez app, and implemented a full CRUD interface for our first resource — and we haven't had to write very much code to do it.

We've learned a bit about the *declarative* nature of Ash. We didn't need to write functions that accepted parameters, processed them, saved the records, etc — we haven't needed to write any functions at all. We declared what our resource should look like, where data should be stored, and what our actions should do. Ash has handled the actual implementations for us.

We've also seen how AshPhoenix provides a tidy Form pattern for integration with web forms, allowing for a really streamlined integration with very little code.

In the next chapter, we'll look at building a second resource and how the two can be integrated together!

Extending Resources with Business Logic

In the first chapter, we learned how to set up Ash within our Phoenix app, created our first resource for Artists within a domain, and built out a full web interface so that we could create, read, update and delete Artist records. This would be a great starting point for any application, to pick your most core domain model concept and build it out.

Now we can start fleshing out the domain model for Tunez a little bit more, because one resource does not a full application make. Having multiple resources, and connecting them together, will allow us to do things like querying and filtering based on related data. So, in the real world, artists release albums, right? Let's build a second resource representing an Album with a more complex structure, link them together, and learn some other handy features of working with declarative resources.

Resources and Relationships

Like we generated our Artist resource, we can start by using Ash's generators to create our basic Album resource. It's music-related, so it should also be part of the `Tunez.Music` domain:

```
$ mix ash.gen.resource Tunez.Music.Album --extend postgres
```

This will generate the resource file in `lib/tunez/music/album.ex`, as well as adding the new resource to the list of resources in the `Tunez.Music` domain module.

The next step, just like when we built our first resource, is to consider what kinds of attributes our new resource needs. What information should we record about an Album? Right now we probably care about:

- The artist who released the album
- The album name
- The year the album was released

- And an image of the album cover. That'll make Tunez look really nice!

Ash has a lot of inbuilt data types,¹ that can let you model just about anything. If we were building a resource for a product in a clothing store, we might want attributes for things like the item size, colour, brand name, and price. A listing on a real estate app might want to store the property address, the number of bedrooms and bathrooms, and the property size.



If none of the inbuilt data types cover what you need, you can also create custom or composite data types.² These can neatly wrap logic around discrete units of data, such as phone numbers, URLs, or latitude/longitude co-ordinates.

In the attributes block of the Album resource, we can start adding our new attributes:

```
02/lib/tunez/music/album.ex
attributes do
  uuid_primary_key :id

  attribute :name, :string do
    allow_nil? false
  end

  attribute :year_released, :integer do
    allow_nil? false
  end

  attribute :cover_image_url, :string
  create_timestamp :inserted_at
  update_timestamp :updated_at
end
```

The name and year_released attributes will be required, but the cover_image_url will be optional. We might not have a high-quality photo on hand for every album, but we can add them in later when we get them.

We haven't added any field to represent the artist though — and that's because it's not going to be just a normal attribute. It's going to be a *relationship*.

Defining Relationships

Relationships, also known as *associations*, are how we describe connections between resources in Ash. There are a couple of different relationship types

1. <https://hexdocs.pm/ash/Ash.Type.html#module-built-in-types>
 2. <https://hexdocs.pm/ash/Ash.Type.html#module-defining-custom-types>

we can choose from, based on the numbers of resources involved on each side:

- `has_many` relationships relate *one* resource to *many* other resources. These are really common — a User can have many Posts, or a Book can have many Chapters. These don't store any data on the *one* side of the relationship, but each of the items on the *many* side will have a reference back to the *one*.
- `belongs_to` relationships relate *one* resource to *one* parent/containing resource. They're usually the inverse of a `has_many` — in the examples above, the resource on the *many* side would typically belong to the *one* resource. A Chapter belongs to a Book, and a Post belongs to a User. The resource belonging to another will have a reference to the related resource, eg. a Chapter will have a `book_id` attribute, referencing the `id` field of the Book resource.
- `has_one` relationships are less common, but are similar to `belongs_to` relationships. They relate *one* resource to *one* other resource, but differ in which end of the relationship holds the reference to the related record. For a `has_one` relationship, the related resource will have the reference. A common example of a `has_one` relationship is Users and Profiles — a User could have one Profile, but the Profile resource is what holds a `user_id` attribute.
- `many_to_many` relationships, as the name suggests, relate many resources to many other resources. These are where you have two pools of different objects, and can link any two resources between the pools. Tags are a really common example — a Post can have many Tags applied to it, and a Tag can also apply to many different Posts.

In our case, we'll be using `belongs_to` and `has_many` relationships — an Artist `has_many` albums, and an Album `belongs_to` an artist.

In code, we define these in a separate top-level relationships block in each resource. In the Artist resource, we can add a relationship with Albums:

```
02/lib/tunez/music/artist.ex
relationships do
  has_many :albums, Tunez.Music.Album
end
```

And in the Album resource, a relationship back to the Artist resource:

```
02/lib/tunez/music/album.ex
relationships do
  belongs_to :artist, Tunez.Music.Artist do
```

```

    allow_nil? false
  end
end

```

Now that our resource is set up, we can generate a database migration for it, using the `ashcodegen` mix task.

```
$ mix ashcodegen create_albums
```

This will generate a new Ecto migration in `priv/repo/migrations/[timestamp]_create_albums.exs` to create the `albums` table in the database, including a foreign key representing the relationship — this will link an `artist_id` field on the `albums` table to the `id` field on `artists` table. A snapshot JSON file will also be created, representing the current state of the `Album` resource.

If you're happy with the migration, run it:

```
$ mix ash.migrate
```

And now we can start adding functionality. A lot of this will seem pretty familiar from building out the `Artist` interface, so we'll cover it quickly. There are a few new interesting parts, however, due to the added relationship, so let's dig right in.

Album actions

If we look at an `Artist`'s profile page in the app, we can see a list of their albums, so we're going to need some kind of `read` action on the `Album` resource, to read the data to display. There's also a button to add a new album, at the top of the album list, so we'll need a `create` action; and each album has `Edit` and `Delete` buttons next to the title, so we'll write some `update` and `destroy` actions as well.

We can add those to the `Album` resource pretty quickly:

```
02/lib/tunez/music/album.ex
actions do
  defaults [:read, :destroy]

  create :create do
    accept [:name, :year_released, :cover_image_url, :artist_id]
  end

  update :update do
    accept [:name, :year_released, :cover_image_url]
  end
end
```

We don't have any customizations to make to the default implementation of `read` or `destroy`, so we can define those as default actions. You might be thinking,

but won't we need to customize the read action, to only show albums for a specific artist? We actually don't! When we load an artist's albums on their profile page, which we'll see how to do shortly, we won't be calling this action directly — we'll be asking Ash to load the albums through the albums relationship on the Artist resource, which will automatically apply the correct filter.

We do have tweaks for the create and update actions, though — for the accept list, of attributes that can be set when calling those actions. When creating a record, it makes sense to need to set the `artist_id` for an album, otherwise it won't be set at all! When updating an album, though, does it really need to be changeable? Can we see ourselves creating an album via the wrong artist profile and then needing to change it later? It seems unlikely, so we don't need to accept the `artist_id` attribute in the update action.

We'll also add code interface definitions for our actions, to make them easier to use in an iex console, and easier to read in our liveviews. Again, these go in our `Tunez.Music` domain module, with the resource definition.

```
02/lib/tunez/music.ex
resources do
  # ...
  resource Tunez.Music.Album do
    define :create_album, action: :create
    define :get_album_by_id, action: :read, get_by: :id
    define :update_album, action: :update
    define :destroy_album, action: :destroy
  end
end
```

Similar to artists, we've provided some sample album content for you to play around with — you can import it into your database by running the following on the command line:

```
$ mix run priv/repo/seeds/02-albums.exs
```

This will populate a handful of albums for each of the sample artists we seeded in [code on page 15](#). And now we can start connecting the pieces, to view and manage the album data in our liveviews.

Creating and updating albums

Our Artist page has a button on it to add a new album for that artist. This links to the `TunezWeb.Albums.FormLive` liveview module, and renders a form template similar to the artist form, with text fields for entering data. We can use AshPhoenix to make this template functional, the same way we did for artists.

First, we construct a new form for the `Album.create` action, in `mount/3`:

```
02/lib/tunez_web/live/albums/form_live.ex
def mount(_params, _session, socket) do
  form = Tunez.Music.form_to_create_album()

  socket =
    socket
    |> assign(:form, to_form(form))
    ...
  
```

We validate the form data and update the form in the liveview's state, in the “validate” handle_event/3 event handler:

```
02/lib/tunez_web/live/albums/form_live.ex
def handle_event("validate", %{"form" => form_data}, socket) do
  socket =
    update(socket, :form, fn form ->
      AshPhoenix.Form.validate(form, form_data)
    end)

  {:noreply, socket}
end
  
```

We submit the form in the “save” handle_event/3 event handler, and process the return value:

```
02/lib/tunez_web/live/albums/form_live.ex
def handle_event("save", %{"form" => form_data}, socket) do
  case AshPhoenix.Form.submit(socket.assigns.form, params: form_data) do
    {:ok, _album} ->
      socket =
        socket
        |> put_flash(:info, "Album saved successfully")
        |> push_navigate(to: ~p"/")

      {:noreply, socket}

    {:error, form} ->
      socket =
        socket
        |> put_flash(:error, "Could not save album data")
        |> assign(:form, form)

      {:noreply, socket}
  end
end
  
```

And finally, we add another function head for mount/3, so we can differentiate between viewing the form to *add* an album, and viewing the form to *edit* an album, based on whether or not album_id is present in the params:

```
02/lib/tunez_web/live/albums/form_live.ex
def mount(%{"id" => album_id}, _session, socket) do
  album = Tunez.Music.get_album_by_id!(album_id)
  
```

```

form = Tunez.Music.form_to_update_album(album)

socket =
  socket
    |> assign(:form, to_form(form))
    |> assign(:page_title, "Update Album")

{:ok, socket}
end

def mount(params, _session, socket) do
  form = Tunez.Music.form_to_create_album()
  ...

```

If this was a bit *too* fast, there is a much more thorough rundown on how this code works in [Creating artists with AshPhoenix.Form, on page 25](#).

Using artist data on the album form

There's one thing missing from this form, that will stop it from working as we expect to manage Album records — there's no mention at all of the Artist that the album should belong to. There's a field to enter an artist on the form, but it's disabled.

We do *know* which artist the album should belong to, though. We clicked the button to add an album on a specific artist page, the album should be for that artist! In the server logs in your terminal, you'll see that we do have the artist ID as part of the params to the FormLive liveview:

```
[debug] MOUNT TunezWeb.Albums.FormLive
Parameters: %{"artist_id" => "an-artist-id"}
```

We can use this ID to load the artist record, show the artist details on the form, and relate the artist to the album in the form.

In the second mount/3 function head, for the create action, we can load the artist record using Music.get_artist_by_id, like we do on the artist profile page. The artist can be assigned to the socket alongside the form.

```
02/lib/tunez_web/live/albums/form_live.ex
def mount(%{"artist_id" => artist_id}, _session, socket) do
  artist = Tunez.Music.get_artist_by_id!(artist_id)
  form = Tunez.Music.form_to_create_album()

  socket =
    socket
      |> assign(:form, to_form(form))
      |> assign(:artist, artist)
  ...
```

In the first mount/3 function head, for the update action, we have the artist ID stored on the album record we load. We can use it to load the Artist record in a similar way:

```
02/lib/tunez_web/live/albums/form_live.ex
def mount(%{"id" => album_id}, _session, socket) do
  album = Tunez.Music.get_album_by_id!(album_id)
  artist = Tunez.Music.get_artist_by_id!(album.artist_id)
  form = Tunez.Music.form_to_update_album(album)

  socket =
    socket
    |> assign(:form, to_form(form))
  > |> assign(:artist, artist)
  ...
  ...
```

Now that we have an artist record assigned in the liveview, we can show the artist name in the disabled field, in render/3:

```
02/lib/tunez_web/live/albums/form_live.ex
<.input name="artist_id" value={@artist.name} label="Artist" disabled />
```



New Album

Artist

Valkyrie's Fury

This doesn't actually add the artist info to the form params, though, so we'll still get an error when submitting the form for a new album, even if all of the data is valid. There are two ways we can address this: we *could* manually update the form data before submitting the form, adding the `artist_id` from the artist record already in the socket.

```
def handle_event("save", %{form: form_data}, socket) do
  form_data = Map.put(form_data, "artist_id", socket.assigns.artist.id)
  ...
  ...
```

This is easy to reason about, but feels messy. This code also runs when submitting the form for both creating *and* updating an album, and the update action on our `Album` resource specifically does *not* accept an `artist_id` attribute. Submitting the form won't raise an error — AshPhoenix throws away any

data that won't be accepted by the underlying action — but it's a sign that we're probably doing things wrong.

Instead, we can add a hook when we build the form for creating an album, that will update the changeset pre-validation. It's a similar idea, but because it will only be configured on the create form, not the update form, it makes more sense. AshPhoenix's form builder functions support a number of options³ to configure how the form behaves, and one that we can use here is `transform_params`. It accepts a function that will be called with the form and the params just before the action is run, and there we can make sure the `artist_id` is set to the value we want.

In the `mount/3` function for the `create` action, add the `transform_params` option to set the `artist_id` attribute:

```
02/lib/tunez_web/live/albums/form_live.ex
def mount(%{"artist_id" => artist_id}, _session, socket) do
  artist = Tunez.Music.get_artist_by_id!(artist_id)

  form =
    Tunez.Music.form_to_create_album(
    >   transform_params: fn _form, params, _context ->
    >     Map.put(params, "artist_id", artist.id)
    >   end
    )
  ...
  ...
```

This is a common pattern to use when you want to provide data to an action via a form, but it shouldn't be editable by the user. If we were creating product variants for the clothing store, for example a type of T-shirt that comes in size M in black, we wouldn't want the user to be able to set the parent product ID manually. Even if they're being sneaky in their browser and adding form fields and populating them with data they shouldn't be editing, we overwrite them with the correct values, so nothing nefarious can happen. And now our `TunezWeb.Album.FormLive` form should work properly, for both creating and editing album data.

Loading Related Resource Data

On the profile page for an Artist in `TunezWeb.Artists.ShowLive`, we want to show a list of albums released by that artist. It's currently populated with placeholder data:

3. https://hexdocs.pm/ash_phoenix/AshPhoenix.Form.html#for_create/3



Valkyrie's Fury

[Delete Artist](#)
[Edit Artist](#)

A short biography about the band Valkyrie's Fury

[New Album](#)

Test Album (2023)

[Delete](#)
[Edit](#)


Track data coming soon....



And this data is defined in the handle_params/3 callback:

```
02/lib/tunez_web/live/artists/show_live.ex
def handle_params(%{"id" => artist_id}, _url, socket) do
  artist = Tunez.Music.get_artist_by_id!(artist_id)

  albums = [
    %{
      id: "test-album-1",
      name: "Test Album",
      year_released: 2024,
      cover_image_url: nil
    }
  ]

  socket =
    socket
    |> assign(:artist, artist)
    |> assign(:albums, albums)
    |> assign(:page_title, artist.name)

    {:noreply, socket}
  end
```

Because we've defined albums as a relationship in our Artist resource, we can automatically *load* the data in that relationship, similar to an Ecto preload. All

actions support an extra argument of *options*, and one of the options for read actions⁴ is `load` — a list of relationships we want to load alongside the requested data. This will use the read action we defined on the Album resource, but include the correct filter to only load albums for the artist specified.



By default, loading relationship data with `load` will use the primary read action of the target resource. As we added a primary read action to the Artist resource, first by manually setting `primary?` true inside the action and then changing it to a default action (which is primary by default), this works out of the box.

If we wanted to use a different action for reading data, this can be configured with the `read_action` option⁵ on the relationship.

We can update our call to `get_artist_by_id!` to include loading the `albums` relationship, and remove the hardcoded `albums`:

```
02/lib/tunez_web/live/artists/show_live.ex
def handle_params(%{"id" => artist_id}, _url, socket) do
  artist = Tunez.Music.get_artist_by_id!(artist_id, load: [:albums])
  socket =
    socket
    |> assign(:artist, artist)
    |> assign(:page_title, artist.name)
    {:noreply, socket}
end
```

We do also need to update a little bit of the template, as it referred to the `@albums` assign (which is now deleted). In the `render/1` function, we iterate over the `@albums` and render album details for each:

```
02/lib/tunez_web/live/artists/show_live.ex
<li :for={album <- @albums}>
  <.album_details album={album} />
</li>
```

This can be updated to render albums from the `@artist` instead:

```
02/lib/tunez_web/live/artists/show_live.ex
<li :for={album <- @artist.albums}>
  <.album_details album={album} />
</li>
```

Now when we view the profile page for one of our sample artists, we should be able to see their actual albums, complete with album covers. Neat!

4. <https://hexdocs.pm/ash/Ash.html#read/2>

5. https://hexdocs.pm/ash/dsl-ash-resource.html#relationships-has_many-read_action

We can use `load` to simplify how we loaded the artist for the album on the Album edit form, as well. Instead of making a second request to load the artist after loading the album, we can combine them into one call:

```
02/lib/tunez_web/live/albums/form_live.ex
def mount(%{"id" => album_id}, _session, socket) do
  album = Tunez.Music.get_album_by_id!(album_id, load: [:artist])
  form = Tunez.Music.form_to_update_album(album)

  socket =
    socket
    |> assign(:form, to_form(form))
  >   |> assign(:artist, album.artist)
  ...
  ...
```

The album data on the artist profile looks a little bit funny though — the albums aren't in any kind of order on the page. We should probably show them in chronological order, with the most recent album release listed first. We can look at doing that with a *preparation*.

Setting default ordering with a preparation

Preparations, as the name might imply, *prepare* the stage for a read action to be run on a resource, and are the primary way we can customize the behaviour of a read action. At its most basic, a preparation will take the query that would be run, and return an updated version of that query. They're a perfect place to set sensible defaults for things like ordering of results.



Preparations can also be used to apply default filters for queries.

It can be really tempting to use one for things like filtering out archived record, but we'd caution against doing that, at least globally for a resource. While a default sort can be undone by applying another sort, you can't undo a filter that removes records from a query.

There are a couple of built-in preparations⁶ we can use, or if we needed something more complex, we could write a custom preparation module. Our needs here are pretty straightforward though — the examples in the documentation for `build`⁷ even cover our exact use case!

So to add our default sorting for albums, we can add a new top-level preparations block in our `Album` resource, with a preparation for applying a default sort using `build`:

6. <https://hexdocs.pm/ash/Ash.Resource.Preparation.Builtins.html>

7. <https://hexdocs.pm/ash/Ash.Resource.Preparation.Builtins.html#build/1>

```
02/lib/tunez/music/album.ex
preparations do
  prepare build(sort: [year_released: :desc])
end
```

Now if we reload an artist's profile, we should see the albums being displayed in chronological order, most recent first. That's much more informative!

Structured Data with Validations and Identities

Tunez can now accept form data that should be more structured, instead of just text. We're also looking at data in a smaller scope. Instead of "any artist in the world that ever was", which is a massive data set; we're looking at albums for any individual artist, which is a much smaller and well-defined list.

Let's set some stricter rules for this data, for better data integrity.

Consistent data with validations

With Albums, we need users to enter a valid year for an Album's `year_released` attribute, instead of any old integer; and a valid-looking image URL for the `cover_image_url` attribute. We can enforce these rules with *validations*.

Any defined validations are checked when calling an action, *before* the core functionality (eg. saving or deleting) is run; and if any of the validations fail, the action will abort and return an error. We've seen implicit cases of this already, when we declared that some attributes were `allow_nil? false` — as well as setting the database field for the attribute to be non-null, Ash also validates that the value is present before it even *gets* to the database.

We can add our own validations to our resources, either for an individual action or globally for the entire resource, much like preparations. In our case, we want to ensure that the data is valid at all times, so we can add global validations by adding a new top-level validations block in the Album resource:

```
02/lib/tunez/music/album.ex
defmodule Tunez.Music.Artist do
  # ...

  validations do
    # Validations will go in here
  end
end
```

We'll add two validations to this block, one for `year_released`, and one for `cover_image_url`. Ash provides a lot of built-in validations,⁸ and two of them are relevant here: `numericality` and `match`.

For `year_released`, we want to validate that the user enters a number between, say, 1950 (arbitrarily-chosen number!) and next year (to allow for albums that have been announced but not released), but we should only validate the field if the user has actually entered data. This could be written like so:

```
02/lib/tunez/music/album.ex
validations do
  validate numericality(:year_released,
    greater_than: 1950,
    less_than_or_equal_to: &__MODULE__.next_year/0
  ),
  where: [present(:year_released)],
  message: "must be between 1950 and next year"
end
```

Ash will accept any zero-arity (no-argument) function reference here. The `next_year` function doesn't exist, so we can add it to the very end of the the `Album` module:

```
02/lib/tunez/music/album.ex
def next_year, do: Date.utc_today().year + 1
```

Note that this validation isn't 100% foolproof — if our favorite band from New Zealand is releasing an album on January 1 in the year after next, we might not be able to add that album to Tunez just yet. That's a bit of an edge case outside the scope of this book though!

For `cover_image_url`, we'll add a regular expression to make sure user enters what *looks* like an image URL. This isn't comprehensive by any means — in a real-world app, we'd likely be implementing a fully-featured file uploader, verifying that the uploaded files were valid images, etc. — but for our use case, it'll solve copy-paste mistakes, or entering nonsense.

```
02/lib/tunez/music/album.ex
validations do
  # ...
  validate match(
    :cover_image_url,
    ~r/(^https:\/\|\/images\|/.+\.png|\|.jpg)$/
  ),
  message: "must start with https:// or /images/"
end
```

8. <https://hexdocs.pm/ash/Ash.Resource.Validation.Builtins.html>

We don't need to do anything to integrate these validations into Album actions, or into the forms in our views. Because they're global validations, they apply for every action, and because the forms in our liveviews are built for actions, they will automatically be included. Entering invalid data in the album form will now show validation errors to our users, letting them know what to fix:

Name	Year Released
New Album	2030
	<small>! must be between 1950 and this year</small>
Cover Image URL	
not a real URL	<small>! must start with https:// or /images/</small>

Unique data with identities

There's one last feature we can add, for a better user experience on this form. Some artists have a *lot* of albums, and it would be nice if we could ensure that duplicate albums don't accidentally get entered. Maintaining data integrity, especially with user-editable data, is important — you wouldn't see Wikipedia allowing multiple pages with the same name, for example, they have to be disambiguated in some way. We want Tunez to be an accurate canonical list of artist and album data.

Tunez will consider an album to be a duplicate if it has the same name as another album by the same artist, ie. the combination of `name` and `artist_id` should be unique for every album in the database. (We'll assume that separate versions of albums with the same name get suffixes attached, like "Remastered" or "Live" or "Taylor's Version".) For ensuring this uniqueness, we can use an *identity* on our resource.

As we briefly touched on earlier in [Reading a single record by identity, on page 17](#), Ash defines an *identity*⁹ as any attribute, or combination of attributes, that can uniquely identify a record. A primary key is a natural and automatically-generated identity, but our data may lend itself to other identities as well.

9. <https://hexdocs.pm/ash/identities.html>

To add the new identity to our resource, add a new top-level identities block to the Album resource. An identity has a name, and a list of attributes that make up that identity. We can also specify a message to display on identity violations:

```
02/lib/tunez/music/album.ex
identities do
  identity :unique_album_names_per_artist, [:name, :artist_id],
    message: "already exists for this artist"
end
```

The way identities are handled depends on the data layer being used. Because we're using AshPostgres, the identity will be handled at the database level as a *unique index* on the two database fields, albums.name and albums.artist_id.

To create the index in the database, we can generate migrations after adding the identity to the Album resource:

```
$ mix ashcodegen add_unique_album_names_per_artist
```

This is the first time we've modified a resource and then generated migrations, so it's worth taking a bit of a closer look.

Like the previous times we've generated migrations, AshPostgres has generated a snapshot file representing the current state of the Album resource. It also created a new migration, which has all of the differences between the last snapshot from when we created the resource, and the brand-new snapshot:

```
02/priv/repo/migrations/[timestamp]_add_unique_album_names_per_artist.exs
def up do
  create unique_index(:albums, [:name, :artist_id],
    name: "albums_unique_album_names_per_artist_index"
  )
end

def down do
  drop_if_exists unique_index(:albums, [:name, :artist_id],
    name: "albums_unique_album_names_per_artist_index"
  )
end
```

Ash correctly worked out that the only difference that required database changes was the new identity, so it created the correct migration to add and remove the unique index we need. Awesome!

Run the migration generated:

```
$ mix ash.migrate
```

And now we can test out the changes on the album form. Create an album with a specific name, and then try to create another one for the same artist with the same name — you should get a validation error on the name field, with the message we specified for the identity.

Deleting All of the Things

We'll round out the CRUD interface for Albums with the `destroy` action. We might not need to invoke it too much while using Tunez, but keeping our data clean and accurate is always an important priority.

While building the Album resource, we've also accidentally introduced a bug around Artist deletion, so we should address that as well.

Deleting album data

Deleting albums is done from the artist's profile page, `TunezWeb.Artists.ShowLive`, via a button next to the name of the album.

Clicking the icon will send the “`destroy_album`” event to the liveview. In the event handler, we can fetch the album record from the list of albums we already have in memory, and then delete it. It's a little bit verbose, but it saves another round trip to the database to look up the album record. Like with artists, we also need to handle both the success and error cases:

```
02/lib/tunez_web/live/artists/show_live.ex
def handle_event("destroy_album", %{"id" => album_id}, socket) do
  case Tunez.Music.destroy_album(album_id) do
    :ok ->
      socket =
        socket
        |> update(:artist, fn artist ->
          Map.update!(artist, :albums, fn albums ->
            Enum.reject(albums, &(&1.id == album_id))
          end)
        end
        |> put_flash(:info, "Album deleted successfully")
        {:noreply, socket}
    {:error, error} ->
      Logger.info("Could not delete album '#{album_id}': #{inspect(error)}")
      socket =
        socket
        |> put_flash(:error, "Could not delete album")
        {:noreply, socket}
  end
end
```

We've almost finished the initial implementation for albums! There's a bug in our Album implementation though — if you try to delete an artist that has albums, you'll see what we mean. We'll fix that now!

Cascading deletes with AshPostgres

When we defined our Album resource, we added a `belongs_to` relationship to relate it to Artists:

```
02/lib/tunez/music/album.ex
relationships do
  belongs_to :artist, Tunez.Music.Artist do
    allow_nil? false
  end
end
```

When we generated the migration for this resource in [Defining Relationships, on page 36](#), it created a foreign key in the database, linking the `artist_id` field on the `albums` table to the `id` field on the `artists` table:

```
02/priv/repo/migrations/[timestamp]_create_albums.exs
def up do
  create table(:albums, primary_key: false) do
    # ...
    add :artist_id,
    references(:artists,
      column: :id,
      name: "albums_artist_id_fkey",
      type: :uuid,
      prefix: "public"
    )
  end
end
```

What we *didn't* define, however, was what should happen with this foreign key value when artists are deleted — if there are three albums with `artist_id = "abc123"`, and artist abc123 is deleted, what happens to those albums?

The default behaviour, as we've seen, is to prevent the deletion from happening. This is verified by looking at the server logs when you try to delete one of the artists that this affects:

```
[info] Could not delete artist '«uuid»': %Ash.Error.Invalid{bread_crumbs:
["Error returned from: Tunez.Music.Artist.destroy"], changeset: "#Changese
t<>", errors: [%Ash.Error.Changes.InvalidAttribute{field: :id, message: "would leave
records behind", private_vars: [constraint: :foreign, constraint_name:
"albums_artist_id_fkey"]}, ...]
```

Because an album doesn't make sense without an artist (we can say the albums are *dependent* on the artist), we should delete all of an artist's albums when we delete an artist. There are two ways we can go about this, each with its own pros and cons:

- We can delete the dependent records in code — in the destroy action for an artist, we can call the destroy action on all of the artist's albums as well. It's very explicit what's going on, but it can be *really* slow (relatively speaking). Sometimes it's a necessary evil though, if you need to run business logic in each of the dependent destroy actions.
- Or we can delete the dependent records in the database, by specifying the ON DELETE behaviour¹⁰ of the foreign key that raised the error. This is super-fast, but it can be a little unexpected if you don't know it's happening. You don't get the chance to run any business logic in your app's code — but if you don't need to, this is easily the preferred option.

Which one you use really depends on the requirements of the app you're building, and as the requirements of your app change, you might need to change the behaviour. For now we'll go with the quick ON DELETE option, which is to delete the dependent records in the database (option 2).

AshPostgres lets us specify the ON DELETE behaviour for a foreign key by configuring a custom reference in the postgres block¹¹ of our resource:

```
02/lib/tunez/music/album.ex
postgres do
  table "albums"
  repo Tunez.Repo
  references do
    reference :artist, on_delete: :delete
  end
end
```

This will make a structural change to our database, so we need to generate migrations and run them:

```
$ mix ashcodegen add_reference_for_album_artist_id
$ mix ash.migrate
```

This will generate a migration that modifies the existing foreign key, setting `on_delete: :delete_all`. Running the migration sets the ON DELETE clause on the `artist_id` field:

10. <https://www.postgresql.org/docs/16/sql-createtable.html#SQL-CREATETABLE-PARMS-REFERENCES>

11. https://hexdocs.pm/ash_postgres/dsl-ashpostgres-datalayer.html#postgres-references

```
tunez_dev=# \d albums
«definition of the columns and indexes of the table»
Foreign-key constraints:
  "albums_artist_id_fkey" FOREIGN KEY (artist_id) REFERENCES
    artists(id) ON DELETE CASCADE
```

And now we can delete artists again, even if they have albums; no error occurs and no data is left behind.

Our albums are really shaping up! They're not complete — we'll look at adding track listings in *the (as yet) unwritten content* — but for now they're pretty good, so we can step back and revisit our artist form.

What if we needed to make changes to the data we call an action with, before saving it into the data layer? The UI in our form might not *exactly* match the attributes we want to store, or we might need to format the data, or conditionally set attributes based on other attributes. We can look at making these kinds of modifications with *changes*.

Changing Data Within Actions

We've been using some built-in changes already in Tunez, without even realizing it, for `inserted_at` and `updated_at` timestamps on our resources. We didn't write any code for them, but Ash takes care of setting them to the current time. Both timestamps are set when calling any `create` action, and `updated_at` is set when calling any `update` action.

Like preparations and validations, changes can be defined at both the top-level of a resource, or at an individual action level. The implementation for timestamps *could* look like this:

```
changes do
  change set_attribute(:inserted_at, &DateTime.utc_now/0), on: [:create]
  change set_attribute(:updated_at, &DateTime.utc_now/0)
end
```



By default, global changes will run on any `create` or `update` action, which is why we wouldn't have to specify an action type for `:updated_at` above. They *can* be run on `destroy` actions, but only when opting-in by specifying `on: [:destroy]` on the change.

There's quite a few built-in changes¹² you can use in your resources, or you can add your own, either inline or with a custom module. We'll go through

12. <https://hexdocs.pm/ash/Ash.Resource.Change.Builtins.html>

what it looks like to build one inline, and then how it can be extracted to a module for re-use.

Defining an inline change

Over time, artists go through phases, and sometimes change their names after re-branding, lawsuits, or lineup changes. Let's track updates to an artist's name over time, by keeping a list of all of the previous values that the name field has had, with a new change function.

This list will be stored in a new attribute, called `previous_names`, so we can add it as an attribute in the `Artist` resource. It'll be a list, or *array*, of the previous names, and default to an empty list for new artists:

```
02/lib/tunez/music/artist.ex
attributes do
  # ...
  attribute :previous_names, {:array, :string} do
    default []
  end
  # ...
end
```

Generate a migration to add the new attribute to the database, and run it:

```
$ mix ashcodegen add_previous_names_to_artists
$ mix ash.migrate
```

We only need to run this change when the `Artist` form is submitted to update an `Artist`, so we can add the change within the `:update` action. (If your `Artist` resource is using defaults to define its actions, you'll need to remove `:update` from that list and define the action separately.) The change macro can take a few different forms of arguments, the simplest being a two-argument anonymous function that takes and returns an `Ash.Changeset`:

```
02/lib/tunez/music/artist.ex
actions do
  # ...
  update :update do
    accept [:name, :biography]
    change fn changeset, _context ->
      changeset
  end
end
```

We can make any changes to the `changeset` we want, including deleting data, changing relationships, adding errors, and more. In this way, changes *could*

be thought of as more general versions of validations — if we set an error message in the changeset, it will stop the action from taking place and return the error to the user.

In our anonymous function, we can use some of the functions from `Ash.Changeset`¹³ to read both the old and new name values from the changeset, and update the `previous_names` attribute where applicable:

```
02/lib/tunez/music/artist.ex
change fn changeset, _context ->
  new_name = Ash.Changeset.get_attribute(changeset, :name)
  previous_name = Ash.Changeset.get_data(changeset, :name)
  previous_names = Ash.Changeset.get_data(changeset, :previous_names)

  names =
    [previous_name | previous_names]
    |> Enum.uniq()
    |> Enum.reject(fn name -> name == new_name end)

  Ash.Changeset.change_attribute(changeset, :previous_names, names)
end
```

Like actions, the `change` macro also accepts an optional second argument of options for the change. Because we only need to update `previous_names` if the `name` field is actually being modified, we can add a `changing/1`¹⁴ validation for the `change` function with a `where` check:

```
02/lib/tunez/music/artist.ex
change fn changeset, _context ->
  # ...
end,
where: [changing(:name)]
```

If the validation fails, the `change` function is skipped, and the previous names won't be updated. That'll save a few CPU cycles!

There's one other small change we need to make, for this `change` function to work. By default, `Ash` will try to do as much work as possible in the data layer instead of in memory, via a concept called *atomics*. Because we've written our `change` functionality as imperative code, instead of in a data-layer-compatible way, we'll need to disable atomics for this update action with the `require_atomic?`¹⁵ option.

```
02/lib/tunez/music/artist.ex
update :update do
```

13. <https://hexdocs.pm/ash/Ash.Changeset.html>

14. <https://hexdocs.pm/ash/Ash.Resource.Validation.Builtins.html#changing/1>

15. https://hexdocs.pm/ash/dsl-ash-resource.html#actions-update-require_atomic?

```
require_atomic? false
#
# ...
end
```

We'll dig into atomics, and how to write changes atomically, later in *the (as yet) unwritten content*.

Defining a change module

The inline version of the previous_names change works, but it's a bit long and imperative, smack-bang in the middle of our declarative resource. Imagine if we had a complex resource with a lot of attributes and changes, it'd be really hard to navigate and handle! And what if we wanted to apply this same record-previous-values logic to something else, like users who can change their usernames? We can extract the logic out into a *change module*.

A change module is a standalone module that uses Ash.Resource.Change.¹⁶ Its main access point is the change/3 function, which has a similar function signature as the anonymous change function we defined earlier, but with an added second opts argument. We can move the content of the anonymous change function, and insert it directly into a new change/3 function in a new change module:

```
02/lib/tunez/music/changes/update_previous_names.ex
defmodule Tunez.Music.Changes.UpdatePreviousNames do
  use Ash.Resource.Change

  def change(changeset, _opts, _context) do
    # The code previously in the body of the anonymous change function
  end
end
```

And update the change call in the update action to point to the new module instead:

```
02/lib/tunez/music/artist.ex
update :update do
  require_atomic? false
  accept [:name, :biography]
  >   change Tunez.Music.Changes.UpdatePreviousNames, where: [changing(:name)]
end
```

A shorter and easier-to-read resource isn't the only benefit of extracting changes into their own modules. Change modules can define their own options and interface, and validate their usage at compile time. To reuse the current

16. <https://hexdocs.pm/ash/Ash.Resource.Change.html>

UpdatePreviousNames module, we might want to make the field names configurable instead of hardcoded to name and previous_names, and have a flag for allowing duplicate values or not.

Details on configuring and validating the interface for change modules using the Spark¹⁷ library are a bit too much to go into here, but built-in changes like Ash.Resource.Change.SetAttribute¹⁸ are a great way to see how they can be implemented.

Changes run more often than you might think!

It's really important to note that changes aren't *only* run when actions are called. When forms are tied to actions, like our update action is tied to the Artist edit form in the web interface, the pre-persistence steps like validations and changes are run multiple times:

- When building the initial form
- During any authorization checks (covered in [Introducing Policies, on page 133](#))
- On every validation of the form
- And when actually submitting the form or calling the action.

Because of this, changes that are time-consuming, or have side effects such as calling external APIs, should be wrapped in *hooks* such as Ash.Changeset.before_action or Ash.Changeset.after_action — these will only be called immediately before or after the action is run.

If we wanted to do this for the UpdatePreviousNames change module, it would look like this:

```
def change(changeset, _opts, _context) do
  Ash.Changeset.before_action(changeset, fn changeset ->
    # The code previously in the body of the function
    # It can still use any `opts` or `context` passed in to the top-level
    # change function, as well.
  end)
end
```

The anonymous function set as the before_action would only run once — when the form is submitted — but it would still have the power to set errors on the changeset to prevent the changes from being saved, if necessary.

17. <https://hexdocs.pm/spark/>

18. https://github.com/ash-project/ash/blob/main/lib/ash/resource/change/set_attribute.ex

Setting attributes in a ‘before_action’ hook will bypass validations!

It's super-important to note that a function defined as a `before_action` will only run *right before save — after* validations of the action have been run.



As such, it's possible to get your data into an invalid state in the database. If you validate that an album's `year_released` must be in the past, but then call `Ash.Changeset.change_attribute(changeset, :year_released, 2050)` in your `before_action` function, that year 2050 will happily be saved into the database. Ash will show a warning at runtime if you do this, which is helpful.

If you want to force any validation to run *after* `before_action` hooks, you can use the `before_action?`¹⁹ option on the validation, eg.

```
validate numericality(:year_released, ...), before_action?: true
Or if you simply want to silence the warning because you're fine
with skipping the validation, replace your call to change_attribute
with force_change_attribute instead.
```

Rendering the previous names in the UI

To finish this feature off, we can show any previous names that an artist has had, on their profile page.

In `TunezWeb.Artists.ShowLive`, we can add the names printed out, as part of the `<header>` block in the `render/1` function:

```
02/lib/tunez_web/live/artists/show_live.ex
<.header>
  <.h1>...</.h1>
>  <:subtitle :if={@artist.previous_names != []}>
>    formerly known as: #{@artist.previous_names.map { |name| name.to_s }.join(", ")}
>  </:subtitle>
...
</ .header>
```



Refrain

formerly known as: Refrain of Fire

[Delete Artist](#)

[Edit Artist](#)

19. https://hexdocs.pm/ash/dsl-ash-resource.html#validations-validate-before_action?

And now our real Artist pages, complete with their real Album listings, are complete! We've learnt about the tools Ash provides for relating resources together, and how we can work with related data for efficient data loading, preparations, and data integrity. These are core building blocks, that you can use when building out your own applications and we'll be using more of in the future as well.

And we *still* haven't needed to write a lot of code — the small snippets we've written, like validations and changes, have been very targeted and specific, but have been usable throughout the whole app, from seeding data in the database to rendering errors in the UI.

We're only scratching the surface though — in the next chapter, we'll make the Artist catalog really useful, giving users the ability to search, sort, and page through artists, using more of Ash's built-in functionality. We'll also see how we can use calculations and aggregates to perform some sophisticated queries, without even breaking a sweat. This is where things will *really* get interesting!

Creating a Better Search UI

In the previous chapter, we learned how we can link resources together with relationships, and use validations, preparations and changes to implement business logic within resource actions. With this new knowledge, we could build out a fairly comprehensive data model if we wished. We could make resources for everything from tracks on albums, band members, record labels, anything we wanted, and also make a reasonable CRUD interface for it all. We've covered a lot!

But we're definitely missing some business logic, and UI polish. If we had a whole lot of artists in the app, it would become difficult to use. The artist catalog is just one big list of cards — there's no way to search or sort data, and we definitely don't need the *whole* list at all times. Let's look at making this catalog a lot more user-friendly, using query filtering, sorting, and pagination.

Custom Actions with Arguments

To improve discoverability, we can add search to the Artist catalog, to allow users to look up artists by name.

What might it ideally look like, if we were designing the interface for this function? It'd be great to be able to call it like:

```
iex> Tunez.Music.search_artists("fur")
{:ok, [%Tunez.Music.Artist{name: "Valkyrie's Fury"}, ...]}
```

Can we do it? Yes we can!

Designing a search action

A search action will be reading existing data from the database, so we can add a new read action to the Artist resource to perform this new search.

`03/lib/tunez/music/artist.ex`

```
actions do
  # ...
  >  read :search do
  >    end
end
```

When we covered read actions in [Defining a read action, on page 16](#), we mentioned that Ash will read data from the data layer based on parameters we provide, which can be done as part of the action definition. Our search action will support one such parameter for now, the text to match names on, via an *argument* to the action.

`03/lib/tunez/music/artist.ex`

```
actions do
  # ...
  read :search do
    argument :query, :ci_string do
      constraints allow_empty?: true
      default ""
    end
  end
end
```

Arguments can be anything from scalar values like integers or booleans, to structured maps, to resource structs. In our case, we'll be accepting a case-insensitive string (or *ci_string*), to allow for case-insensitive searching. This argument can then be used in a *filter*, to add conditions to the query, limiting the records returned to only those that match the condition.

`03/lib/tunez/music/artist.ex`

```
actions do
  # ...
  read :search do
    argument :query, :ci_string do
      # ...
    end
  end
  >  filter expr(contains(name, ^arg(:query)))
  end
end
```

Whoa! There's a lot of new stuff in single line of code. Let's break it down a bit.

Filters with expressions

Filters are the where-clauses of our queries, allowing us to only fetch the records that match our query. They use a special SQL-like syntax, inspired by Ecto, but are much more expressive.

In iex you can test out different filters by running some queries inline. You'll need to run require Ash.Query first:

```
iex(1)> require Ash.Query
Ash.Query
iex(2)> Ash.Query.filter(Tunez.Music.Album, year_released == 2024)
#Ash.Query<resource: Tunez.Music.Album,
  filter: #Ash.Filter<year_released == 2024>>
iex(3)> |> Ash.read()
[debug] QUERY OK source="albums" db=0.8ms queue=0.8ms idle=1138.7ms
SELECT a0."id", a0."name", a0."year_released", a0."cover_image_url",
a0."inserted_at", a0."updated_at", a0."artist_id" FROM "albums" AS a0
WHERE (a0."year_released"::bigint = $1::bigint) ORDER BY
a0."year_released" DESC [2024]
{:ok, [%Tunez.Music.Album{year_released: 2024, ...}, ...]}
```

Filters aren't just limited to equality checking — they can use any of the expression syntax¹ including operators and functions. All of the expression syntax listed is data-layer-agnostic, and because we're using AshPostgres, is converted into SQL when we run the query.

Unlike running a filter with Ash.Query.filter, whenever we refer to expressions elsewhere, we need to wrap the body of the filter in a call to expr. The reasons for this are historical — the Ash.Query.filter function predates other usages of expressions, so this will likely be changed in a future version of Ash for consistency — so it's something to keep in mind.

Inside our filter, we can use the contains/2 expression function, which is a substring checker. It checks to see if the first argument, in our case a reference to the name attribute of our resource, contains the second argument, which is a reference to the query argument to the action!

If we were using a similar filter in pure Elixir, to filter a list of structs by a case-insensitive string search, it might look like:

```
def search(query)
  query = String.downcase(query)

  load_all_artists()
  |> Enum.filter(fn %{name: name} ->
    name
```

1. <https://hexdocs.pm/ash/expressions.html>

```

|> String.downcase()
|> String.contains?(query)
end)
end

search("test")

```

But because we're using AshPostgres, it will be converted to use the `ilike` function in PostgreSQL,² and generate a query like:

```

SELECT a0."id", a0."name", a0."biography", a0."previous_names",
a0."inserted_at", a0."updated_at" FROM "artists" AS a0 WHERE
(a0."name"::text ILIKE $1) ["%test%"]

```

Which does exactly what we want — a case-insensitive substring match on the column contents, based on the string we provide.

Speeding things up with custom database indexes

Using an `ilike` query naively over a massive data set isn't exactly performant — it'll run a sequential scan over every record in the table. As more and more artists get added, the search would get slower and slower. To make this query more efficient, we can add a custom database index called a GIN index³ on the name column.

AshPostgres supports the creation of custom indexes⁴ like a GIN index. To create a GIN index specifically, we first need to enable the PostgreSQL `pg_trgm` extension.⁵ AshPostgres handles enabling and disabling PostgreSQL extensions, via the `installed_extensions`⁶ function in the `Tunez.Repo` module. By default it only includes ash-functions, so we can add `pg_trgm` to this list:

```

03/lib/tunez/repo.ex
defmodule Tunez.Repo do
  use AshPostgres.Repo, otp_app: :tunez

  def installed_extensions do
    # Add extensions here, and the migration generator will install them.
    ["ash-functions", "pg_trgm"]
  end

```

Then we can add the index to the `postgres` block of our Artist resource:

```

03/lib/tunez/music/artist.ex
postgres do

```

-
2. <https://www.postgresql.org/docs/current/functions-matching.html#FUNCTIONS-LIKE>
 3. https://pganalyze.com/blog/gin-index#indexing-like-searches-with-trigrams-and-gin_trgm_ops
 4. https://hexdocs.pm/ash_postgres/dsl-ashpostgres-datalayer.html#postgres-custom_indexes
 5. <https://www.postgresql.org/docs/current/pgtrgm.html>
 6. https://hexdocs.pm/ash_postgres/AshPostgres.Repo.html#module-installed-extensions

```
table "artists"
repo Tunez.Repo

custom_indexes do
  index "name gin_trgm_ops", name: "artists_name_gin_index", using: "GIN"
end
end
```

Generally, AshPostgres will be able to generate the names of indexes by itself from the fields, but because we're creating a custom index we have to specify a valid name.

Lastly, generate and run the migration to update the database with the new extension and index.

```
$ mix ash.codegen add_gin_index_for_artist_name_search
$ mix ash.migrate
```

What kind of performance benefits do we actually get for this? We ran some tests, by inserting a million records with various names into our Artists table using the faker library. Without the index, searching for a word like “snow” (which returns 3,041 results in our data set) takes about 150ms.

```
Gather (cost=1000.00..23725.43 rows=10101 width=131) (actual
time=0.633..150.672 rows=3041 loops=1)
  Workers Planned: 2
  Workers Launched: 2
    -> Parallel Seq Scan on artists a0 (cost=0.00..21715.33 rows=4209
        width=131) (actual time=0.570..136.405 rows=1014 loops=3)
        Filter: (name ~~* '%snow%':text)
        Rows Removed by Filter: 332320
Planning Time: 0.351 ms
Execution Time: 150.789 ms
```

With the index, however, the numbers look a *lot* different:

```
Bitmap Heap Scan on artists a0 (cost=118.28..15299.90 rows=10101
width=131) (actual time=1.571..13.443 rows=3041 loops=1)
  Recheck Cond: (name ~~* '%snow%':text)
  Heap Blocks: exact=2759
    -> Bitmap Index Scan on artists_name_idx (cost=0.00..115.76 rows=10101
        width=0) (actual time=1.104..1.105 rows=3041 loops=1)
        Index Cond: (name ~~* '%snow%':text)
Planning Time: 0.397 ms
Execution Time: 14.691 ms
```

That's a huge saving, the query now only takes about 10% of the time! It might not seem like such a big deal when we're talking about milliseconds, but for every artist query being run, it all adds up. Faster responses lead to happier users and less server load. Everybody wins!

Integrating search into the UI

Now that we have the first version of our search action built, we can make the tidy interface we imagined and integrate it into the Artist catalog.

A code interface with arguments

We previously imagined a search function API like:

```
iex> Tunez.Music.search_artists("fur")
{:ok, [%Tunez.Music.Artist{name: "Valkyrie's Fury"}, ...]}
```

We can create a new code interface in our domain for this, one that supports passing arguments to the action.

```
03/lib/tunez/music.ex
resource Tunez.Music.Artist do
  # ...
  define :search_artists, action: :search, args: [:query]
end
```

Defining a list of arguments with names that match the arguments defined in our function makes that link that we're after — the first parameter passed when calling the `Tunez.Music.search_artists` function will now be assigned to the `query` argument in the action.

You can verify that this link has been made by checking out the function signature in iex:

```
iex(1)> h Tunez.Music.search_artists
def search_artists(query, params_or_opts \\ %{}, opts \\ [])
```

Calls the search action on `Tunez.Music.Artist`.

Any action arguments *not* listed in the `args` list on the code interface will be placed into the next argument, the map of `params`. If we didn't specify `args: [:query]`, we would need to call the search function like this:

```
Tunez.Music.search_artists(%{query: "fur"})
```

Which works, but isn't anywhere near as nice!

Searching from the catalog

In the Artist catalog, searches should be repeatable and sharable, and we'll achieve this by making the searched-for text part of the query string, in the page URL. If a user visits a URL like <http://localhost:4000/?q=test>, Tunez should run a search for the string “test” and show only the matching results.

We currently read the list of artists to display in the `handle_params/3` function definition, in `Tunez.Artists.IndexLive`:

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_params(_params, _url, socket) do
  artists = Tunez.Music.read_artists!()

  socket =
    socket
    |> assign(:artists, artists)
  # ...
```

If the params to the page (from the page route/query string) contain a `q` value, we can read it out and call our new `search_artists` function instead:

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_params(params, _url, socket) do
  query_text = Map.get(params, "q", "")
  artists = Tunez.Music.search_artists!(query_text)

  socket =
    socket
    |> assign(:query_text, query_text)
    |> assign(:artists, artists)
  # ...
```

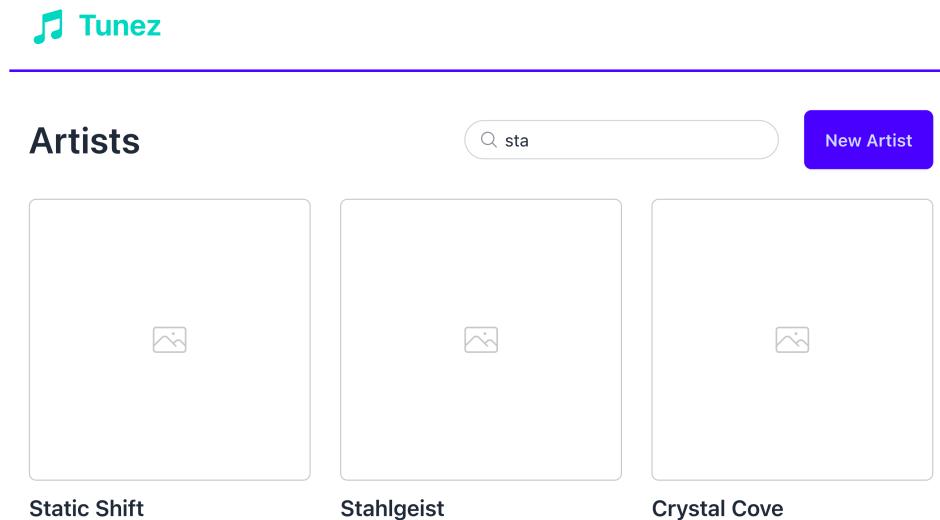
Of course, users don't search by editing the URL — they search by typing text in a search box. We can add another action slot to the `.header` component in the `render/3` function of `IndexLive`, to render a search box function component.

```
03/lib/tunez_web/live/artists/index_live.ex
<.header responsive={false}>
  <.h1>Artists</ .h1>
  >  <:action>
  >    <.search_box query={@query_text} method="get"
  >      data-role="artist-search" phx-submit="search" />
  >  </:action>
  >  <:action>
  >    <.button_link
  # ...
```

When a user types something in the box and presses Enter to submit, the “`search`” event will be sent to the liveview. The logic for this has already been implemented:

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_event("search", %{"query" => query}, socket) do
  params = remove_empty(%{q: query})
  {:noreply, push_patch(socket, to: ~p"/?#{params}")}
```

The event handler takes the entered text and patches the liveview — updating the URL with the new query string and calling `handle_params/3` with the new params, which then re-runs the search and will re-render the catalog.



That's really neat! There's another big thing we can add, to make the artist catalog more awesome — the ability to sort the artists on the page. We'll start with some basic sorts, like sorting them alphabetically or by most recently updated, and then later in the chapter we'll look at some really *amazing* ones.

Dynamically Sorting Artists

Our searching functionality is fairly limited — Tunez doesn't have a concept of "best match" when searching text, artists either match or they don't. To help users potentially surface what they want to see more easily, we'll let them sort their search results. Maybe they want to see most recently added artists listed first? Maybe they want to see artists that have released the most albums listed first? (oops that's a bit of a spoiler!) Let's dig in.

Letting users set a sort method

We'll start from the UI layer — how can users select a sort method? Usually it's by a dropdown of sort options at the top of the page, so we'll drop one next to the search box. In `Tunez.Artists.IndexLive`, we can add another action to the actions list in the header function component:

```
03/lib/tunez_web/live/artists/index_live.ex
<.header responsive={false}>
```

```
>   <.h1>Artists</ .h1>
>   <:action><.sort_changer selected={@sort_by} /></ :action>
<:action>
# ...
```

The @sort_by assign doesn't yet exist, but it will store a string defining what kind of sort we want to perform. We can add this to the list of assigns, in handle_params/3:

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_params(params, _url, socket) do
  >   sort_by = nil
  # ...

  socket =
    socket
  >   |> assign(:sort_by, sort_by)
  >   |> assign(:query_text, query_text)
  # ...
```

The actual sort_changer function component has already been defined further down in the liveview — it reads a set of option tuples for the sort methods we'll support, with internal and display representations, and embeds them into a form, with a phx-change event handler.

```
03/lib/tunez_web/live/artists/index_live.ex
def sort_options do
  [
    {"updated_at", "recently updated"},
    {"inserted_at", "recently added"},
    {"name", "name"}
  ]
end
```

When the user selects a sort option, the “change_sort” event will be sent to the liveview. The handle_event/3 function head for this event looks pretty similar to the function head for the “search” event, right below it, except we now have an extra sort_by parameter in the query string. Let's add sort_by to the params list in “search” event handler as well, by reading it from the socket assigns. This will let users search then sort, or sort then search, and the result will be the same because both parameters will always be part of the URL.

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_event("change_sort", %{sort_by: sort_by}, socket) do
  params = remove_empty(%{q: socket.assigns.query_text, sort_by: sort_by})
  {:noreply, push_patch(socket, to: ~p"/?#{params}")}
```

```
# Context: Reacting to changes in the search query and updating the URL
def handle_event("search", %{"query" => query}, socket) do
  params = remove_empty(%{q: query, sort_by: socket.assigns.sort_by})
  {:noreply, push_patch(socket, to: ~p"/?#{params}")}
```

► end

Test it out in your browser — now changing the sort dropdown should navigate to a URL with the sort method in the query string, like http://localhost:4000/?q=the&sort_by=name.

Now that we have the sort method in the query string, we can read it when the page loads just like we read pagination parameters, in `handle_params/3`. We'll do some validation to make sure that it's a valid option from the list of options, and then store it in the socket like before.

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_params(params, _url, socket) do
  sort_by = Map.get(params, "sort_by") |> validate_sort_by()
  # ...

  socket =
    # ...
```

That's the full loop of what we need to implement from a UI perspective — we have a default sort method defined, the user can change the selected value, that value gets reflected back to them in the URL and on the page. Now we can look at how to *use* that value, to change the way the data is returned when our user runs a search.

The base query for a read action

When we run any read action on a resource, we always have to start from some base, onto which we can build a query and start layering extras like filters, loads and so on. We've seen examples of this throughout the book, all the way back to our very first example of how to run a read action:

```
► iex(2)> Tunez.Music.Artist
► Tunez.Music.Artist
► iex(3)> |> Ash.Query.for_read(:read)
► #Ash.Query<resource: Tunez.Music.Artist>
iex(4)> |> Ash.Query.sort(name: :asc)
#Ash.Query<resource: Tunez.Music.Artist, sort: [name: :asc]>
iex(5)> |> Ash.Query.limit(1)
#Ash.Query<resource: Tunez.Music.Artist, sort: [name: :asc], limit: 1>
iex(6)> |> Ash.read()
[debug] QUERY OK source="artists" db=3.4ms queue=1.2ms idle=1863.7ms
SELECT a0."id", a0."name", a0."biography", a0."inserted_at", a0."updated_at"
FROM "artists" AS a0 ORDER BY a0."name" LIMIT $1 [1]
{:ok, [#Tunez.Music.Artist<...>]}
```

The core thing needed to create a query for a read action is knowing which resource needs to be read. By default this is what Ash does when we call a read action, or a code interface that points to a read action — it uses the resource module itself as the base and builds the query from there.

We *can* change this though. We can pass in our own hand-rolled query when calling a code interface for a read action, or pass a list of options to be used with the resource module when constructing the base query, and these will be used instead.

We mentioned earlier in [Loading Related Resource Data, on page 43](#) that every action can take an optional set of arguments, but it's worth reiterating. These don't have to be defined as arguments to the action, they're added at the end, and they can radically change the behaviour of the action. For code interfaces for read actions, this list of options⁷ includes the query option, and that's what we'll use to provide a query in the form of a keyword list.

The query keyword list can include any of the opts that `Ash.Query.build/3`⁸ supports, and in our case we're interested in setting a sort order, so we'll pick `sort_input`.

Using `sort_input` for succinct yet expressive sorting

A few of you are probably already wondering, `sort_input`, when `sort` is right there? What's the difference? Both could be used for our purposes, but one is much more useful than the other when sorts come from query string input.

`sort` is the traditional way of specifying a sort order with field names and sort directions, eg. to order records alphabetically by name, A to Z, you would specify `[name: :asc]`. To order alphabetically by name and then by newest created record first (to consistently sort artists that have the same name), you would specify `[name: :asc, inserted_at: :desc]`. Which is fine, this works, you can test it out with iex and our `Tunez.Music.search_artists` code interface function:

```
iex(6)> Tunez.Music.search_artists("the", [query: [sort: [name: :asc]]])
{:ok,
%Ash.Page.Offset{
  results: [
    #Tunez.Music.Artist<name: "Nights in the Nullarbor", ...>,
    #Tunez.Music.Artist<name: "The Lost Keys", ...>
  ],
  ...
}}
```

7. <https://hexdocs.pm/ash/code-interfaces.html#using-the-code-interface>

8. <https://hexdocs.pm/ash/Ash.Query.html#build/3>

`sort_input` is a bit different — instead of a list, we can specify a single comma-separated string of fields to sort on. Sorting is ascending by default, but can be inverted by prefixing a field name with a `-`. So our example from before, sorting alphabetically by name and then newest first would be `name,-inserted_at`. Heaps better!

To use `sort_input`, we do need to make one change to our resource though — as it's intended to let users specify their own sort methods, it will only permit sorting on *public* attributes. We don't want users trying to hack our app in any way, after all. All attributes are *private* by default, for the highest level of security, so we'll have to explicitly mark those we want to be publicly accessible. This can be done by adding `public? true` as an option on each of the attributes we want to be sortable:

```
03/lib/tunez/music/artist.ex
attributes do
  # ...
  attribute :name, :string do
    allow_nil? false
  >   public? true
  end
  # ...
  >   create_timestamp :inserted_at, public?: true
  >   update_timestamp :updated_at, public?: true
end
```

Once the attributes are marked public, then `sort_input` will be usable the way we want:

```
iex(6)> Tunez.Music.search_artists("the", [query: [sort_input: "-name"]])
{:ok, %Ash.Page.Offset{results: [...], ...}}
```

Because we've condensed sorting down to specifying a single string value at runtime, it's perfect for adding as an option when we run our search, in the `handle_params/3` function in `TunezWeb.Artists.IndexLive`:

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_params(params, _url, socket) do
  # ...
  >   artists = Tunez.Music.search_artists!(query_text, query: [sort_input: sort_by])
  # ...
```

It's actually more powerful than our UI needs — it supports supporting on multiple columns, when we only have a single dropdown for one field, but that's okay. There's just one little tweak to make in our `sort_options` function —

when we want recently added or updated records to be shown first, they should be sorted *descending*, so prefix those two field names with a `-`.

```
03/lib/tunez_web/live/artists/index_live.ex
def sort_options do
  [
    {"-updated_at", "recently updated"},
    {"-inserted_at", "recently added"},
    {"name", "name"}
  ]
end
```

We can now search and sort, or sort and search, and everything works just as expected. There's still too much data to display on the page, though. Even if searching through All Of The Artists That Ever Was returns only a few hundred or thousand results, that would take too long to render. We'll split up the results with pagination, and let users browse artists at their own pace.

Pagination of Search Results

Pagination is the best way of limiting the amount of data on an initial page load, to the most important things a user would want to see. If they want more data (more search results, more tweets/posts, more artists), they can request more data, either by scrolling to the bottom of the page and having more data load automatically (usually called *infinite scroll*), or by clicking a button to load more.

We'll implement the more traditional method of having distinct pages of results, and letting users go backwards/forwards between pages via buttons at the bottom of the catalog.

Adding pagination support to the search action

Our first step in implementing pagination is to update our search action to support it. Ash supports automatic pagination of read actions⁹ using the pagination¹⁰ macro, so we'll add that to our action definition.

```
03/lib/tunez/music/artist.ex
read :search do
  # ...
  >   pagination offset?: true, default_limit: 12
end
```

9. <https://hexdocs.pm/ash/read-actions.html#pagination>

10. <https://hexdocs.pm/ash/dsl-ash-resource.html#actions-read-pagination>

We've elected to use offset pagination (eg. "show the next 20 records after the 40th record") over keyset pagination (eg. "show the next 20 records after record ID=12345"). It's easier to understand, and well-suited for when the data isn't frequently being updated. When the data is frequently being updated, such as for news feeds or timelines, or you want to implement infinite scrolling, then keyset pagination would be the better choice.

Adding the pagination macro immediately changes the return type of the search action. You can see this if you run a sample search in iex:

```
iex(1)> Tunez.Music.search_artists!("cove")
<<database query>>
%Ash.Page.Offset{
  results: [#Tunez.Music.Artist<name: "Crystal Cove", ...>],
  limit: 12,
  offset: 0,
  count: nil,
  rerun: {#Ash.Query<...>, [<<opts>>]},
  more?: false
}
```

The list of artists resulting from running the text search is now wrapped up in an Ash.Page.Offset struct, which contains extra pagination-related information — how many results there are in total, whether there are more results to display, and information about how this same query can be re-run if necessary. If we were using keyset pagination, you'd get back an Ash.Page.Keyset struct instead, but the data within would be similar.

This means we'll need to update the liveview as well, to support the new data structure.

Showing paginated data in the catalog

In TunezWeb.Artists.IndexLive, we load a list of artists and assign them directly to the socket, for the template to iterate over:

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_params(params, _url, socket) do
  # ...
  artists = Tunez.Music.search_artists!(query_text, query: [sort_input: sort_by])
  socket =
    socket
    |> assign(:query_text, query_text)
    |> assign(:artists, artists)
  # ...
```

Now calling `search_artists!` will return a `Page` struct, so rename the variable and `socket assign` to better reflect what is being stored.

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_params(params, _url, socket) do
  # ...
  page = Tunez.Music.search_artists!(query_text, query: [sort_input: sort_by])
  socket =
    socket
    |> assign(:query_text, query_text)
  ➤   |> assign(:page, page)
  # ...
```

We also need to update the template code, to use the new @page assign and iterate through the page results.

```
03/lib/tunez_web/live/artists/index_live.ex
➤ <div :if={@page.results == []} class="p-8 text-center">
  <.icon name="hero-face-frown" class="w-32 h-32 bg-base-300" />
  <br /> No artist data to display!
</div>
#
# -----
# Context: Updating the template code to use the paginated artist results
<ul class="gap-6 lg:gap-12 grid grid-cols-2 sm:grid-cols-3 lg:grid-cols-4">
  ➤   <li :for={artist <- @page.results}>
    <.artist_card artist={artist} />
  </li>
</ul>
```

This works — now we only have one page worth of artists showing in the catalog, but no way of navigating to other pages.

We can add some dummy pagination links to the bottom of the artist catalog template, with the pagination_links function component. As pagination info will also be kept in the URL for easy sharing/reloading, the component will use the query text, the current sort and the page to construct URLs to link to, for changing pages.

```
03/lib/tunez_web/live/artists/index_live.ex
def render(assigns)
  ~H"""
  # ...
  <.pagination_links page={@page} query_text={@query_text} sort_by={@sort_by} />
  """
end
```

To make the pagination links functional, we'll look at another of AshPhoenix's modules — `AshPhoenix.LiveView`.¹¹ It contains a handful of really useful helper

11. https://hexdocs.pm/ash_phoenix/AshPhoenix.LiveView.html

functions for inspecting a Page struct to see if there's a previous page, next page, what the current page is, and so on. We can use these to add links to for the next/previous pages in the pagination_links function component, conditionally disabling them if there is no valid page to link to.

```
03/lib/tunez_web/live/artists/index_live.ex
<div
  >   :if={AshPhoenix.LiveView.prev_page?(@page) ||
  >         AshPhoenix.LiveView.next_page?(@page)}
  >     class="flex justify-center pt-8 join"
  >
  >       <.button_link data-role="previous-page"
  >         patch={"/?#{query_string(@page, @query_text, @sort_by, "prev")}"}
  >         disabled={!AshPhoenix.LiveView.prev_page?(@page)}
  >         class="join-item" kind="primary" outline
  >
  >           « Previous
  </.button_link>
  <.button_link data-role="next-page"
  >         patch={"/?#{query_string(@page, @query_text, @sort_by, "next")}"}
  >         disabled={!AshPhoenix.LiveView.next_page?(@page)}
  >         class="join-item" kind="primary" outline
  >
  >           Next »
  </.button_link>
</div>
```

The query_string helper function doesn't yet exist, but we can quickly write it. It will take some pagination info from the @page struct, and use it to generate a keyword list of data to put in the query string:

```
03/lib/tunez_web/live/artists/index_live.ex
def query_string(page, query_text, sort_by, which) do
  case AshPhoenix.LiveView.page_link_params(page, which) do
    :invalid -> []
    list -> list
  end
  |> Keyword.put(:q, query_text)
  |> Keyword.put(:sort_by, sort_by)
  |> remove_empty()
end
```

We're using offset pagination, so when you call AshPhoenix.LiveView.page_link_params/2 it will generate limit and offset parameters.

```
iex(1)> page = Tunez.Music.search_artists!("a")
%Ash.Page.Offset{results: [#Tunez.Music.Artist<...>, ...], ...}
iex(2)> TunezWeb.Artists.IndexLive.query_string(page, "a", "name", "prev")
[sort_by: "name", q: "a"]
iex(3)> TunezWeb.Artists.IndexLive.query_string(page, "a", "-inserted_at", "next")
```

```
[sort_by: "-inserted_at", q: "a", limit: 12, offset: 12]
```

When interpolated by Phoenix into a URL on the “Next” button link, it will become http://localhost:4000/?q=a&sort_by=name&limit=12&offset=12.



The last step in the process is to use these limit/offset parameters, to make sure we load the right page of data. At the moment even if we click “Next”, the URL changes but we still only see the first page of artists. To do that, we’ll use another one of the helpers from `AshPhoenix.LiveView`, to parse the right data out of the params before we load artist data in `handle_params/3`.

```
03/lib/tunez_web/live/artists/index_live.ex
def handle_params(params, _url, socket) do
  # ...
  >>> page_params = AshPhoenix.LiveView.page_from_params(params, 12)
  page =
    Tunez.Music.search_artists!(query_text,
  >>>   page: page_params,
  >>>   query: [sort_input: sort_by]
  )
```

We *could* pluck out the limit and offset values from params ourselves, but by doing it this way, if we wanted to change the pagination type — from offset to keyset, or vice versa — we wouldn’t have to touch this view code at all. We’d only have to change one line of code, the pagination definition in the search action, and everything else would still work. If you really want to be wild, you can even support *both* types of pagination in the action — URLs that include params for either type will work. Nifty!

And that’s it! We’ve now got full sorting, searching, and pagination for our artist catalog. It was a lot to go through and understand, but not actually a lot of code. Concerns that belong entirely to our UI, like sorting, stayed in the UI layer of the app. Features that are more in-depth, like text searching, came into the resource layer to be analyzed and optimized.

Now we’d love to talk about a really killer data modelling feature that Ash provides — calculations!

No DB field? No Problems, with Calculations

Calculations¹² are an awesome way of defining a special type of attribute that isn't stored in your database, but is *calculated* on-demand from other information. What types of other information can you use? Well, anything really! You can use data from related resources, data from files on the filesystem, from external sources, or even just some way of tweaking, deriving or reformatting data you already store for that resource.

Calculations have to be specifically loaded when reading data, the same way you load a relationship; but once they're loaded, they can be treated like any other attribute of a resource.

Calculating data with style

Let's say we wanted to display how many years ago each album was released, on an artist's profile page. That'll make all Tunez's users feel really old! (We won't *actually* do this, because it's a terrible idea, but for demonstration purposes.)

Like a lot of the functionality we've seen before, we can add calculations to a resource by defining a top-level calculations block in the resource.

```
defmodule Tunez.Music.Album do
  # ...
  ▶  calculations do
  ▶    end
end
```

Inside the calculations block, we can use the calculate¹³ macro to define individual calculations. A calculation needs three things — a name for the resulting attribute, the type of the resulting attribute, and some method of generating the value to store in the attribute.

```
calculations do
  ▶  calculate :years_ago, :integer, expr(2025 - year_released)
end
```

Calculations use the expression syntax¹⁴ that we saw earlier with filters, to make for really terse code. These are SQL-ish, so we can't use arbitrary Elixir functions in them (hence we're hardcoding 2025 for the year), but we can write some very complex conditions. If we wanted to use some logic that *can't*

12. <https://hexdocs.pm/ash/calculations.html>

13. <https://hexdocs.pm/ash/dsl-ash-resource.html#calculations-calculate>

14. <https://hexdocs.pm/ash/expressions.html>

be written as an expression, such as dynamically using the current year, or converting a string of minutes and seconds to a number of seconds, we could define a separate calculation module. We'll see this exact example later in *the (as yet) unwritten I don't know how to generate a cross reference to ch08.minutes_to_seconds!*

Once you've added a calculation, you can test it out in iex by loading the calculation as part of the data for an album. We've seen `load: [:albums]` when loading artist data before, and to load nested data, each item in the load list can be a keyword list of nested things to load.

```
iex(1)> Tunez.Music.get_artist_by_id(<<uuid>>, load: [:albums: [:years_ago]])
{:ok, #Tunez.Music.Artist<
  albums: [
    #Tunez.Music.Album<year_released: 2022, years_ago: 3, ...>,
    #Tunez.Music.Album<year_released: 2012, years_ago: 13, ...>
  ],
  ...
>}
```

You could then use this `years_ago` attribute when rendering a view, or in an API response, like any other attribute. And because they *are* just like any other attribute, you can even use them within other calculations:

```
calculations do
  calculate :years_ago, :integer, expr(2025 - year_released)
  > calculate :string_years_ago,
  >           :string,
  >           expr("wow, this was released " <> years_ago <> " years ago!")
end
```

If you load the `string_years_ago` calculation, you don't need to specify that it depends on another calculation so that should be loaded too — Ash can work that out for you.

```
iex(1)> Tunez.Music.get_artist_by_id(<<uuid>>,
...> load: [:albums: [:string_years_ago]])
{:ok, #Tunez.Music.Artist<
  albums: [
    #Tunez.Music.Album<
      year_released: 2022,
      string_years_ago: "wow, this was released 3 years ago!",
      years_ago: #Ash.NotLoaded<:calculation, field: :years_ago>,
      ...
    >,
    #Tunez.Music.Album<
      year_released: 2012,
      string_years_ago: "wow, this was released 13 years ago!",
      years_ago: #Ash.NotLoaded<:calculation, field: :years_ago>,
      ...
    >,
  ],
  ...
>}
```

```
>
],
...
>}
```

You only get back what you request!

One important thing to note here is that Ash will only return the calculations you requested, even if some extra calculations are evaluated as a side-effect.



In the above example, Ash will calculate the `years_ago` field for each artist record, because it's needed to calculate `string_years_ago` — but `years_ago` will *not* be returned as part of the Artist data.

This is to avoid accidentally relying on these implicit side-effects. If we changed how `string_years_ago` is calculated, to not use `years_ago`, it would break any usage of `years_ago` in our views. This can lead to hard-to-debug issues!

Calculations are a really, really powerful tool. They can be used for simple data formatting, like our `string_years_ago` example here, or complex tasks like building tree data structures out of a flat data set, or pathfinding in a graph. Calculations can also work with resource relationships and their data, and here we get to what we actually want to build for Tunetz.

Calculations with related records

Tunetz is recording all this interesting album information for each artist, but not showing any of it in the artist catalog. So we'll use calculations to surface some of it as part of the loaded Artist data, and display it on the page.

There are three pieces of information we really want:

- The number of albums each artist has released,
- The year that each artist's latest album was released in, and
- The most recent album cover for each artist

Let's look at how we can build each of those with calculations!

Counting albums for an artist

Ash provides the `count/2` expression function,¹⁵ also known as an inline *aggregate* function (we'll see why shortly), that we can use to count records in a relationship.

15. <https://hexdocs.pm/ash/expressions.html#sub-expressions>

So to count each artist's albums as part of a calculation, we could add it as a calculation in the Artist resource:

```
defmodule Tunez.Music.Artist do
  # ...
  calculations do
  >    calculate :album_count, :integer, expr(count(albums))
  end
end
```

Testing this in iex, you can see it makes a pretty efficient query, even when querying multiple records. There's no n+1 query issues here, it's all handled in one query through a clever join:

```
iex(1)> Tunez.Music.search_artists("a", load: [:album_count])
SELECT a0."id", a0."name", a0."biography", a0."previous_names",
a0."inserted_at", a0."updated_at", coalesce(s1."aggregate_0", $1::bigint)
::bigint::bigint::bigint FROM "artists" AS a0 LEFT OUTER JOIN LATERAL
(SELECT sa0."artist_id" AS "artist_id", coalesce(count(sa0."id"::uuid),
$2::bigint)::bigint AS "aggregate_0" FROM "public"."albums" AS sa0 WHERE
(a0."id" = sa0."artist_id") GROUP BY sa0."artist_id") AS s1 ON TRUE WHERE
(a0."name"::text ILIKE $3) ORDER BY a0."id" LIMIT $4 [0, 0, "%a%", 13]
{:ok, %Ash.Page.Offset{...}}
```

It's a little bit icky with some extra type-casting that doesn't need to be done, but we'll address that shortly (this isn't the final form of our calculation!)

Finding the most recent album release year for an artist

This one is a little bit trickier. We're working with relationship data again, so we'll use another inline aggregate function. Because we've ensured that albums are always ordered according to release year in [Setting default ordering with a preparation, on page 46](#), the first album in the list of related albums will always be the most recent.

The first aggregate function is used to fetch a specific attribute value from the first record in the relationship, so we can use it to pull out just the year_released value from the album, and store it in a new attribute on the Artist.

```
calculations do
  calculate :album_count, :integer, expr(count(albums))
  >    calculate :latest_album_year_released, :integer,
  >      expr(first(albums, field: :year_released))
end
```

Finding the most recent album cover for an artist

This is very similar to, but a slight twist on, the previous calculation. Again we want the most recent album, but only out of albums that have the

optional `cover_image_url` attribute specified. We *could* add this extra condition using the `filter` option on the base query, like we did when we set a sort order in [The base query for a read action, on page 70](#), but we don't actually need to — for convenience, Ash will filter out `nil` values automatically. Note that the calculation can still return `nil` if an artist has no albums at all, or has no albums with album covers.

Everything combined, our calculation can look like this:

```
calculations do
  calculate :album_count, :integer, expr(count(albums))
  calculate :latest_album_year_released, :integer,
    expr(first(albums, field: :year_released))

  > calculate :cover_image_url, :string,
  >   expr(first(albums, field: :cover_image_url))
end
```



If you don't want this convenience, if you really do want the cover for the most recent album even if it's `nil`, you can add the `include_nil?: true` option to the first inline-aggregate function call.

And this works! We can specify any or all of these three calculation names, `:album_count`, `:latest_album_year_released`, and `:cover_image_url`, when loading artist data and get the attributes properly calculated and returned. This is really powerful; and we've only scratched the surface of what you can do with calculations.

The code here is starting to get a little bit gnarly, though. Our three calculations have one thing in common — they all use *inline aggregate* functions, to surface some attribute or derived value from relationships. Instead of defining the aggregates inline, we can look at extracting them into full aggregates, and see how that cleans up the code.

Relationship Calculations as Aggregates

Aggregates are a specialized type of calculation, as we've seen above. All aggregates are calculations; but a calculation like `years_ago` in our Album example was not an aggregate.

Aggregates perform some kind of calculation on records in a relationship — it could be a simple calculation like `first` or `count`, or a more complicated calculation like `min` or `avg` (average), or you can even provide a fully-custom implementation if the full list of aggregate types¹⁶ doesn't have what you need.

16. <https://hexdocs.pm/ash/aggregates.html#declaring-aggregates-on-a-resource>

To start adding aggregates to our Artist resource, we first need to add the aggregates block at the top level of the resource. (You might be sensing a pattern about this, by now.)

```
03/lib/tunez/music/artist.ex
defmodule Tunez.Music.Artist do
  # ...
  aggregates do
    end
  end
```

Each of the three inline-aggregate calculations we defined, can be rewritten to be an aggregate within this block. An aggregate needs at least three things — the type of aggregate, the name of the attribute to be used for the result value, and the relationship to be used for the aggregate.

So our example of the album_count calculation:

```
calculate :album_count, :integer, expr(count(albums))
```

Could be written more appropriately as a count aggregate:

```
03/lib/tunez/music/artist.ex
aggregates do
  > count :album_count, :albums
end
```

We don't need to specify the type of the resulting attribute — Ash knows that a count is always an integer, it can't be anything else, even if it's zero. This also simplifies the generated SQL a little bit, there's no need for repeatedly casting things as bigints.

Our latest_album_year_released calculation can be rewritten similarly:

```
03/lib/tunez/music/artist.ex
aggregates do
  count :album_count, :albums
  # old: calculation
  # calculate :latest_album_year_released, :integer,
  #   expr(first(albums, field: :year_released))
  > first :latest_album_year_released, :albums, :year_released
end
```

We've dropped a little bit of the messy syntax, and the result is a lot easier to read. We don't need to define that latest_album_year_released is an integer — that can be inferred because the Album resource already defines the year_released attribute as an integer. If the syntax seems a bit mysterious (where

did the field key go?), the options available for each type of aggregate are fully laid out in the Ash.Resource DSL documentation.¹⁷

The final calculation, for cover_image_url, is the same as for latest_album_year_released. The include_nil?: true option can be used here too, if you really want a cover that might be nil, but we'll rely on the default value of false. If a given artist has *any* awesome album covers, we want the most recent one.

```
03/lib/tunez/music/artist.ex
aggregates do
  count :album_count, :albums
  first :latest_album_year_released, :albums, :year_released
  # old: calculation
  # calculate :cover_image_url, :string,
  #   expr(first(albums, field: :cover_image_url))
  > first :cover_image_url, :albums, :cover_image_url
end
```

In this way, we can put all the logic of how to calculate a latest_album_year_released or a cover_image_url for an artist where it belongs, in the domain layer of our application, and our front-end views don't have to worry about where it might come from. On that note, let's integrate these aggregates in our artist catalog.

Using aggregates like any other attribute

It would be amazing if the artist catalog looked like a beautiful display, of album artwork and artist information.

The screenshot shows a web-based artist catalog. At the top, there's a header with the title "Artists", a "sort by:" dropdown set to "recently updated", a search bar with a magnifying glass icon, and a blue "New Artist" button. Below the header, there are three artist profiles, each featuring a large thumbnail image and the artist's name.

- Cielarko**: Shows a vibrant, colorful landscape painting of rolling hills under a rainbow sky. Below the image, it says "3 albums, latest release 2023".
- Zombie Kittens!!**: Shows a dark, atmospheric illustration of a black cat with glowing orange eyes sitting in a dimly lit alleyway at night, with a full moon in the background. Below the image, it says "3 albums, latest release 2023".
- Crystal Cove**: Shows a dramatic, moody illustration of a sailing ship on a choppy sea under a bright, glowing moon in a cloudy sky. Below the image, it says "2 albums, latest release 2024".

17. <https://hexdocs.pm/ash/dsl-ash-resource.html#aggregates>

In Tunez.Artists.IndexLive, the cover image display is handled by a call to the cover_image function component within artist_card:

```
03/lib/tunez_web/live/artists/index_live.ex
<div id={"artist-#{@artist.id}"} data-role="artist-card" class="relative mb-2">
  <.link navigate={~p"/artists/#{@artist}"}>
    <.cover_image />
  </.link>
</div>
```

Because we can use and reference aggregate attributes like any other attributes on a resource, we can add an image argument to the cover_image function component, to replace the default placeholder image with our cover_image_url calculation:

```
03/lib/tunez_web/live/artists/index_live.ex
<div id={"artist-#{@artist.id}"} data-role="artist-card" class="relative mb-2">
  <.link navigate={~p"/artists/#{@artist}"}>
    >   <.cover_image image={@artist.cover_image_url} />
  </.link>
</div>
```

Refreshing the artist catalog after making the change might not be what you expect — why aren't the covers displaying? Because we aren't loading them! Remember that we need to specifically load calculations/aggregates if we want to use them, they won't be generated automatically.

Because we'll want all of our new aggregates loaded every time we run the search action on the Artist resource, and it's not a dynamic list in any way, we can add a preparation¹⁸ to that action to load our data.

```
03/lib/tunez/music/artist.ex
read :search do
  # ...
  prepare build(load: [:album_count, :latest_album_year_released,
    :cover_image_url])
end
```

Reloading the artist catalog will now populate the data for all of the aggregates we listed, and look at the awesome artwork appear! Special thanks to Midjourney, for bringing our imagination to life!

For the album count and latest album year released fields, we can add those details to the end of the artist_card function, using the previously-unused artist_card_album_info component defined right below it:

18. <https://hexdocs.pm/ash/Ash.Resource.Preparation.Builtins.html#build/1>

```
03/lib/tunez_web/live/artists/index_live.ex
def artist_card(assigns) do
  ~H"""
  # ...
  <.artist_card_album_info artist={@artist} />
  """
end
```

And behold! The artist catalog is now in its full glory!

Earlier in the chapter, we looked at sorting artists in the catalog, via three different attributes - name, inserted_at and updated_at. We've explicitly said a few times now, that calculations and aggregates can be treated like any other attribute — does that mean we might be able to sort on them too???

You bet you can!

Sorting based on aggregate data

Around this point is where Ash really starts to shine, and you might start feeling a bit of a tingle with the power at your fingertips. Hold that thought, because it's going to get even better. Let's add some new sort options for our aggregate attributes, to our list of available sort options in Tunez.Music.IndexLive:

```
03/lib/tunez_web/live/artists/index_live.ex
def sort_options do
  [
    {"-updated_at", "recently updated"},
    {"-inserted_at", "recently added"},
    {"name", "name"},
    {"-album_count", "number of albums"},
    {"--latest_album_year_released", "latest album release"}
  ]
end
```

We want artists with the most albums and with the most recent albums listed first, so we'll sort them descending by prefixing the attribute name with a -. Using -- is a bit special - it'll put any nil values (if an artist hasn't released any albums!) at the end of the list.

To allow the aggregates to be sorted on, we do need to mark them as public? true, like we did with our initial set of sortable attributes in [Using sort_input for succinct yet expressive sorting, on page 71](#):

```
03/lib/tunez/music/artist.ex
aggregates do
  count :album_count, :albums do
    public? true
  end
```

```
>   first :latest_album_year_released, :albums, :year_released do
    public? true
  end

# ...
end
```

And then we'll be able to sort in our artist catalog, to see which artists have the most albums, or have released albums most recently:

The screenshot shows a web application interface for managing artists. At the top, there's a header with the title "Artists". Below the header, there's a search bar with a dropdown menu set to "sort by: number of albums" and a search icon. To the right of the search bar is a blue button labeled "New Artist".

Three artist profiles are displayed in a row:

- Violet Depths**: An abstract painting with red, orange, and purple hues.
- Tangerine**: An abstract painting with swirling blue and teal colors.
- Infiniverse**: An abstract painting with a mix of orange, yellow, blue, and purple.

Below each profile, there's a brief description:

- Violet Depths**: 7 albums, latest release 2023
- Tangerine**: 5 albums, latest release 2022
- Infiniverse**: 4 albums, latest release 2007

This is all *amazing!* We've built some really excellent functionality over the course of this chapter, to let users search, sort, and paginate through data.

And in the next one, we'll see how we can use the power of Ash to build some neat APIs for Tunez, using our existing resources and actions. Reduce, re-use and recycle code!

Generating APIs Without Writing Code

In the previous chapter, we looked at making the artist catalog a lot more user-friendly by letting users search, sort and filter artist data. This was a big boost to Tunez's popularity, so much so that some users are asking how they can use the data from Tunez in their own apps.

We can give users access to a Tunez *application programming interface*, or *API* — a way of letting their apps talk to Tunez, to fetch or modify data. APIs are everywhere, whenever we build apps that can communicate with other apps, we're doing it via an API. If you're connecting to Facebook to read a user's friends list, or built an app that upload photos to a image-hosting service like Cloudinary, you're using those services's APIs.

Let's look at how we can build an API for Tunez to let other apps talk to *us*, using the resources and actions we've defined so far.

Model Your Domain, Derive The Rest

One of the core design principles¹ of Ash is its *declarative* and *derivable* nature. By themselves, resources are static configuration files that Ash can interpret and generate code from. We've seen examples of this with code interfaces for our actions — we *declared* that we should have an interface for our Artist :search action that accepts one argument for the query text, and Ash generated the function for us to call.

This can be taken further — Ash can generate a lot more than functions. It can generate entire APIs around the existing resources and actions in your app, hence the name of this chapter.

1. <https://hexdocs.pm/ash/design-principles.html>

It sounds wild, but what *is* an API, really? A set of functions, that map from an input URL to a response in a format like JSON.² Our web UI is an API, albeit a heavily customized one that returns HTML. An API using something like GraphQL³ or REST⁴ is a lot more standardized. Both the incoming requests and the outgoing responses have a very strict format to adhere to, and that can be generated for us using Ash.

We'll build two APIs in this chapter, using both REST and GraphQL. In a real app you'd probably want one or the other, but we'll show off a little bit here and add both. Let's go!

Building a JSON REST Interface

A REST (or RESTful) API can be generated by Ash using the `ash_json_api` package. This will accept requests over HTTP, and return data formatted as JSON. As a bonus, APIs generated with `ash_json_api` are compatible with the JSON:API⁵ specification, and can optionally also generate OpenAPI⁶ schemas, opening up a whole world of supporting tooling options.

Setup

You can add `ash_json_api` to Tunez using the `igniter.install` Mix task:

```
$ mix igniter.install ash_json_api
```

This will add a few new pieces to your app, that have a *lot* of power. The additions include:

- The `ash_json_api` Hex package, as well as its sibling dependency `open_api_spx` (in `mix.exs` and `mix.lock`)
- Code formatting and configuration to support a new application/vnd.api+json media type, needed for JSON:API compatibility⁷ (in `config/config.exs`)
- A new scope in your Phoenix router to accept web requests for the `/api/json/` URL (in `lib/tunez_web/router.ex`)
- A new `TunezWeb.AshJsonApiRouter` module that uses `AshJsonApi.Router`. This will process the web requests and return responses in the correct format (in `lib/tunez_web/ash_json_api_router.ex`)

2. <https://stackoverflow.blog/2022/06/02/a-beginners-guide-to-json-the-data-format-for-the-internet/>

3. <https://graphql.org/>

4. <https://www.ibm.com/topics/rest-apis>

5. <https://jsonapi.org/>

6. <https://www.openapis.org/>

7. <https://jsonapi.org/format/#introduction>

This takes care of a lot of the boilerplate around a REST API, leaving us to handle the implementation of what our API should actually *do*.

Adding Artists to the API

What we primarily want to expose in our API is the CRUD interface for our resources to let users manage artist and album data over the API. Each of our resources can be exposed as a type/schema/definition and each action on a resource exposed as an operation.

By default, the API is empty — we have to manually include each resource and action we want to make public. To add a resource to the API, we can use Ash's patch generator to *extend* the resource with the `AshJsonApi.Resource` extension:

```
$ mix ash.patch.extend Tunez.Music.Artist json_api
```

This will make some handy changes in our app:

- `AshJsonApi.Resource` will be added as an extension to the `Tunez.Music.Artist` resource
- A default API “type” will be added to the resource, in a new `json_api` block in the resource. Each record in a API response is identified by an `id` and a `type` field,⁸ the type usually being a string version of the resource name.

And because this is the first resource in the `Tunez.Music` domain to be configured for `AshJsonApi`, the patch generator will also connect pieces in the domain:

- `AshJsonApi.Domain` will be added as an extension to the `Tunez.Music` domain
- And the `Tunez.Music` domain will be added to the list of domains configured in the `TunezWeb.AshJsonApiRouter` module.

You could make all the changes yourself manually, but there's a few moving parts there and it can be easy to miss a connection. The generators are a convenient way of making sure everything is set up as it should be.

Next, to make the actions on the `Artist` resource available in the API, we need to set up routes for them. Like code interfaces, this can be done either on the resource or the domain. However, to keep the domain as the solid boundary between our domain model and the outside world, we'll add them on the domain.

8. <https://jsonapi.org/format/#document-resource-object-identification>

In a new top-level json_api block in the Tunez.Music domain module, configure the routes using the DSL provided by AshJsonApi:⁹

```
04/lib/tunez/music.ex
defmodule Tunez.Music do
  # ...
  json_api do
    routes do
      base_route "/artists", Tunez.Music.Artist do
        get :read
        index :search
        post :create
        patch :update
        delete :destroy
      end
    end
  end
end
```

This code will connect a GET request to read a single artist by a given ID to the read action of the Tunez.Music.Artist resource, automatically applying the correct filter. A POST request will be connected to the create action, and so on.

As Ash also holds *introspectability* as one of its core design principles,¹⁰ you can fetch a list of generated routes and the actions they connect to in iex using AshJsonApi.Domain.Info:

```
iex(1)> AshJsonApi.Domain.Info.routes(Tunez.Music)
iex(2)> |> Enum.map(fn r -> {r.route, r.method, r.resource, r.action} end)
[
  {"/artists/:id", :get, Tunez.Music.Artist, :read},
  {"/artists", :get, Tunez.Music.Artist, :search},
  {"/artists", :post, Tunez.Music.Artist, :create},
  {"/artists/:id", :patch, Tunez.Music.Artist, :update},
  {"/artists/:id", :delete, Tunez.Music.Artist, :destroy}
]
```

So how can we actually *use* the API? For GET requests, you can access the endpoints provided in a browser like any other URL. Alternatively, you could use a dedicated API client app such as Bruno,¹¹ shown here making a GET request to /api/json/artists:

9. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-domain.html#json_api-routes

10. <https://hexdocs.pm/ash/design-principles.html>

11. <https://www.usebruno.com/>

```

1  {
2    "data": [
3      {
4        "attributes": {
5          "name": "Crystal Cove",
6          "inserted_at": "2024-08-13T00:17:20.933370Z",
7          "updated_at": "2024-08-13T00:22:07.215726Z"
8        },
9        "id": "465ef0f7-bb52-4dbd-a377-e1fabea9a720",
10       "links": {},
11       "meta": {},
12       "type": "artist",
13       "relationships": {}
14     }
15   ],
16   "links": {},
17   "meta": {},
18   "jsonapi": {}
19 }
20
21
22
23
24
25
26
27
28

```

Don't forget the custom headers!



While not strictly required for GET requests, you should configure your API client to add the correct Content-Type and Accept headers when making any requests to your API. These can usually be set on either the individual request level or the collection level, depending on your API client.

The value for both Content-Type and Accept headers should be application/vnd.api+json.

The /api/json section of the URL matches the scope that our AshJsonApi router is mounted in, the Phoenix router, and /artists matches the base route for the Tunez.Music.Artist resource, meaning this request will connect to the search action of the resource.

The action accepts a query argument that can be passed in as a query string parameter, and the search results are returned in a neat JSON format. Links for pagination are automatically included, because the action supports pagination. And we barely needed to lift a finger!

What data gets included in API responses?

You might notice that some attributes are missing in the API response — the artist name is shown, but the biography and previous_names are missing, as are the aggregates for album_count, cover_image_url and latest_album_year_released that we added in the last chapter.

This is because only *public attributes* (attributes that are specifically marked `public?: true`) are returned in API queries, by default. This is for security reasons — if all attributes were included by default, it would be really easy to accidentally leak information as you add more data to your resources, if you didn't also explicitly *remove* them from your API.

Some of the attributes are already public, such as those we used for sorting in the previous chapter. To add biography and previous_names to the API response, you can also mark them as `public?: true` in the Tunez.Music.Artist resource:

```
04/lib/tunez/music/artist.ex
attributes do
  # ...
  attribute :biography, :string do
    >   public? true
    end

  attribute :previous_names, {:array, :string} do
    default []
    >   public? true
    end
  # ...
end
```

Aggregates are a little different. For this usage, they are *not* treated like every other attribute, and aren't included by default even if they're public. This is because calculations and aggregates can be computationally expensive, and if they aren't specifically needed by users of the API, you can save time and effort by not calculating and returning them.

There are still two ways that you can make calculations and aggregates visible in your API:

- If you *do* want them to be calculated and returned by default, you can use the `default_fields` config option, eg. `default_fields [:id, :name, :biography, :album_count]`. This can be set at the resource level,¹² to apply any time an instance of the resource is returned in a response, or for any specific API route (either in the domain¹³ or in a resource).¹⁴ This will replace the default “return all public attributes” behaviour, though, so you'll have to list *all* fields that should be returned by default, including any public attributes.

12. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-resource.html#json_api-default_fields

13. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-domain.html#json_api-routes-base_route-get-default_fields

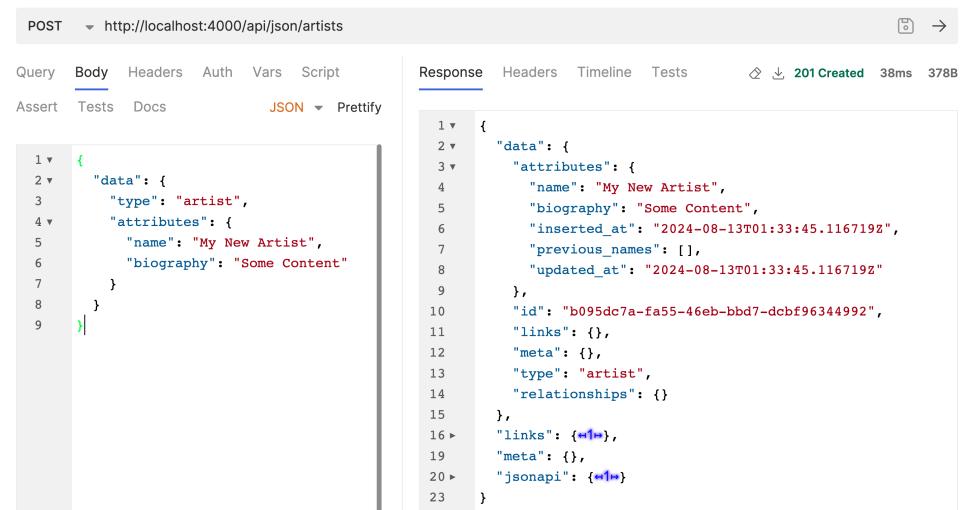
14. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-resource.html#json_api-routes-get-default_fields

- Alternatively, part of the JSON:API spec¹⁵ states that users can request which specific fields they want to fetch as part of their API request. Our API is JSON:API-compliant, so users can add the `fields` query string parameter and list only the fields they need in a comma-separated list. The fields can be any public fields, including aggregates and calculations, so a URL like `http://localhost:4000/api/json/artists?fields=name,album_count` would return only names and the number of albums for each artist in the search results.

Creating artist records

We won't cover *every* endpoint we created, but it's worth a quick look at how data can be created, as well as how it can be read.

As our introspection showed earlier, we can make POST requests to the same URL we used for searching, to access the create action of our resource. We can post a JSON object in the format specified in the JSON:API specification,¹⁶ containing the content for the artist record to be created.



```

POST http://localhost:4000/api/json/artists
Content-Type: application/vnd.api+json

{
  "data": {
    "type": "artist",
    "attributes": {
      "name": "My New Artist",
      "biography": "Some Content"
    }
  }
}

HTTP/1.1 201 Created
Content-Type: application/vnd.api+json; charset=UTF-8
Content-Length: 449
Date: Mon, 13 Aug 2024 13:45:11 GMT
Server: Apache/2.4.46 (Ubuntu)
Set-Cookie: PHPSESSID=1234567890; expires=Mon, 13-Aug-2024 13:45:11 UTC; path=/; secure; HttpOnly
X-Powered-By: PHP/8.2.14

{
  "data": {
    "id": "b095dc7a-fa55-46eb-bbd7-dcbf96344992",
    "type": "artist",
    "attributes": {
      "name": "My New Artist",
      "biography": "Some Content"
    },
    "links": {},
    "meta": {}
  }
}

```

In your Phoenix server logs, you can see the create request being handled by the `AshJsonApiRouter` module and processed:

```
[info] POST /api/json/artists
[debug] Processing with TunezWeb.AshJsonApiRouter
Parameters: %{"data" => %{"attributes" => %{"biography" => "Some Content", "name" => "My New Artist"}, "type" => "artist"}}
```

15. <https://jsonapi.org/format/#fetching-sparse-fieldsets>

16. <https://jsonapi.org/format/#crud-creating>

```
Pipelines: [:api]
[debug] QUERY OK source="artists" db=0.1ms
INSERT INTO "artists" ("id", "name", "biography", "inserted_at",
"previous_names", "updated_at") VALUES ($1,$2,$3,$4,$5,$6) RETURNING
"updated_at", "inserted_at", "previous_names", "biography", "name", "id"
[«uuid», "My New Artist", "Some Content", «timestamp», [], «timestamp»]
```

Because this API endpoint connects to the create action in our `Tunez.Music.Artist` resource, it accepts all of the same data as the action does. Posting additional data (e.g., an attribute that the action doesn't accept or a non-existent attribute) or invalid data (e.g., a missing required field) will return an error message, and the record won't be created.

Other requests can be made in a similar way — a PATCH request to update an existing artist, and a DELETE request to delete an artist record.

Adding Albums to the API

We can add album management to the JSON API in much the same way we added artists, by extending the `Tunez.Music.Album` resource:

```
$ mix ash.patch.extend Tunez.Music.Album json_api
```

And adding our routes in the `json_api` block in the domain:

```
04/lib/tunez/music.ex
json_api do
  routes do
    # ...
  >   base_route "/albums", Tunez.Music.Album do
  >     post :create
  >     patch :update
  >     delete :destroy
  >   end
  end
end
```

This closely resembles the web UI. We don't have an API endpoint to list all albums, but we do have endpoints to manage individual album records. This will create URLs like `/api/json/albums` and `/api/json/albums/:album_id`, with various HTTP methods to connect to the different actions in the resource.

Because we're not in the web UI, though, we don't have the nice pre-filled hidden artist ID when submitting a HTTP request to create an album — we need to provide a valid one with the attributes of the album to be created, like this:

```
{
  "data": {
```

```

    "type": "album",
    "attributes": {
      "name": "New Album",
      "artist_id": [a-valid-uuid],
      "year_released": 2022
    }
  }
}

```



You can also create endpoints like `/api/json/artists/[artist-uuid]/relationships/albums` for creating albums, using the `post_to_relationship`¹⁷ function for the route. This is a little bit different though, and is what we'd typically use to match the UI if the UI allowed you to update multiple album records from the artist edit form. We'll see an example of this *the (as yet) unwritten content*.

As part of this, we can also mark some of the attributes on the `Tunez.Music.Album` resource as `public?: true`, such as `name`, `year_released`, and `cover_image_url`, so they can be returned in API responses.

```
04/lib/tunez/music/album.ex
attributes do
  uuid_primary_key :id

  attribute :name, :string do
    allow_nil? false
    public? true
  end

  attribute :year_released, :integer do
    allow_nil? false
    public? true
  end

  attribute :cover_image_url, :string do
    public? true
  end

  # ...

```

There's just one part we're missing now — listing an artist's albums.

Showing albums for a given artist

The JSON:API spec allows for two methods of fetching related resources¹⁸ for a given resource. We'll cover both methods; you can choose the one that suits you when building your own APIs.

17. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-domain.html#json_api-routes-base_route-post_to_relationship

18. <https://jsonapi.org/format/#fetching-relationships>

Both methods require the relationship to be *public* to be accessible over the API, so you'll need to mark it as `public?: true` in `Tunez.Music.Artist`:

```
04/lib/tunez/music/artist.ex
relationships do
  has_many :albums, Tunez.Music.Album do
    public? true
  end
end
```

Including related records

This is the easiest way to provide relationship data, and it mirrors what we see in the web UI when we view an Artist profile page. You can allow related records to be *included* when fetching the parent resource — returning the artist record, and their albums, in one request. This is convenient for consumers of the API as they only need to make a single request, but the responses can be large and they may overfetch data.

To enable this in our API, edit the `json_api` block in the `Tunez.Music.Artist` resource to list which relationships can be included from this resource:

```
04/lib/tunez/music/artist.ex
json_api do
  type "artist"
  includes [:albums]
end
```

You'll then be able to fetch album data by adding `include=albums` to the query string of any request for artist data, such as <http://localhost:4000/api/json/artists?query=cove&include=albums>. The response will have a list of record identifiers under the `relationships` key of the fetched data, and then a separate list of the full records under the top-level `included` key. The format is a little quirky, but it's the JSON:API way!

Linking to a list of related records

For a different approach, you can return a link to fetch relationship data via a separate `related`¹⁹ route in your domain, and also specify which action should be used when fetching the related data.

```
04/lib/tunez/music.ex
base_route "/artists", Tunez.Music.Artist do
  # ...
  related :albums, :read, primary?: true
end
```

19. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-domain.html#json_api-routes-base_route-related

This will add a related relationship link to the response when generating any artist API response:

```
{
  "data": {
    // ...
    "relationships": {
      "albums": {
        "links": {
          "related": "http://localhost:4000/api/json/artists/«id»/albums"
        }
      }
    }
  }
}
```

Accessing the related URL will provide a list of the related records, in JSON format:

```
{
  "data": [
    {
      "attributes": {
        "name": "Shadows of Catnip",
        "year_released": 2023,
        "cover_image_url": "/images/albums/zombie_kittens/shadows_of_catnip.png",
        // ...
      }
    }
  ]
}
```

Which is great! Now we've got all the same functionality from our web UI, accessible over a JSON API. As we make further changes to our actions to add more functionality, they'll all automatically flow through to our API endpoints as well.

Now how can we get the word out about Tunez's fantastic new API?

We can auto-generate some documentation that we can share publicly to show people how to integrate with the API really easily!

Generating API documentation with OpenApiSpecx

When we installed AshJsonApi, it also added a package called `open_api_spepx`²⁰ to our `mix.exs` file, and this is how we can generate OpenAPI specifications automatically for our API. We don't have to do anything to set it up — the `AshJsonApi` installer did so when it created the JSON API router in `lib/tunez_web/ash_json_api_router.ex`:

```
04/lib/tunez_web/ash_json_api_router.ex
defmodule TunezWeb.AshJsonApiRouter do
  use AshJsonApi.Router,
  domains: [Module.concat([Tunez.Music])],
```

20. https://hexdocs.pm/open_api_spepx/

```
>     open_api: "/open_api"
end
```

This one line of code will give you a full OpenAPI specification document at the provided route, http://localhost:4000/api/json/open_api by default.

We can use this specification document with any tool or library that works with OpenAPI (and there are many!)²¹ One that OpenApiSpex provides support for out of the box is Swagger UI,²² to generate full documentation for our API and even let users try out endpoints directly from the docs.

OpenApiSpex's SwaggerUI plug has already been set up in our router, in lib/tunez_web/router.ex:

```
04/lib/tunez_web/router.ex
scope "/api/json" do
  pipe_through [:api]

  > forward "/swaggerui",
  >           OpenApiSpex.Plug.SwaggerUI,
  >           path: "/api/json/open_api",
  >           default_model_expand_depth: 4

  forward "/", TunezWeb.AshJsonApiController
end
```

This sets up the /api/json/swaggerui URL, with a full set of Swagger UI API documentation:

The screenshot shows the Swagger UI interface. At the top, it displays the URL http://localhost:4000/api/json/open_api. Below the URL is a green header bar with the text "Open API Specification 1.1 OAS 3.0". Underneath the header, there is a "Servers" dropdown menu set to "http://localhost:4000" and an "Authorize" button with a lock icon. The main content area is titled "album" and lists four API operations: POST /api/json/albums, DELETE /api/json/albums/{id}, PATCH /api/json/albums/{id}, and GET /api/json/artists/{id}/albums. Each operation is represented by a button with its method and URL, followed by a lock icon and a dropdown arrow.

21. <https://tools.openapis.org/>

22. <https://swagger.io/tools/swagger-ui/>

Totally for free! And it'll stay up to date as you update your API, adding or updating resources or actions.

If Swagger UI isn't to your liking, Redoc²³ is a good alternative. It can be installed in your app via the `redoc_ui_plug`²⁴ Hex package, and configured in your Phoenix router in a very similar way to Swagger UI.

And if you decide you don't want the documentation after all, you only need to remove the `SwaggerUI` plug from your router.

Customizing the generated API

Now that we have a great overview of our API, and we can see it the way a user would, we can see some places that it can be improved. These certainly aren't the *only* ways, but they're low-hanging fruit that will give quick wins.

Adding informative descriptions

While it's amazing that we can get all this for free, some of the defaults in the generated content can be a bit lacking. Our API probably shouldn't be called "Open API Specification". And `AshJsonApi` doesn't know what we really *mean* when we say "Get artists", so the default description of the search API endpoint is "/artists operation on artist resource". Not too great.

Ash allows us to add description metadata in a few different places, that will be picked up by the OpenAPI schema generator and added to the documentation. This includes:

- A description for a resource as a whole. This can be added as part of a top-level resource block, such as this in `Tunez.Music.Artist`:

```
defmodule Tunez.Music.Artist do
  use Ash.Resource, ...

  resource do
    description "A person or group of people that makes and releases music."
  end
```

- A description for an action, or argument for an action, in a resource. These can be added in the action declaration itself, such as:

```
read :search do
  description "List Artists, optionally filtering by name."
  argument :query, :ci_string do
    description "Return only artists with names including the given value."
```

23. <https://github.com/Redocly/redoc>

24. https://hexdocs.pm/redoc_ui_plug/index.html

```
# ...
```

As a nice bonus, these descriptions should be picked up by any Elixir-related language server packages in your text editor, such as ElixirLS or elixir-tools in VSCode.

Basic information that is OpenAPI-specific, such as the name of the API, can be customized via options to use `AshJsonApi.Router`, in your JSON API router module. If you need to make more specific changes, you can also add a `modify_open_api` hook function,²⁵ to be called when generating the OpenAPI spec. This function will have access to the whole generated spec, and there are a lot of things²⁶ that can be changed or overwritten, so be careful!

```
04/lib/tunez_web/ash_json_api_router.ex
defmodule TunezWeb.AshJsonApiRouter do
  use AshJsonApi.Router,
  domains: [Module.concat([Tunez.Music])],
  open_api: "/open_api",
  open_api_title: "Tunez API Documentation",
  open_api_version: to_string(Application.spec(:tunez, :vsn))
end
```

Once you've made any changes like descriptions, refreshing the Swagger UI docs will immediately reflect the changes, and they're looking a lot better now.

Removing unwanted extras

Looking through the docs carefully shows that our API can actually do a little bit more than we thought. Expanding the section for GET /api/json/artists, our artist search, shows the endpoint will allow data to be filtered via a filter parameter in the URL. This is pretty cool, but we already have our own specific filtering set up, to search artists by name. So while it sounds like a waste, we'll disable the generated filtering for parity with the web interface.

`AshJsonApi` provides both generated filtering *and* sorting of data, for any index actions in our API router. These can be disabled either at the resource level or per-action. For `Tunez`, we want to keep the generated sorting of artists because we allow that via the web, but disable the generated filtering. We can do that with the `derive_filter?` config option²⁷ in the `Tunez.Music.Artist` resource:

```
04/lib/tunez/music/artist.ex
json_api do
  type "artist"
```

25. https://hexdocs.pm/ash_json_api/open-api.html#customize-values-in-the-openapi-documentation

26. https://hexdocs.pm/open_api_spex/OpenApiSpex.OpenApi.html#t:t/0

27. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-resource.html#json_api-derive_filter?

```

    includes [:albums]
  >  derive_filter? false
end

```

And that's our JSON REST API, fully complete! It packs a lot of punch, for not a lot of code. We didn't have to write any endpoints, generate any JSON, worry about error handling — everything is handled by AshJsonApi, which generates API endpoints and controllers to connect the actions in our resources to the outside world. It's pretty nifty.

If JSON and REST aren't to your liking, maybe you're in the GraphQL camp. We can build a GraphQL API for Tunez in a very similar way!

Building a GraphQL Interface

A GraphQL API can be generated by Ash using the `ash_graphql` package. It's built on top of the excellent `absinthe`²⁸ library, so it's rock-solid and ready for production use. This will create a standard GraphQL endpoint, accepting GET requests over HTTP using GraphQL syntax and returning JSON responses.

GraphQL APIs are a little more flexible than REST APIs — though with the JSON:API specification, the gap is smaller than you might think. We won't debate the pros and cons of each type of API here, but both approaches can create well-defined, well-structured, and well-documented interfaces for your users to work with.

Setup

You can add `ash_graphql` to Tunez using the `igniter.install` Mix task:

```
$ mix igniter.install ash_graphql
```

This will add a few new pieces to your app, that have a *lot* of power. The changes include:

- A new `graphql` pipeline and scope in your Phoenix router, to accept requests for both `/gql`, the GraphQL endpoint, and `/gql/playground`, a GraphiQL API client (in `lib/tunez_web/router.ex`)
- A new `TunezWeb.GraphqlSchema` module, that uses `AshGraphql` and `Absinthe.Schema` and is seeded with a sample runnable query (in `lib/tunez_web/graphql_schema.ex`)
- Code formatting and configuration for Absinthe to support requests made with the `application/graphql` media type (in `lib/tunez_web/endpoint.ex`)

28. <https://hexdocs.pm/absinthe/>

This takes care of all of the boilerplate around setting up a GraphQL API. After restarting your Phoenix server, you can test out the GraphiQL playground by accessing <http://localhost:4000/gql/playground> — there isn't a lot to see there at the moment, but we do have a generated schema with the sample query that AshGraphql provides when no other queries are present. Hello, AshGraphql!

Now we can look at what the API needs to actually *do*.

Adding Artists to the API

What we primarily want to expose in our GraphQL API is the CRUD interface for our resources, to let users manage artist and album data over the API. Each of our resources can be exposed as a type, and each action on a resource exposed as either a query or a mutation.

By default, the API is empty — we have to manually include each resource and action that we want to make public. To add a resource to the API, we can use Ash's patch generator to *extend* the resource with the `AshGraphql.Resource` extension:

```
$ mix ash.patch.extend Tunez.Music.Artist graphql
```

This will make some handy changes to our app:

- `AshGraphql.Resource` will be added as an extension to the `Tunez.Music.Artist` resource
- A default GraphQL type will be added to the resource, in a new `graphql` block in the resource. This is usually a simplified atom version of the resource name.

And because this is the first resource in the `Tunez.Music` domain to be configured for AshGraphql, the patch generator will also connect pieces in the domain:

- `AshGraphql.Domain` will be added as an extension to the `Tunez.Music` domain
- And the `Tunez.Music` domain will be added to the list of domains configured in the `TunezWeb.GraphqlSchema` module.

You could make all the changes yourself manually, but there's a few moving parts there and it can be easy to miss a connection. The generators are a convenient way of making sure everything is set up as it should be.

Next, to make the actions on the `Artist` resource available in the API, we need to create queries and mutations for them. Like code interfaces, this can be done either on the resource or the domain, but to keep the domain as the solid boundary between our domain model and the outside world, we'll add them on the domain.

In the top-level `graphql` block defined in the `Tunez.Music` domain model, we can add queries for read actions of our `Artist` resource. AshGraphql provides macros like `get` and `list`²⁹ for this, which describe what kind of responses we expect from the queries.

```
04/lib/tunez/music.ex
defmodule Tunez.Music do
  # ...
  graphql do
    queries do
      get Tunez.Music.Artist, :get_artist_by_id, :read
      list Tunez.Music.Artist, :search_artists, :search
    end
  end
end
```

This will create GraphQL queries named `getArtistById`, connecting to the `read` action of the `Tunez.Music.Artist` resource and automatically applying an ID filter; and `searchArtists` connecting to the `search` action.

We can do the same for the non-read actions in our resource, which will all be mutations in the API:

```
04/lib/tunez/music.ex
graphql do
  # ...
  mutations do
    create Tunez.Music.Artist, :create_artist, :create
    update Tunez.Music.Artist, :update_artist, :update
    destroy Tunez.Music.Artist, :destroy_artist, :destroy
  end
end
```

This gives us a lot out of the box. In the GraphiQL playground, expanding the Schema tab now shows the queries and mutations we just defined, and they're fully-typed — a `getArtistById` query will return an `Artist` type, with all public attributes of the resource also typed and available to be requested. We can run any query, and fetch data in the shape we want.

29. https://hexdocs.pm/ash_graphql/dsl-ashgraphql-domain.html#graphql-queries



```

1 ▾ {
2 ▾   searchArtists(query: "cove") {
3     results {
4       name
5       albumCount
6     }
7   }
8 }
```

```

    ▾   "data": {
    ▾     "searchArtists": {
    ▾       "results": [
    ▾         {
    ▾           "albumCount": 2,
    ▾           "name": "Crystal Cove"
    ▾         }
    ▾       ]
    ▾     }
    ▾   }
}
```

The search action accepts a query argument, which means that the generated searchArtists query also accepts a query argument. Because the action also supports pagination, the request and response both support pagination, and it's all right there in the generated types. We barely needed to lift a finger!

What data gets included in API responses?

If you skipped over the JSON API section because GraphQL is much more interesting, you might be surprised to see that fields like biography and previousNames aren't defined in the GraphQL Artist type.

It was hinted at in the previous section, but it's really important to note that only *public attributes* (attributes that are specifically marked `public?: true`) can be requested and returned in GraphQL responses. This is for security reasons — if all attributes were included by default, it would be really easy to accidentally leak information as you add more data to your resources, if you didn't also explicitly *remove* them from your API.

By adding `public? true` to those attributes in the Tunez.Music.Artist resource:

```
04/lib/tunez/music/artist.ex
attributes do
  # ...
  attribute :biography, :string do
    >   public? true
  end

  attribute :previous_names, {:array, :string} do
    >   default []
    >   public? true
  end
  # ...
end
```

They'll then be added to the GraphQL Artist type and can be requested like any other field.

Aggregates and calculations must also be marked as `public: true` if you want to make them accessible in the API.

Creating artist records

We won't cover every operation we created, but it's worth a quick look at how data can be created, as well as how it can be read.

Expanding the Schema tab in the playground shows that we can call a mutation named `createArtist` for creating new Artist records. Because it connects to the `create` action in the `Tunez.Music.Artist` resource, the attributes the action accepts is matched by the typing of the input to the mutation.

```

1 mutation {
2   createArtist(input: {
3     name: "Unleash the Rangers",
4     biography: "A great Canadian band"
5   }) {
6     errors { fields message }
7     result { name albumCount }
8   }
9 }
```

In your Phoenix server logs, you can see the mutation being handled by Absinthe using the `TunezWeb.Schema` module, and processed:

```
[debug] ABSINTHE schema=TunezWeb.Schema variables=%{}
---
mutation {
  createArtist(input: {
    name: "Unleash the Rangers",
    biography: "A great Canadian band"
  }) {
    errors { fields message }
    result { name albumCount }
  }
}
---
[debug] QUERY OK source="artists" db=2.9ms
INSERT INTO "artists" ("id", "name", "biography", "inserted_at", "previous_names", "updated_at") VALUES ($1,$2,$3,$4,$5,$6) RETURNING "updated_at", "inserted_at", "previous_names", "biography", "name", "id" [«uuid», "Unleash the Rangers", "A great Canadian band", «timestamp», [], «timestamp»]
```

If the submitted data passes the input type checking, but fails resource validation (such as an empty name value), the mutation will return a nicely typed error message and the record won't be created. And because the action and mutation will return the record being created, on success, we can use all the usual GraphQL ideas of requesting only the fields we need in the response.

Adding Albums to the API

We can add album management to the GraphQL API in much the same way we added artists, by extending the Tunez.Music.Album resource:

```
$ mix ash.patch.extend Tunez.Music.Album graphql
```

And adding our mutations to the graphql block in the domain:

```
04/lib/tunez/music.ex
graphql do
  mutations do
    # ...
  >  create Tunez.Music.Album, :create_album, :create
  >  update Tunez.Music.Album, :update_album, :update
  >  destroy Tunez.Music.Album, :destroy_album, :destroy
  end
end
```

This closely resembles the web UI. We don't have a query to list all albums, but we do have mutations like `createAlbum` and `updateAlbum` to manage individual album records. Mutations for existing records have their arguments split into `id`, for the ID of the artist/album to be updated, and `input` for the data to update the record with.

Because we're not in the web UI, though, we don't have the nice pre-filled hidden artist ID when submitting a HTTP request to create an album — we need to provide a valid one with the attributes of the album to be created, like this:

```
mutation {
  createAlbum(input: { name: "New Album Name",
                      artistId: [an-artist-uuid],
                      yearReleased: 2022
                    }) {
    result { id }
  }
}
```

There's just one part we're missing — listing an artist's albums.

Showing albums for a given artist

If you've followed the JSON API section of this chapter, you may have already made the changes necessary to get this working.

By adding two resources with an existing relationship to our API, the flexible nature of GraphQL means that we'll automatically be able to load related records — *as long as the relationship is public*. A relationship is just another field to GraphQL, we don't need to do anything fancy to support them other than making the relationship public. We can do this by adding the option `public? true` to the relationship, in the `Tunez.Music.Artist` resource:

```
04/lib/tunez/music/artist.ex
relationships do
  has_many :albums, Tunez.Music.Album do
    public? true
  end
end
```

This will add the `albums` field to the `Artist` type in the GraphQL API, letting you load related albums anywhere an artist is loaded. Super nifty!

Note that privacy settings on relationships are one way — to be able to load a related artist for an album in the API, you would also need to make the artist relationship in the `Tunez.Music.Album` resource public.

Customizing the generated API

With introspection in the GraphiQL playground, we now have a great overview of our API and we can see it the way a user might. We can also see that there are a few places it can be improved upon — a lot of it was covered in [Customizing the generated API, on page 101](#) for the JSON API, but it applies equally here for GraphQL.

Adding informative descriptions

While it's amazing that we can get all the functionality in our API basically for free, some of the defaults can be a bit lacking. While query and mutation names are usually pretty self-explanatory, sometimes we want to add some more details, or be explicit about what a query argument is when searching artists. Ash allows us to add description metadata to our app in a few different places, that will be picked up by the GraphQL schema generator and added to the documentation. This includes:

- A description for a resource as a whole. This can be added as part of a top-level resource block, such as this in `Tunez.Music.Artist`:

```
defmodule Tunez.Music.Artist do
```

```
use Ash.Resource, ...

resource do
  description "A person or group of people that makes and releases music."
end
```

- A description for an action, or argument for an action, in a resource. These can be added in the action declaration itself, such as:

```
read :search do
  description "List Artists, optionally filtering by name."
  argument :query, :ci_string do
    description "Return only artists with names including the given value."
    # ...
```

As a nice bonus, these descriptions should be picked up by any Elixir-related language server packages in your text editor, such as ElixirLS or elixir-tools in VSCode. And if you extend your API in future, such as adding a REST JSON API, the same descriptions will be used to improve both.

Removing unwanted extras

Looking through the schema carefully, shows that our API can actually do a little bit more than we thought. Expanding the section for the searchArtists query shows that it also accepts arguments named filter, for filtering data, and sort for sorting data. The sort option will let us sort on any public attribute, either ascending or descending. And the filter will let us write complex conditions using greaterThan, lessThanOrEqualTo, ilike, notEq comparisons and so on, and then combine them with and, or and not clauses. And AshGraphql will generate this for *any* list action in our API, for free.

It's customizable, too. Our searchArtists query already has a query argument to filter by name, so we probably don't want to be able to filter on that in the filter too. Some fields also don't make too much sense to filter on, like biography. To customize the list of fields that we can filter on, we can add the filterable_fields config option³⁰ to our graphql block, in the Tunez.Music.Artist resource:

```
04/lib/tunez/music/artist.ex
graphql do
  type :artist
  > filterable_fields [:album_count, :cover_image_url, :inserted_at,
  >   :latest_album_year_released, :updated_at]
end
```

We've removed the values that don't make sense, but we'd still allow users to search by name and also apply filters like:

30. https://hexdocs.pm/ash_graphql/dsl-ashgraphql-resource.html#graphql-filterable-fields

- artists that don't have a cover image URL (`{ coverImageUrl: { isNil: true } }`)
- artists that haven't released an album since 2010 (`{ latestAlbumYearReleased: { lessThan: 2010 } }`)
- artists that were added to Tunez in the last week (`{ insertedAt: { greaterThan: "2024-07-14T17:03:00Z" } }`) (insert your own timestamp)

That's pretty powerful!

If you wanted to take the nuclear option and disable the automatic filtering or sorting, you can also do that with the `derive_filter?` and `derive_sort?` options in your resource - set them to false.

And that's the our GraphQL API, fully done up to match the functionality of our web UI. It's pretty powerful, given how little code we needed to write to support it. We didn't have to define our own GraphQL resolvers, or types, or worry about error handling — everything is handled by AshGraphql. Awesome!

We'll be revisiting our two APIs over the rest of this book, as we add more functionality to Tunez — we want to keep full feature parity with the web, and we also want to see if growing the API organically over time will be difficult to do. In the meantime, we'll look at something a bit different.

Some bad actors have started polluting Tunez with bad data, oh no! This won't do! Tunez has to be the best, most accurate source of high-quality information — and that means locking down certain types of access to only people that we trust to not do anything dodgy. Before we can start limiting access, though, we need to know who people are, and that means some kind of authentication process. Onward march!

Authentication: Who Are You?

In chapter 4, we expanded Tunez with APIs — we now have HTML in the browser, REST JSON, and GraphQL. It was fun seeing how Ash's declarative nature could be used to generate everything for us, using the existing domains, resources and actions in our app.

But now it's time to get down to serious business. The world is a scary place, and unfortunately we can't trust everyone in it to have free reign over the data in Tunez. We need to start locking down access to critical functionality, to only trusted users — but we don't yet have any way of knowing who those users *are*.

We can solve this by adding authentication to our app, and requiring users to log in before they can create or modify any data. Ash has a library that can help with this, called...

Introducing AshAuthentication

It's a very imaginative package name, I know.

There are two parts to AshAuthentication — the core `ash_authentication` package, and the `ash_authentication_phoenix` Phoenix extension to provide things like signup and registration forms. We'll start with the basic library, to get a feel for how it works, and then add the web layer afterwards.

This chapter will be a little different than everything we've covered so far, because we really won't have to write much code until the later stages. The AshAuthentication installer will generate most of the necessary code into our app for us, and while we won't have to modify a lot of it, it's important to understand it. (And it's there if we *do* need to modify it!)

You can install AshAuthentication with Igniter:

```
$ mix igniter.install ash_authentication
```

This will generate a *lot* of code in several stages — so let's break it down bit by bit.

After confirming that we really do want to install the package, our first new interesting thing is that you'll also be prompted to add the simple_sat dependency (if you're using Windows) or the picosat_elixir dependency (if you're using any other operating system). These are both SAT solvers¹ — Ash's policy authorizer uses a SAT solver to evaluate authorization policies. We won't be writing any policies ourselves until [Chapter 6, Authorization: What Can You Do?, on page 133](#), but AshAuthentication provides some basic protections for its own resources out of the box and so we need a SAT solver installed.

Once a SAT solver is installed, and Ash sets up the familiar formatting options, we get into the real meat of the code that AshAuthentication generates.

New domain, who's this?

We're now working with a whole different section of our domain model. Previously we were building music-related resources, so we created a domain named Tunez.Music. Authentication is part of a separate system, an account management system, and so the generator will create a new domain called Tunez.Accounts. This domain will be populated with two new resources — Tunez.Accounts.User and Tunez.Accounts.Token.

The Tunez.Accounts.User resource, in lib/tunez/accounts/user.ex, is what will represent, well, *users* of your app. It comes pre-configured with AshPostgres as its data layer, so each user record will be stored in a row of the users database table.

By itself, the user resource doesn't do much yet. It doesn't even have any attributes, except an id. It does have some authentication-related configuration, in the top-level authentication block — linking the resource with *tokens*, and this is what makes up most of the rest of the generated code.

Tokens and secrets and config, oh my!

Tokens, via the Tunez.Accounts.Token resource and the surrounding config, are the secret sauce to a basic AshAuthentication installation. Tokens are how we securely identify users — from an authentication token provided on every request (“I am logged in as rebecca”), to password reset tokens appended to links in emails, and more.

1. <https://codingnest.com/modern-sat-solvers-fast-neat-underused-part-1-of-n/>

This is the part you *really* don't want to get wrong when building a web app, because the consequences could be pretty bad. So AshAuthentication generates all of the token-related code we need right up front, before we even do anything else. For basic uses, we shouldn't need to touch anything in the generated token code, but it's there if we need to.

So how do we actually use all this code? We need to set up at least one authentication *strategy*.

Setting Up Password Authentication

AshAuthentication supports a number of authentication strategies² — ways we can identify users in our app. Traditionally, we think of logging in to an app via entering an email address and password, which is one of the supported strategies (the password strategy), but there are several more. We can authenticate via different types of OAuth, or even via magic links sent to a user's email address.

Let's set the password strategy up and get a feel for how it works. AshAuthentication comes with igniters to add strategies to our existing app, so you can run the following command:

```
$ mix ash_authentication.add_strategy password
```

This will add a lot *more* code to our app. We now have:

- Two new attributes for the Tunez.Accounts.User resource — email and hashed_password. The email attribute is also marked as an identity, so it must be unique.
- A strategies block added to the authentication configuration in the Tunez.Accounts.User resource. This states that we want to use the email attribute as the identity field for this strategy, and it also sets up the resettable option to allow users to reset their passwords.
- An add_ons block also added to the authentication configuration in the Tunez.Accounts.User resource. This will require users to confirm their email addresses by clicking on links in emails, when registering for an account or changing their email address. Super handy!
- A whole set of actions in our Tunez.Accounts.User resource, around signing in, registering, and resetting passwords.
- And lastly, some template email modules that will be used when it comes to actually sending password reset/email confirmation emails.

2. https://hexdocs.pm/ash_authentication/get-started.html#choose-your-strategies-and-add-ons

That's a lot of goodies!

Because the tasks have created a few new migrations, run `ash.migrate` to get our database up to date:

```
$ mix ash.migrate
```

There will be a few warnings from the email modules about the routes for password reset/email confirmation not existing yet — that's okay, we haven't looked at setting up AshAuthenticationPhoenix yet! But we can still test out our code with the new password strategy in an `iex` session, to see how it works.

Don't try this in a real app!



Note that we'll skip AshAuthentication's built-in authorization policies for this testing, by passing the `authorize?: false` option to `Ash.create`. This is only for testing purposes — the real code in our app won't do this.

Testing authentication actions in iex

One of the generated actions in the `Tunez.Accounts.User` resource is a `register_with_password` create action, which takes `email`, `password`, and `password_confirmation` arguments and creates a user record in the database. It doesn't have a code interface defined, but you can still run it by generating a changeset for the action, and submitting it. Try using your own email address, and the password `supersecret`:

```
iex(1)> Tunez.Accounts.User
Tunez.Accounts.User
iex(2)> |> Ash.Changeset.for_create(:register_with_password, %{email: <<email>>,
... (2)> password: "supersecret", password_confirmation: "supersecret"})
#Ash.Changeset<
  domain: Tunez.Accounts,
  action_type: :create,
  action: :register_with_password,
  ...
>
iex(3)> |> Ash.create!(authorize?: false)
[debug] QUERY OK source="users" db=2.4ms
INSERT INTO "users" ("id", "email", "hashed_password") VALUES ($1,$2,$3)
RETURNING "hashed_password", "email", "id", "confirmed_at" [<<uuid>>,
#Ash.CiString<<email>>, <<hashed password>>]
[debug] QUERY OK source="tokens" db=2.5ms
INSERT INTO "tokens" AS t0 ("id", "created_at", "expires_at", "extra_data",
"inserted_at", "purpose", "subject", "updated_at", "jti") VALUES ($1,$2,$3,
$4,$5,$6,$7,$8,$9) ON CONFLICT ("id", "jti") DO UPDATE SET "expires_at"
= EXCLUDED."expires_at", "extra_data" = EXCLUDED."extra_data",
```

```

"purpose" = EXCLUDED."purpose", "subject" = EXCLUDED."subject",
"updated_at" = COALESCE(EXCLUDED."updated_at", $10) RETURNING
"updated_at","inserted_at","extra_data","purpose","expires_at","subject",
"jti","id","created_at" [«uuid», «now», «three days from now», %{email:
«email»}, «now», "confirm_new_user", "user?id=«uuid»", «now»,
«string», «now»]

```

Click this link to confirm your email:

```
http://localhost:4000/auth/user/confirm_new_user?confirm=«token»
```

```

#Tunez.Accounts.User<
  __meta__: #Ecto.Schema.Metadata<:loaded, "users">,
  confirmed_at: nil,
  id: «uuid»,
  email: #Ash.CiString<«email»>,
  ...
>

```

Note that there are no code interfaces defined for the actions in the Tunez.Accounts.User resource, so you'll need to construct the changesets manually.

Calling this action has done a few things:

- Inserted the new user record into the database, including securely hashing the provided password;
- Created a new confirm_new_user token record, for the new user to use when confirming their email address, and
- Generated an “email” to send to the user to actually confirm their email address (the “Click this link to confirm your email” text in the output above). Note that it hasn't actually *sent* an email — that requires a bit of extra mailer-specific configuration — but the basis is there.

What can we do with our new user record? We can try to authenticate them, using the created sign_in_with_password action. This mimics what a user would do on a login form, by entering their email address and password:

```

iex(9)> Tunez.Accounts.User
Tunez.Accounts.User
iex(10)> |> Ash.Query.for_read(:sign_in_with_password, %{email: «email»,
... (10)> password: "supersecret"})
#Ash.Query<
  resource: Tunez.Accounts.User,
  arguments: %{password: "***redacted***", email: #Ash.CiString<«email»>},
  filter: #Ash.Filter<email == #Ash.CiString<«email»>>
>
iex(11)> |> Ash.read(authorize?: false)
[debug] QUERY OK source="users" db=0.7ms idle=1472.4ms
SELECT u0."id", u0."confirmed_at", u0."email", u0."hashed_password" FROM

```

```
"users" AS u0 WHERE (u0."email"::citext = ($1::citext)) [«email»]
{:ok,
 [
  #Tunez.Accounts.User<
  __meta__: #Ecto.Schema.Metadata<:loaded, "users">,
  confirmed_at: nil,
  email: «email»,
  ...
]
```

And it works! AshAuthentication has validated that the credentials are correct, by fetching any user records with the provided email, hashing the provided password, and verifying that it matches what is stored in the database. You can also try it with different credentials like an invalid password; AshAuthentication will properly return an error.

Calling `sign_in_with_password` with correct credentials has also generated an authentication token in the returned user's metadata, to be stored in the browser and used to authenticate the user in future.

```
iex(12)> {:ok, [user]} = v()
{:ok, [#Tunez.Accounts.User<...>]}
iex(13)> user.__metadata__.token
"eyJhbGciOi..."
```

This token is a JSON Web Token, or JWT.³ It's cryptographically signed by our app to prevent tampering — if a malicious user has a token and edits it to attempt to impersonate another user, the token will no longer verify. To test out the verification, we can use some of the built-in AshAuthentication functions like `AshAuthentication.Jwt.verify/2` and `AshAuthentication.subject_to_user/2`:

```
iex(14)> AshAuthentication.Jwt.verify(user.__metadata__.token, :tunetz)
{:ok,
 %{
  "aud" => "~> 4.2",
  "exp" => 1731813618,
  "iat" => 1730604018,
  "iss" => "AshAuthentication v4.2.7",
  "jti" => <<string>>,
  "nbf" => 1730604018,
  "purpose" => "user",
  "sub" => "user?id=<<uuid>>"}
}, Tunez.Accounts.User}
```

The interesting parts of the decoded token here are the `sub` (subject) and the `purpose`. JWTs can be created for all kinds of purposes, and this one is for user

3. <https://jwt.io/>

authentication, hence the purpose “user”. The subject is a specially-formatted string with a user ID in it, which we can verify belongs to a real user:

```
iex(15)> {:ok, claims, resource} = v()
{:ok, %{}, Tunez.Accounts.User}
iex(16)> AshAuthentication.subject_to_user(claims["sub"], resource)
[debug] QUERY OK source="users" db=1.6ms queue=0.8ms idle=1848.6ms
SELECT u0."id", u0."confirmed_at", u0."hashed_password", u0."email" FROM
"users" AS u0 WHERE (u0."id":::uuid = $1::uuid) [«uuid»]
{:ok,
#Tunez.Accounts.User<
  email: #Ash.CiString<«your email»>,
...
>}
```

We don’t really need to muck around with all of this, though. It’s good to know how AshAuthentication works, and how to verify that it works, but we’re building a web app — we want forms that users can fill out to register or sign-in. For that, we can look at AshAuthentication’s sister library, AshAuthenticationPhoenix.

Automatic UIs With AshAuthenticationPhoenix

As the name suggests, AshAuthenticationPhoenix is a library that connects AshAuthentication with Phoenix, providing a great LiveView-powered UI that we can tweak a little bit to fit our site look and feel, but otherwise don’t need to touch.

Like other libraries, you can install it with Igniter:

```
$ mix igniter.install ash_authentication_phoenix
```

Ignoring the same warnings about some routes not existing (this will be the last time we see them!), the AshAuthenticationPhoenix installer will set up:

- A basic Igniter config file in `.igniter.exs` — this is the first generator we’ve run that needs specific configuration (for `Igniter.Extensions.Phoenix`), so it gets written to a file
- A `TunezWeb.AuthOverrides` module, that we can use to customize the look and feel of the generated liveviews a bit (in `lib/tunez_web/auth_overrides.ex`)
- A `TunezWeb.AuthController` module, to securely process sign-in requests (in `lib/tunez_web/controllers/auth_controller.ex`). This is due to a bit of a quirk in how LiveView works; it doesn’t have access to the user session to store data on successful authentication.

- A `TunezWeb.LiveUserAuth` module providing a set of hooks we can use in live-views to require a certain authentication status (in `lib/tunez_web/live_user_auth.ex`)
- And lastly, the most important part — updating our web app router in `lib/tunez_web/router.ex` to add plugs and routes for all of our authentication-related functionality.

Before we can test it out, there's one manual change we need to make, as Igniter doesn't (yet) know how to patch JavaScript files. Because AshAuthenticationPhoenix's liveviews are styled with Tailwind CSS, we need to add its liveview paths to Tailwind's content lookup paths. Add the `"..deps/ash_authentication_phoenix/**/*.*ex"` line to the content array in `assets/tailwind.config.js`, restart your mix phx.server, and then we can see what kind of UI we get, by visiting the sign-in page at <http://localhost:4000/sign-in>:

Ash Framework

Email

Password

[Forgot your password?](#)

[Need an account?](#)

Sign in

It's pretty good! Out of the box we can sign-in, register for new accounts, and request password resets, without lifting a finger.

After signing in, we get redirected back to the Tunez homepage — but there's no indication that we're now logged in, and no link to log out. We'll fix that now.

Showing the currently-authenticated user

Most web apps show their current user indicator in the top-right corner of the page, so that's what we'll implement as well. The main Tunez navigation is part of the application layout, in `lib/tunez_web/components/layouts/app.html.heex`, so we can edit to add a new rendered `user_info` component:

```
05/lib/tunez_web/components/layouts/app.html.heex
<div class="flex flex-col w-full max-w-6xl m-auto">
  <div class="w-full navbar px-4 pb-0">
    <% # ... %>
  >  <.user_info current_user={@current_user} />
  </div>
  <div class="divider divider-primary m-0"></div>
  <% # ... %>
</div>
```

This is an existing function component located in the TunezWeb.Layouts module, in lib/tunez_web/components/layouts.ex, and shows sign in/register buttons if there's no user logged in, and a dropdown of user-related things if there is. Refreshing the app after making this change shows a big error, however:

```
key :current_user not found in: %{
  socket: #Phoenix.LiveView.Socket<...>,
  sort_by: "-updated_at",
  __changed__: %{}...},
  page_title: "Artists",
  inner_content: %Phoenix.LiveView.Rendered{...},
  ...
}
```

What gives? Fixing this will require digging a bit into how the new router code works, so let's take a look.

Digging into AshAuthenticationPhoenix's generated router code

We didn't really go over the changes to our router in lib/tunez_web/router.ex, after installing AshAuthenticationPhoenix — we just assumed everything was all good. For the most part it is, but there are one or two things we need to tweak.

The igniter added plugs to our pipelines to load the current user — `load_from_bearer` for our API pipelines, and `load_from_session` for our browser pipeline. This works for traditional “deadview” web requests, that receive a request and send the response in the same process.

LiveView works differently, though. When a new request is made to a liveview, it spawns a new process and keeps that active websocket connection open for that realtime data transfer. This new process doesn't have access to the session, so although our base request knows who the user is, the spawned process does not.

Enter `live_session`, and how its wrapped by AshAuthentication, `ash_authentication_live_session`. This macro will ensure that when new processes are spawned, they get copies of the data in the session so the app will continue working as expected.

What does this mean for Tunez? It means that all our liveview routes that are expected to have access to the current user, need to be moved into the `ash_authentication_live_session` block in the router.

```
05/lib/tunez_web/router.ex
scope "/", TunezWeb do
  pipe_through :browser

  > # This is the block of routes to move
  > live "/", Artists.IndexLive
  > #
  > live "/albums/:id/edit", Albums.FormLive, :edit

  auth_routes AuthController, Tunez.Accounts.User, path: "/auth"
  sign_out_route AuthController
  ...

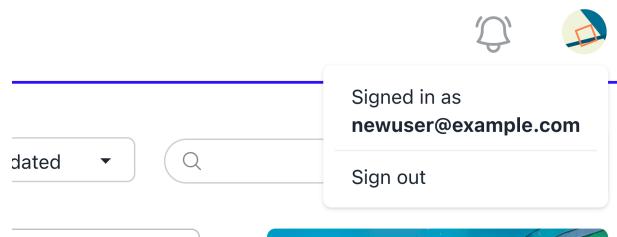
```

The `ash_authentication_live_session` helper is in a separate scope block in the router, earlier on in the file:

```
05/lib/tunez_web/router.ex
scope "/", TunezWeb do
  pipe_through :browser

  ash_authentication_live_session :authenticated_routes do
    > # This is the location that the block of routes should be moved to
    > live "/", Artists.IndexLive
    > #
    > live "/albums/:id/edit", Albums.FormLive, :edit
  end
end
```

With this change, our app should be renderable, and we should see information about the currently logged in user in the top-right corner of the main navigation.



Now we can turn our attention to the generated liveviews themselves. We want them to look totally seamless in our app, like we wrote and styled them ourselves. While we don't have control over the HTML that gets generated, we can customize a lot of the styling and some of the content, using *overrides*.

Stylin' and profilin' with overrides

Each liveview component in AshAuthenticationPhoenix's generated views has a set of overrides configured, that we can use to change things like component class names and image URLs.

When we installed AshAuthenticationPhoenix, a base `TunezWeb.AuthOverrides` module was created in `lib/tunez_web/auth_overrides.ex`. This shows the syntax that we can use to set different attributes, that will then be used when the liveview is rendered:

```
05/lib/tunez_web/auth_overrides.ex
override AshAuthenticationPhoenix.Components.Banner do
  set :image_url, "https://media.giphy.com/media/g7GKcSzWQfugw/giphy.gif"
  set :text_class, "bg-red-500"
end
```

As well a link to the complete list of overrides⁴ you can use, in the documentation.

Let's test it out, by changing the colour of the "Sign In" button on the sign in page. Buttons in Tunez are purple, not blue! It can be a bit tricky to find exactly which override will do what you want, but in this case, the submit button is an `input`, and under `AshAuthenticationPhoenix.Components.Password.Input` is an override for `submit_class`. Perfect.

In the overrides file, set a new override for that `Input` component:

```
05/lib/tunez_web/auth_overrides.ex
defmodule TunezWeb.AuthOverrides do
  use AshAuthenticationPhoenix.Overrides

  > override AshAuthenticationPhoenix.Components.Password.Input do
  >   set :submit_class, "phx-submit-loading:opacity-75 btn btn-primary"
  > end
```

As any overrides we set will completely override the default styles, there may be more of a change than you expect. If you're curious about what the default values for each override are, or you want to copy and paste them so you can only change what you need, you can see them in the AshAuthenticationPhoenix source code.⁵

4. https://hexdocs.pm/ash_authentication_phoenix/ui-overrides.html#reference

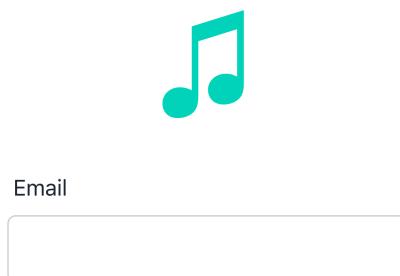
5. https://github.com/team-alembic/ash_authentication_phoenix/blob/main/lib/ash_authentication_phoenix/overrides/default.ex

The generated classes include styles for dark mode using the Tailwind dark-class modifier,⁶ so if you're building an app that specifies darkMode: "media" (or doesn't specify darkMode at all) as part of its Tailwind config in assets/tailwind.config.js, you'll see light or dark theming depending on your system theme. For ease of use, Tunez has disabled automatic dark mode by setting darkMode: "selector", but it's good to be aware of.

We won't bore you with every single class change to make, to turn a default AshAuthenticationPhoenix form into one matching the rest of the site theme, so we've provided a set of overrides in the starter app in lib/tunez_web/auth_overrides_sample.txt. You can take the contents of that file and replace the contents of the TunezWeb.AuthOverrides module, like so:

```
05/lib/tunez_web/auth_overrides.ex
defmodule TunezWeb.AuthOverrides do
  use AshAuthentication.Phoenix.Overrides
  alias AshAuthentication.Phoenix.Components
  ...
  override Components.Banner do
    set :image_url, nil
  end
  ...
end
```

And it should look like this:



Feel free to tweak the styles the way you like - Tunez is your app, after all!

Why do users always forget their passwords!?

For the last part of the password strategy implementation, we'll look at filling in the gaps on letting users reset their passwords. This means we'll be generating and sending emails!

We briefly saw email-related things earlier when we tested registering for an account in iex - text like this was printed to the console:

Click this link to confirm your email:

6. <https://tailwindcss.com/docs/dark-mode>

```
http://localhost:4000/auth/user/confirm_new_user?confirm=<<token>>
```

This was for email address confirmation, not password resetting, but they’re almost identical in implementation so you can implement both at the same time.

The modules responsible for generating the emails, *senders* in AshAuthentication jargon, live in `lib/tunez/accounts/user/senders`. There’s one for `SendNewUserConfirmationEmail`, and one for `SendPasswordResetEmail`. Currently, they don’t actually *send* any emails — they only print messages out to the console like we saw earlier. Very deceptive.

If we connect these modules up to mailers, they’ll be able to send emails like a boss. Phoenix apps come with a Swoosh⁷ integration built in for sending email, so we can leverage that. We’ve provided a basic mailer with the Tunez app, in `lib/tunez/emails.ex` — it defines two public functions named `deliver_email_confirmation_email/2` and `deliver_password_reset_email/2` (and one other that’s a bit of a spoiler for what’s coming next) that do exactly what they say on the tin, generate HTML emails and send them with Swoosh.

To connect the email confirmation sender to the mailer, you can edit the `send/2` function in `Tunez.Accounts.User.Senders.SendNewUserConfirmationEmail` to call the `deliver_email_confirmation_email/2` function:

```
05/lib/tunez/accounts/user/senders/send_new_user_confirmation_email.ex
defmodule Tunez.Accounts.User.Senders.SendNewUserConfirmationEmail do
  # ...
  def send(user, token, _) do
    >   Tunez.Emails.deliver_email_confirmation_email(
    >     user,
    >     url(~p"/auth/user/confirm_new_user?#[{confirm: token}]"))
    >
  end
end
```

The URL was already part of the module; we’ve only needed to move it to pass it to the mailer. You can connect the `SendPasswordResetEmail` module up in a very similar way.

We don’t need to set up an email provider to send real emails, while working in development. Swoosh provides a “mailbox” we can use — any emails sent, no matter the target email address, will be delivered to the dev mailbox (instead of actually being sent!). This dev mailbox is added to our router in dev mode only, and can be accessed at <http://localhost:4000/dev/mailbox>.

7. <https://hexdocs.pm/swoosh/>

The mailbox is empty by default, but if you register for a new account and then refresh the mailbox:

Mailbox	From	"Tunez" <tunez@sevenseacat.net>
1 message(s)	To	newuser@example.com
Tunez	Subject	Confirm your email address
Confirm your email address	Sent at	05/11/2024, 5:17:25 pm

The email contains a link to confirm the email address, which, sure, that's totally my email address and I did sign up for the account, so click the link.

You'll be redirected back to the app homepage, with a flash message letting us know that our email address is confirmed. It works!

Setting Up Magic Link Authentication

Word going around nowadays is that some users think that passwords are just *so passé*, and they'd much prefer to be able to log in using magic links instead — enter their email address, click the login link that gets sent straight to their inbox, and they're in. That's no problem!

AshAuthentication doesn't limit our apps to just *one* method of authentication, we can add as many as we like from the supported strategies⁸ or even write our own. So there's no problem with adding the magic link strategy to our existing password-strategy-using app, and users can even log in with either strategy depending on their mood. Let's go.

To add the strategy, we can run the `ash_authentication.add_strategy` Mix task:

```
$ mix ash_authentication.add_strategy magic_link
```

This will add:

- A new `magic_link` authentication strategy block to our `Tunez.Accounts.User` resource, in `lib/tunez/accounts/user.ex`
- Two new actions named `sign_in_with_magic_link` and `request_magic_link`, also in our `Tunez.Accounts.User` resource,
- And finally, a new sender module responsible for generating the magic link email, in `lib/tunez/accounts/user/senders/send_magic_link_email.ex`.

The `magic_link` config block in the `Tunez.Accounts.User` resource lists some sensible values for the strategy configuration, such as the name of the identity attribute

8. https://hexdocs.pm/ash_authentication/get-started.html#choose-your-strategies-and-add-ons

(email by default). There are more options⁹ that can be set, such as how long generated magic links are valid for (`token_lifetime`), but we won't need to add anything extra to what is generated here.

Wait... that's it? Yep, that's it. The initial setup of AshAuthentication generates a lot of code for the initial resources, but adding subsequent strategies typically only needs a little bit.

Once you've added the strategy, visiting the sign-in page will have a nice surprise:



Is it really that simple?? Unfortunately not, but we're about 90% of the way there without breaking a sweat.

Connecting the magic link email

Like we did with password reset and email confirmation, we'll need to connect the generated sender module to a mailer. The default behaviour of the sender is to print the email content to the logs, which can't be used by our users viewing the app in their browsers.

In the generated sender `Tunez.Accounts.User.Senders.SendMagicLinkEmail` module, we can replace the `IO.puts` call with a call to the pre-existing `Tunez.Emails.deliver_magic_link_email` function:

```
05/lib/tunez/accounts/user/senders/send_magic_link_email.ex
defmodule Tunez.Accounts.User.Senders.SendMagicLinkEmail do
  # ...

  def send(user_or_email, token, _) do
    # ...

    email =
      case user_or_email do
        %{email: email} -> email
        email -> email
      end

  >   Tunez.Emails.deliver_magic_link_email(
  >     email,
```

9. https://hexdocs.pm/ash_authentication/dsl-ashauthentication-strategy-magiclink.html#options

```
>     url(~p"/auth/user/magic_link/?token=#{token}")  
>   )  
end  
end
```

It's a little bit different from the other mailers — because our magic link form can be used to either create new accounts or log in to existing accounts, we may have a user record or just an email address to send the email to. And again, the URL was already present in this sender, we only need to pass it on to the mailer.

Now if we fill out the magic link form with the same email address we confirmed earlier, an email will be delivered to our dev mailbox, with a sign-in link to click. One step closer! Click the link, and... we get an error message, a generic “incorrect email or password” that doesn’t give away any secrets. How can we tell what’s happening behind the scenes?

Debugging when authentication goes wrong

While showing a generic failure message is good in production for security reasons (we don’t want to give away information like if an email address definitely belonged to an account, to potentially-bad actors), it’s not good in development while you’re trying to make things work, or debugging issues.

To get more information about what’s going on, we can enable authentication debugging for our development environment *only*, by placing the following at the bottom of config/dev.exs:

```
05/config/dev.exs  
config :ash_authentication, debug_authentication_failures?: true
```

Restart your mix phx.server to apply the new config change, and then send another magic link to the same email address. Click the magic link, and you should see a big yellow warning in your server logs:

```
[warning] Authentication failed:  
Bread Crumbs:  
  > Error returned from: Tunez.Accounts.User.sign_in_with_magic_link  
Invalid Error  
* attribute hashed_password is required  
(ash 3.x.xx) lib/ash/error/changes/required.ex:4: Ash.Error.Changes...  
(ash 3.x.xx) lib/ash/changeset/changeset.ex:3337: Ash.Changeset...  
...
```

Other possible errors that may occur here

This may not be the exact error you see here, there are two other options:

- Invalid magic_link token: this is what you will see if you've clicked on an expired magic link. Magic link is only valid for 10 minutes by default, so if you've clicked on one after that time it will fail and you'll see this error in your logs.
- Attempted to update stale record of Tunez.Accounts.User with filter: not is_nil(confirmed_at): for security reasons, you can't use a magic link to an existing account that has an unconfirmed email address, as this would allow account hijacking.¹⁰ This will show a different flash message in the browser, instead of the generic "incorrect email or password" — you'll be told that the email address already belongs to an account.



If you try with a different email address, either with an email belonging to a confirmed account or not belonging to an existing account, this error will go away!

That makes sense conceptually — when we added the password strategy to our Tunez.Accounts.User resource, AshAuthentication marked the generated email and hashed_password attributes as required fields that could not be null. Now that we've added magic link authentication, we still need email to identify our users, but they might not have a password.

To fix this issue, we need to update the Tunez.Accounts.User resource to mark the hashed_password attribute as allow_nil? true:

```
05/lib/tunez/accounts/user.ex
defmodule Tunez.Accounts.User do
  # ...
  attributes do
    # ...
    attribute :hashed_password, :string do
      >     allow_nil? true
      sensitive? true
    end
  end
end
```

10. https://hexdocs.pm/ash_authentication/confirmation.html#important-security-notes

As this will also change the field in the database, we need to generate a migration for it and run it:

```
$ mix ashcodegen allow_nil_user_passwords
$ mix ash.migrate
```

Now when we click the magic link in the email again (as long as it's not expired!) you should be successfully logged in to Tunez, all without needing a password!

Without turning the AshAuthentication debugging on, these kinds of issues would be near-impossible to fix. It's safe to leave it enabled in development, as long as you don't mind the warning about it during server start. If the warning is too annoying, feel free to turn debugging off, but don't forget that it's available to you!

And that's all we need to do, to implement magic link authentication in our apps. Users will be able to create accounts via magic links, and also log into their existing accounts that were created with an email and password. Our future users will thank us!

Can we allow authentication over our APIs?

In the previous chapter, we built two shiny APIs that users can use to programmatically access Tunez and its data. In order to make sure the APIs have full feature parity with the web UI, we need to make sure they can register and sign in via the API as well. When we start locking down access to critical parts of the app, we don't want API users to be left out!

Let's give it a try, and see how far we can get. We'll start with adding registration support, in our JSON API.

To add JSON API support to our `Tunez.Accounts.User` resource, we can extend it using Ash's extend patcher:

```
$ mix ash.patch.extend Tunez.Accounts.User json_api
```

This will configure our JSON API router, domain module, and resource with everything we need to start connecting routes to actions. To create a POST request to our `register_with_password` action, we can add a new route to the domain,¹¹ like we did with [Adding Albums to the API, on page 96](#). We've customized the actual URL with the `route` option, to create a full URL like `/api/json/users/register`.

11. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-domain.html#json_api_routes-base_route-post

```
05/lib/tunez/accounts.ex
defmodule Tunez.Accounts do
  use Ash.Domain, extensions: [AshJsonApi.Domain]

  ▶  json_api do
  ▶    routes do
  ▶      base_route "/users", Tunez.Accounts.User do
  ▶        post :register_with_password, route: "/register"
  ▶      end
  ▶    end
  ▶  end

  # ...
end
```

Looks good so far! But if you try it in an API client, or using cURL, correctly supplying all the arguments that the action expects... it won't work, it always returns a forbidden error. Drat.

This is because at the moment, the `Tunez.Accounts.User` resource is really tightly secured. All of the actions are restricted, to only be accessible via `AshAuthenticationPhoenix`'s form components. (Or if we skip authorization checks, like we [did earlier on page 116](#). That was for test purposes only!)

This is good for security reasons — we don't want *any* old code to be able to do things like change people's passwords! It makes our development lives a little bit harder, though, because to understand how to allow the functionality we want, we need to dive into our next topic, *authorization*. Buckle up, this may be a long one...

Authorization: What Can You Do?

We left Tunez in a good place at the end of the previous chapter. Visitors to the app can now register accounts, or login with either an email and password or a magic link. Now we can identify *who* is using Tunez.

However, we couldn't allow users to authenticate or register an account for Tunez via either of our APIs. The app also doesn't behave any differently depending on whether a user is logged in or not. Anyone can still create, edit and delete data, without needing to be logged in. This is what we want to prevent, for better data integrity — unauthenticated users should have a total read-only view of the app, and authenticated users should be able to perform only the actions they are granted access to. We can enforce this by implementing *access control* in the app, using a component of Ash called *policies*.

Introducing Policies

Policies define who has access to resources within our app, and what actions they can run. Each resource can have its own set of policies, and each policy can apply to one or more actions defined in that resource.

Policies are checked internally by Ash *before* any action is run — if *all* policies that apply to a given action *pass* (return authorized), then the action is run. If one or more of the policies *fail* (return unauthorized), then the action is *not* run and an error is returned.

Because policies as part of resource definitions, they will automatically be checked on all calls to actions in those resources. You only have to write them once, and they will apply everywhere — in our web UI, our REST and GraphQL APIs, an iex REPL, and any other interfaces we add in future. You also don't have to worry about the *how* of policy checking, the *when*, you're freed up to

focus on the actual business logic of who can access what. This makes access control simple, straightforward, and fast to implement.

A lot of policy checks will naturally depend on the entity calling the action, or the *actor*. This is usually (but not always!) the person using the app, clicking buttons and links in their browser. This is why it's a pre-requisite to know who our users are!

At its core, a policy is made up of two things:

- One or more *policy conditions*, to determine whether or not the policy applies to a specific action request
- A set of *policy checks*, that are run to see if a policy passes or fails. Each policy check is itself made up of a check condition, and an action to take if the check condition matches the action condition.

There's a lot of conditions to consider, so an example will hopefully make it clearer. If we were building a Blog application, and wanted to add policies for a Post in our blog, a policy might look like the following:

```
defmodule Blog.Post do
  use Ash.Resource, authorizers: [Ash.Policy.Authorizer]

  policies do
    policy action(:publish) do
      forbid_if expr(published == true)
      authorize_if actor_attribute_equals(:role, "admin")
    end
  end
end
```

The single condition for this policy is `action(:publish)`, meaning the policy will apply to any calls to the hypothetical `publish` action in this `Blog.Post` resource. It has two policy checks — one that will *forbid* the action *if* the `published` attribute on the resource is true, and one that will *authorize* the action *if* the actor calling the action has the attribute `role` set to “admin”.

Or in human terms, *admin* users can *publish* a blog post if it's not already *published*. Any other cases would return an error.

As mentioned previously in [Loading Related Resource Data, on page 43](#), all actions (including code interfaces) support an extra argument of options to customize the behaviour of the action. One of the options that all actions support is `actor` (as seen in the list of options for each action type),¹ so we can modify each action call to add the `actor` option:

1. <https://hexdocs.pm/ash/Ash.html>

```
# Instead of this...
Blog.Post.publish(post)

# We can use this!
Blog.Post.publish(post, actor: current_user)
```

Decisions, decisions

When evaluating policy checks, the first check that successfully *makes a decision* determines the overall result of the policy. Policies could be thought of like a `cond` statement — checks are evaluated in order until one makes a decision, and the rest are ignored — so the order of checks within a policy is really important, perhaps more than it first appears.

In our `Blog.Post` example, if we attempt to publish a post that has the attribute `published` equal to true, the first check will make a decision because the check condition is true (the attribute `published` is indeed equal to true), and the action is to `forbid_if` the check condition is true.

The second check for the admin role is irrelevant — a decision has already been made, so no-one can publish an already-published blog post, period.

If the order of the checks in our policy was reversed:

```
policy action(:publish) do
  authorize_if actor_attribute_equals(:role, "admin")
  forbid_if expr(published == true)
end
```

Then the logic is actually a bit different. Now we only check the `published` attribute if the first check *doesn't* make a decision. To not make a decision, the actor would have to have `role` attribute that *doesn't* equal “admin”. In other words, this policy would mean that admin users can publish *any* blog post. Subtly different, enough to cause unintended behaviour in your app!

Authorizing API Access for Authentication

Now that we have a bit of context around policies and how they can restrict access to actions, we can check out the policies that were automatically generated in the `Tunez.Accounts.User` resource when we installed AshAuthentication:

```
06/lib/tunez/accounts/user.ex
defmodule Tunez.Accounts.User do
  # ...
  policies do
    bypass AshAuthentication.Checks.AshAuthenticationInteraction do
      authorize_if always()
    end
```

```

policy always() do
  forbid_if always()
end
end
# ...
end

```

There are two policies defined in the resource — one *bypass* (more on bypasses [on page 149](#)) and one standard policy. The `AshAuthenticationInteraction` module, used in the the bypass, contains a custom policy condition that applies to any actions called from `AshAuthenticationPhoenix`'s liveviews, and it will *always* authorize them. So the web UI works out of the box, and allows actions like `register_with_password` or `sign_in_with_magic_link` to be run.

The second policy *always* applies to any other action call (the policy condition), and it will *forbid if*, well, *always!* This includes action calls from any of our generated API endpoints, and explains why we always got a forbidden error when we tried to register an account via the API. The `User` resource is open for `AshAuthentication`'s internal action calls, and firmly closed for absolutely everything else. Let's update this and write some new policies, to allow access to the actions we want.

Writing our first user policy

If you test running any actions on the `Tunez.Accounts.User` resource in `iex` (without skipping authorization with `authorize?: false`, like we [did earlier on page 116](#)), you'll immediately see the forbidden errors that were getting processed by `AshJsonApi`, and returned in our API responses.

```

iex(1)> Tunez.Accounts.User
Tunez.Accounts.User
iex(2)> |> Ash.Changeset.for_create(:register_with_password, %{email:
... (2)> <>email>, password: "password", password_confirmation: "password"})
#Ash.Changeset<...>
iex(3)> |> Ash.create()
{:error, %Ash.Error.Forbidden{...}}

```

It wouldn't hurt to have *some* of the actions on the resource, such `register_with_password` and `sign_in_with_password`, accessible over the API. To do that, you can remove the policy `always()` policy — if there are no policies in a `policies` block that apply to an action, it will be forbidden by default. We can replace it with a new policy, one that will *authorize* calls to those actions:

```

06/lib/tunez/accounts/user.ex
policies do
  bypass AshAuthentication.Checks.AshAuthenticationInteraction do
    authorize_if always()

```

```

    end
  ▶  policy action([:register_with_password, :sign_in_with_password]) do
  ▶    authorize_if always()
  ▶  end
end

```

This is the most permissive type of policy check. As it says on the tin, it *always* authorizes any action that meets the policy condition — any action that has one of those two names.

Authorizing a sign-in action does *not* mean the sign-in will be successful!



It's important to note the distinction between an action being *authorized* ("this user is allowed to run this action"), and an action returning a *successful result* ("this user provided valid credentials and is now signed in"). It can get a little murky, especially around authentication-related actions.

In our example, if a user attempts to sign-in with an incorrect password, the action will now be authorized, so the sign-in action will run... and return an authentication failure because of the incorrect password.

With that policy in place, if you re-run the action to register a user in iex, it will now have a very different result:

```

iex(1)> Tunez.Accounts.User
Tunez.Accounts.User
iex(2)> |> Ash.Changeset.for_create(:register_with_password, %{email:
... (2)> <>email>, password: "password", password_confirmation: "password"})
#Ash.Changeset<...>
iex(3)> |> Ash.create()
{:ok, #Tunez.Accounts.User<...>}

```

This is what we want! Is that all we needed to do, to make it accessible over the API?

Authenticating via JSON

When we left the JSON API at the end of the previous chapter, we'd configured a route in our `Tunez.Accounts` domain for the `register_with_password` action of the `User` resource, but it didn't quite work:

```

05/lib/tunez/accounts.ex
defmodule Tunez.Accounts do
  use Ash.Domain, extensions: [AshJsonApi.Domain]

  ▶  json_api do
  ▶    routes do

```

```
>     base_route "/users", Tunez.Accounts.User do
>       post :register_with_password, route: "/register"
>     end
>   end
>
# ...
end
```

Now that we've updated the policies to authorize the action, will accessing this API endpoint work, and successfully create a new user? We can test it with an API client again:

```
POST http://localhost:4000/api/json/users/register
```

Query	Body	Headers	Auth	Vars	Script	Assert
Tests	Docs				JSON	Prettify

```

1 ▶  {
2 ▶    "data": {
3 ▶      "attributes": {
4 ▶        "email": "apiuser@example.com",
5 ▶        "password": "password",
6 ▶        "password_confirmation": "password"
7 ▶      }
8 ▶    }
9 ▶ }
```

```

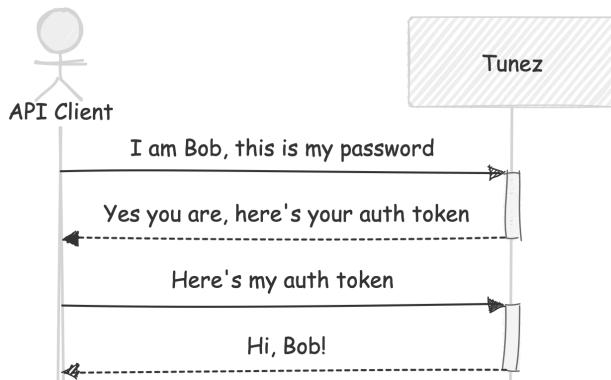
1 ▶  {
2 ▶    "data": {
3 ▶      "attributes": {
4 ▶        "email": "apiuser@example.com"
5 ▶      },
6 ▶      "id": "8642c1e9-d524-427c-80ee-28ce54532b0e",
7 ▶      "links": {},
8 ▶      "meta": {},
9 ▶      "type": "user",
10 ▶      "relationships": {}
11 ▶    },
12 ▶    "links": { },
13 ▶    "meta": {},
14 ▶    "jsonapi": { }
15 ▶ }
```

It does create a user! But the response isn't quite right. When we tested out authentication earlier in [Testing authentication actions in iex, on page 116](#), we saw that the user data included an authentication token as part of its metadata, to be used to authenticate future requests.

```
iex(11)> |> Ash.read(authorize?: false)
{:ok, [#Tunez.Accounts.User<...>]}
iex(12)> {:ok, [user]} = v()
{:ok, [#Tunez.Accounts.User<...>]}
iex(13)> user.__metadata__.token
"eyJhbGciOi..."
```

That was for sign-in, not registration, but the same principle applies — when you register an account, you're typically automatically signed into it.

This token needs to be included as part of the response, so API clients can store it and send it as a request header with future requests to the API.



This is handled for us with `AshAuthenticationPhoenix`, with cookies and sessions, but for an API the clients must handle it themselves.

In `AshJsonApi`, we can attach *metadata* to an API response² — this can include any arbitrary data from anywhere. The user's authentication token sounds like a good fit! The metadata option for a route takes a three-argument function that includes the data being returned in the response (the created user), and returns a map of data to include, so we can extract the token and return it:

```
06/lib/tunex/accounts.ex
post :register_with_password do
  route "/register"
  metadata fn _subject, user, _request ->
    %{token: user.__metadata__.token}
  end
end
```

And now it will work as we expect! The token gets returned, and can now be used to authenticate subsequent responses.

The exact same process can be used for creating the sign-in route — give it a nice URL, and add the token to the response:

```
06/lib/tunex/accounts.ex
base_route "/users", Tunex.Accounts.User do
  # ...
  post :sign_in_with_password do
    route "/sign_in"
    metadata fn _subject, user, _request ->
      %{token: user.__metadata__.token}
    end
  end
end
```

2. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-resource.html#json_api-routes-post-metadata

```
end
```

Other authenticated-related actions, such as magic links or password resetting, don't make as much sense to perform over an API — they all require following links in emails that go to the app to complete. It *can* be done, but it's a much less common use case so we'll leave it as an exercise to the reader!

(Hint: for generic actions, such as `request_password_reset_with_password`, you'll need to use the route DSL³ to create the API route.)

Authenticating via GraphQL

Adding support for authentication to our GraphQL API is a little bit easier than for the JSON API, now that we've granted access to the actions via our policies — some of the abstractions around metadata are already built into the library.

To get started, extend the `Tunez.Accounts.User` resource with `AshGraphql`, using Ash's `extend` patcher:

```
$ mix ash.patch.extend Tunez.Accounts.User graphql
```

This will configure our GraphQL schema, domain module, and resource with everything we need to start creating queries and mutations for actions.

`AshGraphql` is strict in regards to which types of actions can be defined as mutations, and which as queries. Read actions must be queries; creates, updates and destroys must be mutations; and generic actions can be either. So for the two actions we want to make accessible over GraphQL — `create :register_with_password` must be added as a mutation, and `read :sign_in_with_password` must be a query.

To create a mutation for the `register_with_password` action, define a new `graphql` block in the `Tunez.Accounts` domain module and populate it using the `create` macro:⁴

```
06/lib/tunez/accounts.ex
defmodule Tunez.Accounts do
  use Ash.Domain, extensions: [AshGraphql.Domain, AshJsonApi.Domain]
  > graphql do
  >   mutations do
  >     create Tunez.Accounts.User, :register_with_password, :register_with_password
  >   end
  > end
```

3. https://hexdocs.pm/ash_json_api/dsl-ashjsonapi-domain.html#json_api-routes-base_route-route

4. https://hexdocs.pm/ash_graphql/dsl-ashgraphql-domain.html#graphql-mutations-create

```
# ...
```

This is exactly the same as mutations we added for creating or updating artist and album records — the type of action, the name of the module, the name of the action, and the name to use for the generated mutation.

In the GraphQL playground, at <http://localhost:4000/gql/playground>, the new mutation appears as part of the schema:

The screenshot shows the GraphQL playground interface with the schema definition for the 'registerWithPassword' mutation. The schema is organized into sections: QUERIES, MUTATIONS, and TYPE DETAILS. The 'registerWithPassword' mutation is defined under the MUTATIONS section. Its type is 'RegisterWithPasswordResult'. The 'RegisterWithPasswordResult' type has fields for 'errors' (a list of 'MutationError') and 'metadata' (a 'RegisterWithPasswordMetadata' object). The 'RegisterWithPasswordMetadata' object has a 'token' field of type 'String!'. The 'result' field of the 'RegisterWithPasswordResult' type is of type 'User'. A search bar at the top is empty, and the bottom navigation bar includes links for 'Documentation', 'Playground', 'GraphiQL', and 'Logout'.

```

    graph TD
      schema[Search the schema ...]
      queries[QUERIES]
      mutations[MUTATIONS]
      typeDetails[TYPE DETAILS]

      registerWithPassword@registerWithPasswordMutation
      registerWithPasswordResult@registerWithPasswordResultType
      errors@registerWithPasswordResultType
      metadata@registerWithPasswordResultType
      token@registerWithPasswordMetadata
      result@registerWithPasswordResultType

      registerWithPassword --> registerWithPasswordMutation
      registerWithPasswordMutation --> registerWithPasswordResult
      registerWithPasswordResult --> errors
      registerWithPasswordResult --> metadata
      metadata --> token
      registerWithPasswordResult --> result
  
```

And that's all that needs to be done! AshGraphql can tell that the action has metadata, and that metadata should be exposed as part of the response. You can test it out in the GraphQL playground, by calling the mutation with an email, password, and password confirmation, the same way as you would on the web registration form, and reading back the token as part of the metadata:

```

mutation {
  registerWithPassword(input: {email: "test@example.com", password: "mypassword",
                                passwordConfirmation: "mypassword"}) {
    result {
      id
      email
    }
    metadata {
      token
    }
  }
}
  
```

To create a query for the `sign_in_with_password` action, update the `graphql` block in the `Tunez.Accounts` domain module and add the query using the `get` macro:⁵

```
06/lib/tunez/accounts.ex
defmodule Tunez.Accounts do
  use Ash.Domain, extensions: [AshGraphql.Domain, AshJsonApi.Domain]

  graphql do
    queries do
      get Tunez.Accounts.User, :sign_in_with_password, :sign_in_with_password
    end
  end
  # ...
```

This macro is typically for fetching records by their primary key, or `id`, and this is reflected in the GraphQL playground's schema — the query takes an email and password as inputs, but *also* an `id` input. Not what we want! We can disable this `id` input by adding `identity false`⁶ to our query definition.

We can now call the query in the playground, and get either user data back, or an authentication failure if we provided invalid credentials. The successful response doesn't contain the user's authentication token, though, and we need that!

AshGraphql has the hidden ability to generate custom types that include metadata in them, using the `type_name` option⁷ for a query. The default type name is what we've defined in the resource — in this case, the type name `:user` was automatically added to our `Tunez.Accounts.User` resource. If we update that in our query options, to *any* other value, the metadata from the action will now be included in the response, as a new type will be generated that can include it. Not all actions that return users will have metadata, so the metadata fields *can't* be added to the base user GraphQL type!

With a new type name added for this query, it can now return the user's authentication token as part of a successful response:

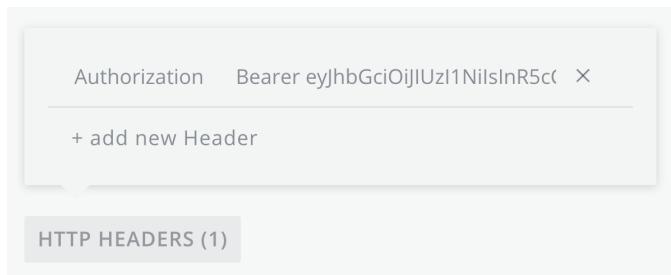
```
06/lib/tunez/accounts.ex
queries do
  get Tunez.Accounts.User, :sign_in_with_password, :sign_in_with_password do
    identity false
  end
  type_name :user_with_token
end
end
```

5. https://hexdocs.pm/ash_graphql/dsl-ashgraphql-domain.html#graphql-queries-get

6. https://hexdocs.pm/ash_graphql/dsl-ashgraphql-domain.html#graphql-queries-get-identity

7. https://hexdocs.pm/ash_graphql/dsl-ashgraphql-domain.html#graphql-queries-get-type-name

The user token returned as part of the API response can then be used to authenticate subsequent API requests, by setting the token in an Authorization HTTP header.



And now users can successfully sign in, register, and authenticate via both of our APIs!

Let's take a look back at how much code we needed to write, to implement this new functionality:

- API configuration in the Accounts domain module, to make our existing `register_with_password` and `sign_in_with_password` User resource actions accessible over our GraphQL and JSON APIs
- A new policy in our User resource, to allow public access to the two actions

And... that's it! With just a few lines of DSL code, we've got a robust authentication system for two different API frameworks!

Now we can look at more general authorization, for the rest of Tunez. This data isn't going to secure itself!

Assigning Roles to Users

Access control — who can access what — is a massive topic. Systems to control access to data can range from the simple (users can perform actions based on a single role), to the more complex (users can be assigned to one or more groups, each with its own permissions; or hold different roles within different groups) to the very very fine-grained (users can be granted specific permissions per piece of data, on top of everything else).

What level of access control you need heavily depends on the specific app you're building, and it may change over time. For Tunez right now, we don't need anything really complicated, we only want to make sure data doesn't get vandalized, so we'll implement a more simple system with *roles*.

Each user will have an assigned *role*, which determines what actions they can take and what data they can modify. We'll have three different roles:

- Basic users won't be able to modify any artist/album data
- Editors will be able to create/update a limited set of data
- Admins will be able to perform any action across the app

This role will be stored in a new attribute in the `Tunez.Accounts.User` resource, named `role`. What type of attribute should it be, though? It could be an atom, and we could add a constraint⁸ to specify that it must be one of our list of valid role atoms:

```
attributes do
  # ...

  attribute :role, :atom do
    allow_nil? false
    default :user
    constraints [one_of: [:admin, :editor, :user]]
  end
end
```

Ash provides a better way to handle enum-type values, though, with the `Ash.Type.Enum`⁹ behaviour. We can define our roles in a separate `Tunez.Accounts.Role` module:

```
06/lib/tunez/accounts/role.ex
defmodule Tunez.Accounts.Role do
  use Ash.Type.Enum, values: [:admin, :editor, :user]
end
```

And then specify that the `role` attribute is actually a `Tunez.Accounts.Role`:

```
06/lib/tunez/accounts/user.ex
attributes do
  # ...

  ▶ attribute :role, Tunez.Accounts.Role do
  ▶   allow_nil? false
  ▶   default :user
  ▶ end
end
```

This has a couple of neat benefits. We can fetch a list of all of the valid roles, with `Tunez.Accounts.Role.values/0`, and we can also specify human-readable values for each role, useful if you wanted a page where you could select a role from a dropdown list.

Generate a migration for the new `role` attribute, and run it:

```
$ mix ashcodegen add_role_to_users
```

8. <https://hexdocs.pm/ash/dsl-ash-resource.html#attributes-attribute-constraints>

9. <https://hexdocs.pm/ash/Ash.Type.Enum.html>

```
$ mix ash.migrate
```

The attribute is created as a text column in the database, and the default option¹⁰ we used is also passed through to be a default in the database. This means that we don't need to manually set the role for any existing users, or any new users that sign up for accounts, they'll all automatically have the role user.

What about custom logic for assigning roles?

Maybe you don't want to hardcode a default role for all users — maybe you want to assign the 'editor' role to users who register with a given email domain, for example.

You can implement this with a custom change module, similar to [UpdatePreviousNames on page 57](#). The wrinkle here is that users can be created either from signing up with a password, or signing in with a magic link (which creates a new account if one doesn't exist for that email address). The change code would need to be used in both actions, and support magic link sign-ins for both new and existing accounts.

We do need some way of changing an user's role though, otherwise Tunez will never have any editors or admins! We can add some utility actions to the Tunez.Accounts.User resource — a read action that will allow fetching a single user by id, and a new update action that *only* allows setting the role attribute for a given user record:

```
06/lib/tunez/accounts/user.ex
actions do
  defaults [:read]
  > update :set_role do
  >   accept [:role]
  > end
  # ...
```

Adding code interfaces for the two new actions will make them easier to run, so we can add those within the Tunez.Accounts.User resource definition in the Tunez.Accounts domain module.

```
06/lib/tunez/accounts.ex
defmodule Tunez.Accounts do
  # ...
  resources do
    # ...
    > resource Tunez.Accounts.User do
    >   define :set_user_role, action: :set_role, args: [:role]
```

10. <https://hexdocs.pm/ash/dsl-ash-resource.html#attributes-attribute-default>

```
>     define :get_user_by_id, action: :read, get_by: [:id]
>   end
```

As we will never be calling these actions from code, only running them manually in an iex console, we don't need to add a policy that will authorize them — we can skip authorization using the `authorize?: false` option when we call them.

Once you've registered an account in your development Tunez app, you can then change it to be an admin in iex:

```
iex(1)> user = Tunez.Accounts.get_user_by_id!(<<uuid>>, authorize?: false)
#Tunez.Accounts.User<...>
iex(2)> Tunez.Accounts.set_user_role(user, :admin, authorize?: false)
{:ok, #Tunez.Accounts.User<role: :admin, ...>}
```

Now that we have users with roles, we can define which roles can perform which actions. We can do this by writing some more policies, to cover actions in our Artist and Album resources.

Writing Policies for Artists

There are two parts we need to complete, when implementing access control in an Ash web app:

- Creating the policies for our resources, and
- Updating our web interface to pass the actor to calls to our actions, as well as niceties like hiding buttons to perform actions if the current user doesn't have permission to run them.

Which order you do them in is up to you, but both will need to be done. Writing the policies first will break the web interface completely (everything will be forbidden, without knowing who is trying to load or modify data!) but doing it first will ensure we catch all permission-related issues when tweaking our liveviews. We'll write the policies first, while they're fresh in our minds.

Creating our first artist policy

To start adding policies to a resource for the first time, we first need to configure it with the Ash policy authorizer. In the `Tunez.Music.Artist` resource, that looks like this:

```
06/lib/tunez/music/artist.ex
defmodule Tunez.Music.Artist do
  use Ash.Resource,
  otp_app: :tunez,
  domain: Tunez.Music,
  data_layer: AshPostgres.DataLayer,
  extensions: [AshGraphql.Resource, AshJsonApi.Resource],
```

- `authorizers: [Ash.Policy.Authorizer]`

Testing a create action

The default behaviour in Ash is to authorize (run policy checks for) any action in a resource that has an authorizer configured, so straight away we get forbidden errors when attempting to run any actions on the `Tunez.Music.Artist` resource:

```
iex(1)> Tunez.Music.create_artist(%{name: "New Artist"})
{:error,
 %Ash.Error.Forbidden{
   changeset: "#Changese<>",
   errors: [%Ash.Error.Forbidden.Policy{...}]}
}
```

There are no policies for the create action, so it is automatically forbidden.

We can add a policy for the action by adding a new policies block at the top level of the `Tunez.Music.Artist` resource, and adding a sample policy that applies to the action:

```
06/lib/tunez/music/artist.ex
defmodule Tunez.Music.Artist do
  # ...
  policies do
    policy action(:create) do
      authorize_if always()
    end
  end
end
```

This is the most permissive type of policy check — it *always* authorizes the given policy. With this policy in our resource, after recompiling we can now create artists again:

```
iex(1)> Tunez.Music.create_artist(%{name: "New Artist"})
{:ok, #Tunez.Music.Artist<...>}
```

There are a whole set of policy checks built into Ash,¹¹ but the most common ones for policy conditions are `action` and `action_type`. Our policy, as written, only applies to the action *named* `create`, but if we wanted to use it for any action of *type* `create`, we could use `action_type(:create)` instead.

Of course, we don't always want to blanket authorize actions, as that would defeat the purpose of authorization entirely. Let's update the policy to only

11. <https://hexdocs.pm/ash/Ash.Policy.Check.Builtins.html>

allow admin users to create artist records, ie. *actors* who have their *role* attribute set to “admin”:

```
06/lib/tunez/music/artist.ex
policies do
  policy action(:create) do
    authorize_if actor_attribute_equals(:role, :admin)
  end
end
```

Because our policy refers to an actor, we have to pass in the actor when calling the action. If we don’t, or pass in `nil` instead, the check won’t make a decision. If none of the checks specifically authorize the policy, it safely defaults to being unauthorized. We can test what the action does in iex, by creating different types of `Tunez.Accounts.User` structs and running the create action:

```
iex(2)> Tunez.Music.create_artist(%{name: "New Artist"})
{:error, %Ash.Error.Forbidden...}

iex(3)> Tunez.Music.create_artist(%{name: "New Artist"}, actor: nil)
{:error, %Ash.Error.Forbidden...}

iex(4)> editor = %Tunez.Accounts.User{role: :editor}
iex(5)> Tunez.Music.create_artist(%{name: "New Artist"}, actor: editor)
{:error, %Ash.Error.Forbidden...}

iex(6)> admin = %Tunez.Accounts.User{role: :admin}
iex(7)> Tunez.Music.create_artist(%{name: "New Artist"}, actor: admin)
{:ok, #Tunez.Music.Artist<...>}
```

The only actor that was authorized to create the artist record was the admin user — just as we intended.

Filling out update and destroy policies

In a similar way, we can write policies for the update and destroy actions in the `Tunez.Music.Artist` resource. Admin users should be able to perform both actions; and we’ll also allow editors (users with `role: :editor`) to update records.

```
06/lib/tunez/music/artist.ex
policies do
  # ...

  policy action(:update) do
    authorize_if actor_attribute_equals(:role, :admin)
    authorize_if actor_attribute_equals(:role, :editor)
  end

  policy action(:destroy) do
    authorize_if actor_attribute_equals(:role, :admin)
  end
end
```

If an editor attempts to update an artist record, the first check won't make a decision — their role isn't "admin", so the next check is looked at. This one makes a decision to authorize, so the policy is authorized and the action will run.

Cutting out repetitiveness with bypasses

When we have some all-powerful role like admin, it can be really repetitive to write `authorize_if actor_attribute_equals(:role, :admin)` in every Single Policy. It'd be much nicer to be like, oh, these users? They're special, just let 'em on through. We can do that by using a *bypass*.¹²

With standard policies defined using `policy`, *all* applicable policies for an action must apply. This allows for cases like the following:

```
policy action_type(:update) do
  authorize_if actor_present()
end

policy action(:force_update) do
  authorize_if actor_attribute_equals(:role, "admin")
end
```

Our hypothetical resource has a few update-type actions that any authenticated user can run (using the built-in `actor_present` policy check),¹³ but also a special `force_update` that only admin users can run. It's not that the `action(:force_update)` policy takes precedence over the `action_type(:update)` policy, it's that both policies apply when calling the `force_update` action, but only one applies to other update actions.

Bypass policies are different. If a bypass policy *authorizes* an action, it can skip all other policies that apply to that action.

```
bypass actor_attribute_equals(:role, "admin") do
  authorize_if always()
end

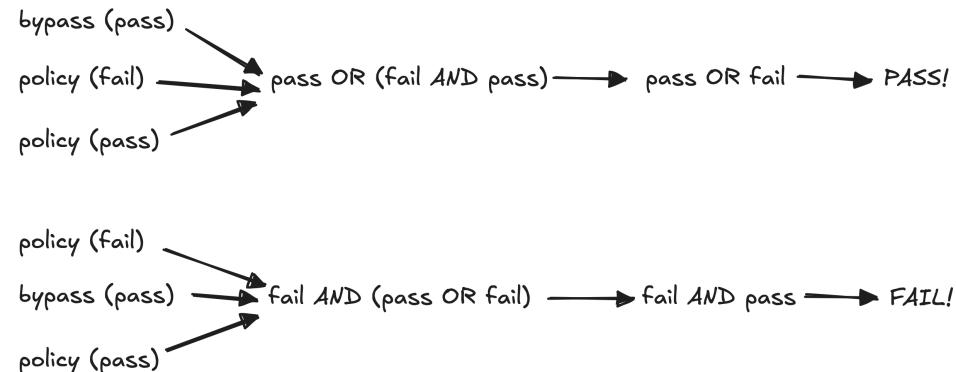
# If the bypass passes, then the result of this policy doesn't matter!
policy action_type(:update) do
  forbid_if always()
end
```

We say that a bypass *can* skip other policies, not that it *will*, because it's a little more complicated than that — order matters when it comes to mixing bypass policies with standard policies.

12. <https://hexdocs.pm/ash/dsl-ash-policy-authorizer.html#policies-bypass>

13. https://hexdocs.pm/ash/Ash.Policy.Check.Builtins.html#actor_present/0

Internally, Ash converts the results of running each policy into a big boolean expression with passing authorization being true and failing authorization being false, which is evaluated by the SAT solver we installed earlier. Standard policies are AND-ed into the expression, so all need to be authorized for the action to be authorized. Bypass policies are OR-ed into the expression, so changing the order of policies within a resource can drastically affect the result. See the following example, where the policy results are the same, but the order is different, leading to a different overall result:



What does this mean when writing policies? Bypasses are really powerful, and allow abstracting common authorization logic into one place. It's possible to encode very complicated logic, but it's also pretty easy to make a mess or have unintended results. We'd recommend the following guidelines:

- Keep all bypass policies together at the start of the policies block, and don't intermingle them with standard policies, and
- Write naive tests for your policies that test as many combinations of permissions as possible, to verify that the behaviour is what you expect. More on testing in the *the (as yet) unwritten content*.

Debugging when policies fail

Policies can get complex, with multiple conditions in each policy check, multiple policy checks of different types within a policy, and as we've just seen, multiple policies that apply to a single action request, including bypasses.

We can tell if authorization fails, because we get an `Ash.Error.Forbidden` struct back from the action request, but we can't necessarily see *why* it might fail.

Ash can display *breakdowns* of how policies were applied to an action request. Similar to how we set a config option to debug authentication failures in the last chapter, we can do the same thing for policies.¹⁴

14. <https://hexdocs.pm/ash/policies.html#policy-breakdowns>

In your dev config in config/dev.exs, turn on policy breakdowns:

```
06/config/dev.exs
config :ash, :policies, show_policy_breakdowns?: true
```

After restarting your iex session, repeat the previous experiment by trying to create an artist with an actor that isn't an admin:

```
iex(1)> editor = %Tunez.Accounts.User{role: :editor}
#Tunez.Accounts.User<...>
iex(2)> Tunez.Music.create_artist!(%{name: "Oh no!"}, actor: editor)
** (Ash.Error.Forbidden)
Bread Crumbs:
  > Error returned from: Tunez.Music.Artist.create
Forbidden Error
* forbidden:
Tunez.Music.Artist.create
Policy Breakdown
  user: %{id: nil}

  Policy | [M]:
    condition: action == :create
    authorize if: actor.role == :admin | ✘ | [M]

SAT Solver statement:
  "action == :create" and
  (( "action == :create" and "actor.role == :admin" )
  or not "action == :create")
```

(Note: The [M]s are actually magnifying glass emojis in the terminal)

Similar output can be seen in a browser, if you attempt to visit the app (because it's all broken right now!) It's a little bit verbose, but it clearly states which policies have applied to this action call, and what the results were — this user was not an admin, so they get a big ✘.

Filtering results in read action policies

The last actions we have to address in the Artist resource are our two read actions — the default read action, and our custom search action.

So far we've only looked at checks for single records, with yes/no answers — can the actor run this action on this record, yes or no? Read actions are built a little differently, as they don't start with an initial record to operate on and they don't modify data. Policies for read actions behave as *filters* — given all of the records that the action would fetch, which ones are the actor allowed to see?

If, for example, we had the following policy that uses a secret attribute:

```
policy action_type(:read) do
  authorize_if expr(secret == false)
end
```

You can call the action as normal, via code like `MyResource.read(actor: user)`, and the results will only include records where the `secret` attribute has the value `false`.

If a policy check would have different answers depending on the record being checked (ie. it checks some property of the record, like the value of the `secret` attribute) we say this is a *filter* check. If it depends only on the actor, or a static value like `always()`, then we say it's a *simple* check.

Filter checks and simple checks can be included in the same policy, eg. to allow `admins` to read all records, but `non-admins` can only read `non-secret` records:

```
policy action_type(:read) do
  authorize_if expr(secret == false)
  authorize_if actor_attribute_equals(:role, "admin")
end
```

Trust, but verify!

One quirk of read policies is distinguishing between “the actor *can't* run the action” and “the actor *can* run the action, but all of the results are filtered out”.

By default, all read actions are runnable and all checks are applied as filters. If you want the whole action to be forbidden on authorization failure, this can be configured in the policy using the `access_type` option.¹⁵



The following policy check would allow admin users to run it and fetch results, but non-admin users would get an forbidden error.

```
policy action(:read_banned_users) do
  access_type :strict
  authorize_if actor_attribute_equals(:role, "admin")
end
```

Tunez won't have any restrictions on reading artists, but we *do* need to have policies for all actions in a resource once we start adding them, so we can add a blanket `authorize_if always()` policy:

15. <https://hexdocs.pm/ash/policies.html#access-type>

```
06/lib/tunez/music/artist.ex
policies do
  # ...
  policy action_type(:read) do
    authorize_if always()
  end
end
```

Removing Forbidden Actions from the UI

At the moment, the Artist resource in Tunez is secure — actions that modify data can only be called if a) we pass in a user record as the actor and b) that actor is authorized to run that action. So far, so good.

The web UI doesn't reflect these changes, though. Even when not logged in to the app, we can still see buttons and forms inviting us to create, edit, or delete data.

The screenshot shows a web page for the artist 'Cielarko'. At the top, there's a navigation bar with a logo, 'Sign In', 'or', and 'Register' buttons. Below the header, the artist's name 'Cielarko' is displayed in a large font, with 'Delete Artist' and 'Edit Artist' buttons to its right. A bio text describes 'Cielarko' as a folk metal band from Belgium that blends Esperanto lyrics with traditional instruments. Below the bio, there's a 'New Album' button and a thumbnail for an album titled 'Verda Horizonto (2023)' with 'Delete' and 'Edit' buttons next to it.

We can't actually *run* the actions, so clicking the buttons and submitting the forms will return an error, but it's not a good user experience to see them at all. And if we *are* logged in, so we *should* have access to manage data, we *still* get an error! Oops.

There are a few things we need to do, to make the UI behave correctly for any kind of user viewing it:

- Update all of our direct action calls to pass the current user as the actor
- Update our forms to ensure we only let the current user see them if they can submit them
- And lastly, update our templates to only show buttons if the current user is able to use them.

It sounds like a lot, but it's only a few changes to make, spread across a few different files. Let's dig in!

Identifying the actor when calling actions

For a more complex app, this would be the biggest change from a functionality perspective — allowing actions to be called by users who *are* authorized to do things. Tunez is a lot simpler, and most of the data management is done via forms, so this isn't a massive change for us. The only actions we call directly are read and destroy actions:

`Tunez.Music.search_artists/2`, in `Tunez.Artists.IndexLive`. We don't strictly *need* to pass the actor in here, as our current policies will allow the action for everyone even if they're not authenticated, but that could change in the future. If it does, we don't want to forget to set the actor, so we may as well do it now!

```
06/lib/tunez_web/live/artists/index_live.ex
def handle_params(params, _url, socket) do
  # ...

  page =
    Tunez.Music.search_artists!(query_text,
      page: page_params,
      query: [sort_input: sort_by],
  ▶   actor: socket.assigns.current_user
  )
```

`Tunez.Music.get_artist_by_id/2`, in `Tunez.Artists.ShowLive`. Again, we don't strictly need it because everyone can read artist information, but like the `search_artists` action, we can't guarantee that this will always be the case in the future. It does no harm to set the actor either, so we'll add it.

```
06/lib/tunez_web/live/artists/show_live.ex
def handle_params(%{"id" => artist_id}, _session, socket) do
  artist =
    Tunez.Music.get_artist_by_id!(artist_id,
      load: [:albums],
  ▶   actor: socket.assigns.current_user
  )
```

`Tunez.Music.destroy_artist/2`, in `Tunez.Artists.ShowLive`. This one we *do* need to pass the actor in to make it work, as only specific types of users can delete artists.

```
06/lib/tunez_web/live/artists/show_live.ex
def handle_event("destroy_artist", _params, socket) do
  case Tunez.Music.destroy_artist(
    socket.assigns.artist,
  ▶   actor: socket.assigns.current_user
  ) do
```

```
# ...
```

Tunez.Music.destroy_album/2, in Tunez.Artists.ShowLive. We haven't added policies for albums yet, but it does no harm to start tweaking our templates to support them now.

```
06/lib/tunez_web/live/artists/show_live.ex
def handle_event("destroy_album", %{"id" => album_id}, socket) do
  case Tunez.Music.destroy_album(
    album_id,
    actor: socket.assigns.current_user
  ) do
    # ...
```

Not too onerous! Moving forward, we'll add the actor to every action we call, to avoid this kind of rework.

Updating forms to identify the actor

This is the biggest change for Tunez, as we create and edit both artists and albums via forms. There are two parts to this: setting the actor when building the forms, and ensuring that the form is submittable.

We don't want to show the form at all if the user wouldn't be able to submit it, so we need to run the submittable check before rendering, in the mount/3 functions of Tunez.Artists.FormLive:

```
06/lib/tunez_web/live/artists/form_live.ex
def mount(%{"id" => artist_id}, _session, socket) do
  artist = Tunez.Music.get_artist_by_id!(artist_id)

  form =
    Tunez.Music.form_to_update_artist(
      artist,
      actor: socket.assigns.current_user
    )
  |> AshPhoenix.Form.ensure_can_submit!()

  # ...

  def mount(_params, _session, socket) do
    form =
      Tunez.Music.form_to_create_artist(
        actor: socket.assigns.current_user
      )
    |> AshPhoenix.Form.ensure_can_submit!()
```

AshPhoenix.Form.ensure_can_submit!/1¹⁶ is a neat little helper function that authorizes the configured action and data in the form using our defined policies, to make

16. https://hexdocs.pm/ash_phoenix/AshPhoenix.Form.html#ensure_can_submit!/1

sure it is submittable. If the authorization fails, then the form can't be submitted, and an exception will be raised.

We can make the same changes to the mount/3 functions in Tunez.Albums.FormLive:

```
06/lib/tunez_web/live/albums/form_live.ex
def mount(%{"id" => album_id}, _session, socket) do
  album = Tunez.Music.get_album_by_id!(album_id, load: [:artist])

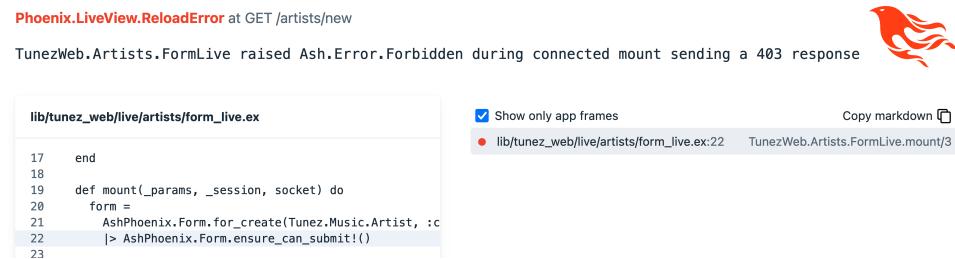
  form =
    Tunez.Music.form_to_update_album(
      album,
    >    actor: socket.assigns.current_user
    )
  > |> AshPhoenix.Form.ensure_can_submit!()

  # ...

def mount(%{"artist_id" => artist_id}, _session, socket) do
  artist = Tunez.Music.get_artist_by_id!(artist_id)

  form =
    Tunez.Music.form_to_create_album(
      transform_params: fn _form, params, _context ->
        Map.put(params, "artist_id", artist.id)
      end,
    >    actor: socket.assigns.current_user
    )
  > |> AshPhoenix.Form.ensure_can_submit!()
```

Now if you click any of the buttons that link to pages focussed on forms, when not logged in as a user with the correct role, an exception will be raised and you'll get a standard Phoenix error page:



The screenshot shows a Phoenix error page. At the top, it says "Phoenix.LiveView.ReloadError at GET /artists/new". Below that, it says "TunezWeb.Artists.FormLive raised Ash.Error.Forbidden during connected mount sending a 403 response". To the right of the text is the Phoenix logo, which is a stylized red bird in flight. The main content area shows the source code for "lib/tunez_web/live/artists/form_live.ex". Lines 17 through 23 are highlighted in blue. The code defines a mount function that creates a form, sets the actor to the current user, and ensures the form is submittable. There are checkboxes for "Show only app frames" and "Copy markdown" with a clipboard icon.

```
lib/tunez_web/live/artists/form_live.ex
17   end
18
19 def mount(_params, _session, socket) do
20   form =
21     AshPhoenix.Form.for_create(Tunez.Music.Artist, :c
22   |> AshPhoenix.Form.ensure_can_submit!
23 
```

That works well for forms — but what about entire pages? Maybe we've built an admin-only area, or we've added an Artist version history page that only editors can see. We can't use the same form helpers to ensure access, but we can prevent users from accessing what they shouldn't.

Blocking pages from unauthorized access

When we installed AshAuthenticationPhoenix, one file that the installer created was the `TunezWeb.LiveUserAuth` module, in `lib/tunez_web/live_user_auth.ex`. We haven't looked at that file yet, but we will now!

It contains several `on_mount` function definitions, that do different things based on the authenticated user (or lack of) — the `live_user_optional` function head will make sure there's always a `current_user` set in the socket assigns, even if it's `nil`; the `live_user_required` function head will redirect away if there's no user logged in, and the `live_no_user` function head will redirect away if there *is* a user logged in!

These are LiveView-specific helper functions,¹⁷ that can be called at the root level of any liveview like so:

```
defmodule Tunez.Accounts.ForAuthenticatedUsersOnly do
  use TunezWeb, :live_view

  # or :live_user_optional, or :live_no_user
  on_mount {TunezWeb.LiveUserAuth, :live_user_required}

  # ...
```

So to block a liveview from unauthenticated users, we could drop that `on_mount` call with `:live_user_required` in that module, and the job would be done!

We can add more function heads to the `TunezWeb.LiveUserAuth` module as well for custom behaviour, such as role-based function heads.

```
06/lib/tunez_web/live_user_auth.ex
defmodule TunezWeb.LiveUserAuth do
  # ...

  def on_mount([role_required: role_required], _, _, socket) do
    current_user = socket.assigns[:current_user]

    if current_user && current_user.role == role_required do
      {:cont, socket}
    else
      socket =
        socket
        |> Phoenix.LiveView.put_flash(:error, "Unauthorized!")
        |> Phoenix.LiveView.redirect(to: ~p"/")

      {:halt, socket}
    end
  end
end
```

17. https://hexdocs.pm/phoenix_live_view/Phoenix.LiveView.html#on_mount/

This would allow us to write `on_mount` calls in a liveview like:

```
defmodule Tunez.Accounts.ForAdminsOnly do
  use TunezWeb, :live_view

  on_mount {TunezWeb.LiveUserAuth, role_required: :admin}
  # ...
```

Now we can secure all of our pages really neatly, both those that are form-based, and those that aren't. We still shouldn't see any shiny tempting buttons for things we can't access, though, so let's hide them if the user can't perform the actions.

Hiding calls to action that the actor can't perform

There are buttons sprinkled throughout our liveviews and components — buttons for creating, editing, and deleting artists; and for creating, updating and deleting albums. We can use Ash's built-in helpers to add general authorization checks to each of them, meaning we don't have to duplicate any policy logic and we won't need to update any templates if our policy rules change.

Ash.can?

The first helper function we'll look at is `Ash.can?`.¹⁸ This is a pretty low-level function, that takes a tuple representing the action to call and an actor, runs the authorization checks for the action, and returns a boolean representing whether or not the action is authorized:

```
iex(1)> Ash.can?({Tunez.Music.Artist, :create}, nil)
false
iex(2)> Ash.can?({Tunez.Music.Artist, :create}, %{role: :admin})
true
iex(3) artist = Tunez.Music.get_artist_by_id!(<<uuid>>
#Tunez.Music.Artist<id: <<uuid>>, ...>
iex(4)> Ash.can?({artist, :update}, %{role: :user})
false
iex(5)> Ash.can?({artist, :update}, %{role: :editor})
true
```

The format of the action tuple looks a lot like how you would run the action manually, as we covered in [Running actions, on page 13](#) — building a changeset for a create action with `Ash.Changeset.for_create(Tunez.Music.Artist, :create, ...)`, or for an update action with `Ash.Changeset.for_update(artist, :update, ...)`.

18. <https://hexdocs.pm/ash/Ash.html#can?/3>

Our liveviews and components don't call actions like this though, we use code interfaces for everything because they're a lot cleaner. Ash also defines some helper functions around authorization for code interfaces that are nicer to read, so we'll look at those next.

`can_*?` code interface functions

We call these `can_*`? functions because the names are dynamically generated based on the name of the code interface. For our `Tunez.Music` domain, for example, `iex` shows a whole set of functions with the `can_` prefix:

```
iex(1)> Tunez.Music.can_
can_create_album/1      can_create_album/2      can_create_album/3
can_create_album?/1     can_create_album?/2     can_create_album?/3
can_create_artist/1     can_create_artist/2     can_create_artist/3
can_create_artist?/1    can_create_artist?/2    can_create_artist?/3
can_destroy_album/2    can_destroy_album/3    can_destroy_album/4
...
...
```

This list includes `can_*`? functions for *all* code interfaces, even ones that don't have policies applied yet like `Tunez.Music.destroy_album`. If authorization isn't configured for a resource, both `Ash.can?` and the `can_*`? functions will simply return true, so we can safely update our templates to add authorization checks for all functions without fear of breaking anything.

One important thing to note is the order of arguments to the code interface helpers. Whereas `Ash.can?` always takes an action tuple and an actor (plus options), because some of the action tuple information is now in the function name, the first argument to a `can_*`? function is always the actor. If the code interface needs extra information, like a record to operate on or params, those come *after* the actor argument.

```
iex(4)> h Tunez.Music.can_create_artist?
def can_create_artist?(actor, params_or_opts \\ %{}, opts \\ [])
```

Runs authorization checks for `Tunez.Music.Artist.create`, returning a boolean.
See `Ash.can?/3` for more information

```
iex(5)> h Tunez.Music.can_update_artist?
def can_update_artist?(actor, record, params_or_opts \\ %{}, opts \\ [])
```

Runs authorization checks for `Tunez.Music.Artist.update`, returning a boolean.
See `Ash.can?/3` for more information

Armed with this new knowledge, we can now update the buttons in our templates, to wrap them in HEEx conditionals to only show the buttons if the relevant `can_*`? function returns true. There's one button in `Tunez.Artists.IndexLive`, for creating an artist:

```
06/lib/tunez_web/live/artists/index_live.ex
<.header responsive={false}>
  # ...
  ➤  <:action :if={Tunez.Music.can_create_artist?(@current_user)}>
    <.button_link <>opts>>New Artist</ .button_link>
  </:action>
</.header>
```

Two buttons in the header of Tunez.Artists.ShowLive, for editing/deleting an artist:

```
06/lib/tunez_web/live/artists/show_live.ex
<.header>
  # ...
  ➤  <:action :if={Tunez.Music.can_destroy_artist?(@current_user, @artist)}>
    <.button_link <>opts>>Delete Artist</ .button_link>
  </:action>
  ➤  <:action :if={Tunez.Music.can_update_artist?(@current_user, @artist)}>
    <.button_link <>opts>>Edit Artist</ .button_link>
  </:action>
</.header>
```

One above the album list in Tunez.Artists.ShowLive, for creating an album:

```
06/lib/tunez_web/live/artists/show_live.ex
<.button_link navigate={"/artists/#{@artist.id}/albums/new"} kind="primary"
  ➤  :if={Tunez.Music.can_create_album?(@current_user)}>
    New Album
  </ .button_link>
```

And two in the album_details function component in Tunez.Artists.ShowLive, for editing/deleting an album:

```
06/lib/tunez_web/live/artists/show_live.ex
<.header class="pl-4 pr-2 !m-0">
  # ...
  ➤  <:action :if={Tunez.Music.can_destroy_album?(@current_user, @album)}>
    <.button_link <>opts>>Delete</ .button_link>
  </:action>
  ➤  <:action :if={Tunez.Music.can_update_album?(@current_user, @album)}>
    <.button_link <>opts>>Edit</ .button_link>
  </:action>
</.header>
```

For these to work, you also need to update the call to the album_details function component, to pass the current_user in:

```
06/lib/tunez_web/live/artists/show_live.ex
<ul class="mt-10 space-y-6 md:space-y-10">
  <li :for={album <- @artist.albums}>
    <.album_details album={album} current_user={@current_user} />
  </li>
</ul>
```

Whew, that was fiddly. Moving forward, we'll wrap everything in authorization checks as we write it in our templates, so we don't have to do this kind of tweaking again.

Beware the policy check that performs queries!

Some policy checks can be devious, and require database queries to figure out if an action is authorized or not. They might reference the actor's group membership, or a count of associated records, or data other than what you've loaded for the page to render.

For a liveview app, if that related data isn't loaded and stored in memory, it'll be re-fetched to recalculate the authorization on every page render, which would be *disastrous* for performance!



You can have Ash make a best guess about authorization using data already in memory, by setting the `run_queries?`¹⁹ option to `Ash.can?/can_*` to false. If a decision can't be made definitively one way or the other, Ash will use the value of the `maybe_is?` option — by default this is true, but you can fail closed by setting it to false.

```
# Authorize based on data in memory, defaulting to unauthorized
Tunez.Music.can_run_complicated_action?(@current_user,
  run_queries?: false, maybe_is?: false)
```

Everything is now in place for artist authorization — you should be able to log in and out of your Tunez dev app as users with different roles, and the app should behave as expected around managing artist data. We've also added authorization checks around album management in our templates — but don't have any policies to go with them. We'll add those now.

Writing Policies for Albums

The rules we want to implement for album management are a little different to those for artist management. Our rules for artists could be summarized as:

- Everyone can read all artist data
- Editors can update (but not create or delete) artists
- Admins can perform any action on artist data

For albums, we want to implement the following:

- Everyone can read all album data

19. <https://hexdocs.pm/ash/Ash.html#can?/3>

- Editors can create album records, or update/delete album records *that they created*
- Admins can perform any action on album data

It's only a small change, but a common use case. In a issue tracker/help desk app, users might be assigned as owners of tickets and thus have extra permissions for those tickets. Or a user might be recorded as the owner of an organization, and have permissions to invite members to the organization.

The key piece of information we need that we're not currently storing, is *who* is creating each album in Tunez. Once we know that, we can write the policies that we want.

Recording who created and last modified a resource

To meet our requirements, we only need to store who created each album, but for a bit of fun we might take it a little bit further, and record who created and last modified records for both artists and albums.

(It'd be pretty neat to show who last updated an artist, on their profile page! Maybe later.)

To record this information for artists, we'll add two new relationships to the Tunez.Music.Artist resource, both pointing at the Tunez.Accounts.User resource — one named `created_by`, and one named `updated_by`.

```
06/lib/tunez/music/artist.ex
relationships do
  # ...
  ▶ belongs_to :created_by, Tunez.Accounts.User
  ▶ belongs_to :updated_by, Tunez.Accounts.User
end
```

We can do the exact same thing to the Tunez.Music.Album resource:

```
06/lib/tunez/music/album.ex
relationships do
  # ...
  ▶ belongs_to :created_by, Tunez.Accounts.User
  ▶ belongs_to :updated_by, Tunez.Accounts.User
end
```

Adding these relationships means an update to the database structure, so we need to generate a migration for the changes, and run it:

```
$ mix ashcodegen add_user_links_to_artists_and_albums
$ mix ash.migrate
```

Because we're now identifying the actor, on page 155 every time we submit a form to create or modify data, we can add changes to some of our actions to store that actor information in our new relationships. One of Ash's built-in change functions is `relate_actor`,²⁰ which does exactly what we want — relate the actor to the record being saved.

We'll do this a little bit differently than previous changes like `UpdatePreviousNames`, though. Storing the actor isn't really related to the business logic of what we want the action to do, it's more of a side effect. We *really* want to implement something like "by the way, whenever you create or update a record, can you also store who made the change? Cheers." So the logic shouldn't be restricted to only the actions *named* `:create` and `:update`, it should apply to *all* actions of type `create` and `update`.

We can do this with a resource-level changes block. Like validations and preparations, changes can be added either to individual actions or to the resource as a whole. In a resource-level changes block, we can also choose one or more action types that the change should apply to, using the `on` option.²¹

In the `Tunez.Music.Artist` resource, add a new top-level changes block and add the changes we want to store.

```
06/lib/tunez/music/artist.ex
defmodule Tunez.Music.Artist do
  # ...
  changes do
    change relate_actor(:created_by, allow_nil?: true), on: [:create]
    change relate_actor(:updated_by, allow_nil?: true), on: [:create]
    change relate_actor(:updated_by, allow_nil?: false), on: [:update]
  end
end
```

The exact same thing can be added to the `Tunez.Music.Album` resource:

```
06/lib/tunez/music/album.ex
defmodule Tunez.Music.Album do
  # ...
  changes do
    change relate_actor(:created_by, allow_nil?: true), on: [:create]
    change relate_actor(:updated_by, allow_nil?: true), on: [:create]
    change relate_actor(:updated_by, allow_nil?: false), on: [:update]
  end
end
```

20. https://hexdocs.pm/ash/Ash.Resource.Change.Builtins.html#relate_actor/

21. <https://hexdocs.pm/ash/dsl-ash-resource.html#changes-change-on>

These changes mean that if we add more create or update actions to the Artist or Album resources in future, they'll automatically have created_by and updated_by tracked — and we won't have to do a thing!

Why allow_nil?: true when creating records?

So that if you want to run or re-run the seed data scripts we provide with the starter Tunez app, they will successfully run both before and after adding these changes!

Depending on your app, you may also want to have nil values representing some kind of "system" action, if data may be created by means other than a user specifically submitting a form.

There's one other change we need to make — as we discovered earlier, the User resource is really locked-down, permission-wise. To relate the actor to a record when updating the record, we need to be able to *read* the actor record from the database, and at the moment we can't. All reading of user data is forbidden, unless being called internally by AshAuthentication.

To get around this, we can add another policy to the User resource, to allow a user to read their *own* record:

```
06/lib/tunez/accounts/user.ex
policies do
  # ...
  > policy action(:read) do
  >   authorize_if expr(id == ^actor(:id))
  > end
end
```

This uses the `^actor` expression template²² to reference the actor calling the action, as part of the policy condition expression. Like pinning a variable in a match or an Ecto query, this is how we can reference outside data that isn't a literal value (like `true` or `:admin`) and isn't an attribute or calculation on the resource (like `created_by_id`).

And that's all we need to do! Now whenever we call `create` or `update` (or their code interfaces) on either an artist or an album, the user ID of the actor will be stored in the `created_by_id` and/or the `updated_by_id` fields of the resource. You can test it out in iex, to make sure you've connected the pieces properly.

```
iex(1)> user = Tunez.Accounts.get_user_by_id!(<>uuid>, authorize?: false)
#Tunez.Accounts.User<id: <>uuid>, role: :admin, ...>
iex(2)> Tunez.Music.create_artist(%{name: "Who Made Me?"}, actor: user)
```

22. <https://hexdocs.pm/ash/expressions.html#templates>

```
{:ok,
#Tunez.Music.Artist<
  name: "Who Made Me?",
  updated_by: #Tunez.Accounts.User<id: <uuid>, ...>,
  created_by: #Tunez.Accounts.User<id: <uuid>, ...>,
  ...
>}
```

Filling out policies

All of the pre-requisite work has been done, the only thing left to do is write the actual policies for albums!

As with artists, the first step is enabling `Ash.Policy.Authorizer` in the `Tunez.Music.Album` resource:

```
06/lib/tunez/music/album.ex
defmodule Tunez.Music.Album do
  use Ash.Resource,
  otp_app: :tunez,
  domain: Tunez.Music,
  data_layer: AshPostgres.DataLayer,
  extensions: [AshGraphql.Resource, AshJsonApi.Resource],
  ➤ authorizers: [Ash.Policy.Authorizer]
```

This means that all of our action calls, including auto-loading albums when loading an artist record (which uses a `read` action!), will now automatically run authorization checks. And because we haven't yet defined any policies, they will all be forbidden by default.

We can reuse some of the policies that we wrote for artists, as the bulk of the rules are the same. In a new policies block in the `Tunez.Music.Album` resource, we can write a bypass for users with the role "admin", as they're allowed to run every action. As this will be the first policy in the policies block, if it passes, all other policies will be skipped.

```
06/lib/tunez/music/album.ex
defmodule Tunez.Music.Album do
  # ...

  policies do
    bypass actor_attribute_equals(:role, :admin) do
      authorize_if always()
    end
  end
end
```

We can also add an allow-all rule for *reading* album data — any user, authenticated or not, should be able to see the full list of albums for any artist.

```
06/lib/tunez/music/album.ex
policies do
  # ...
  policy action_type(:read) do
    authorize_if always()
  end
end
```

The main rules we want to look at are for editors. They'll have limited functionality — we want them to be able to create albums, and update/delete albums *if* they are related to those records via the `created_by` relationship.

The policy for create actions is pretty straightforward, it can use the same `actor_attribute_equals` built-in policy check we've used a few times now:

```
06/lib/tunez/music/album.ex
policies do
  # ...
  policy action(:create) do
    authorize_if actor_attribute_equals(:role, :editor)
  end
end
```

If the actor calling the create action doesn't have the role "admin" (which would be authorized by the bypass) or "editor" (which would be authorized by this create policy), then the action will be forbidden.

Lastly, we can write one policy that covers both the update and destroy actions, as the rules are identical for both. If we only wanted to verify the `created_by` relationship link, we could use the built-in `relates_to_actor_via` policy check,²³ like this:

```
policy action([:update, :destroy]) do
  authorize_if relates_to_actor_via(:created_by)
end
```

We *could* still technically use this! As only editors can create albums (ignoring the admin bypass), then if the album was created by the actor, the actor *must* be an editor! Right?? But... rules can change. Maybe in some months, a new "creator" role will be added that can *only* create records. But by the checks in this policy, they would also be authorized to update and destroy records, if they created them. Not good. Let's make the editor check in this policy explicit.

23. https://hexdocs.pm/ash/Ash.Policy.Check.Builtins.html#relates_to_actor_via/2

We can't combine built-in policy checks (that would be handy though!), so we'll have to fall back to writing an expression, like `expr(published == true)`, to verify both conditions in the same policy check. We end up with a policy like the following:

```
06/lib/tunez/music/album.ex
policies do
  # ...
  policy action_type([:update, :destroy]) do
    authorize_if expr(^actor(:role) == :editor and created_by_id == ^actor(:id))
  end
end
```

It's a little verbose, but it clearly captures our requirements — updates and deletes should be authorized if the actor's role is "editor", and the record was created by the actor.

Test it out in your app! Register a new user, make it an editor with `Tunez.Accounts.set_user_role/2`, and see how it behaves! As we already edited all of the templates to add authorization checks, we don't need to make any other changes. Note that your editor doesn't have access to edit any existing albums, but if they create a new one, they can then edit *that* one. Perfect!

And all of our authorization policies have automatically flowed through to our APIs. Trying to create albums or artists when not being authenticated will now be forbidden, but when user authentication tokens are provided in the request headers, the requests will now operate as normal. And we didn't need to do anything for that! We defined our policies once, in a central place, and they apply everywhere.

All this manual testing is getting a bit tiresome, though. We're starting to get more complicated logic in our app, and we can't keep manually testing everything. In the next chapter, we'll dive into testing — what to test, how to test it, and how Ash can help you get the best bang for your testing buck!

CHAPTER 7

All About Testing

Content to be supplied later.

CHAPTER 8

Fun With Nested Forms

Content to be supplied later.

CHAPTER 9

PubSub and Real-Time Notifications

Content to be supplied later.

Bibliography

- [Jur15] Saša Jurić. *Elixir in Action*. Manning Publications Co., Greenwich, CT, 2015.
- [TD24] Bruce A. Tate and Sophie DeBenedetto. *Programming Phoenix LiveView*. The Pragmatic Bookshelf, Dallas, TX, 2024.
- [TV19] Chris McCord, Bruce Tate and José Valim. *Programming Phoenix 1.4*. The Pragmatic Bookshelf, Dallas, TX, 2019.

Thank you!

We hope you enjoyed this book and that you're already thinking about what you want to learn next. To help make that decision easier, we're offering you this gift.

Head on over to <https://pragprog.com> right now, and use the coupon code BUYANOTHER2025 to save 30% on your next ebook. Offer is void where prohibited or restricted. This offer does not apply to any edition of *The Pragmatic Programmer* ebook.

And if you'd like to share your own expertise with the world, why not propose a writing idea to us? After all, many of our best authors started off as our readers, just like you. With up to a 50% royalty, world-class editorial services, and a name you trust, there's nothing to lose. Visit <https://pragprog.com/become-an-author/> today to learn more and to get started.

Thank you for your continued support. We hope to hear from you again soon!

The Pragmatic Bookshelf



The Pragmatic Bookshelf

The Pragmatic Bookshelf features books written by professional developers for professional developers. The titles continue the well-known Pragmatic Programmer style and continue to garner awards and rave reviews. As development gets more and more difficult, the Pragmatic Programmers will be there with more titles and products to help you stay on top of your game.

Visit Us Online

This Book's Home Page

<https://pragprog.com/book/ldash>

Source code from this book, errata, and other resources. Come give us feedback, too!

Keep Up-to-Date

<https://pragprog.com>

Join our announcement mailing list (low volume) or follow us on Twitter @pragprog for new titles, sales, coupons, hot tips, and more.

New and Noteworthy

<https://pragprog.com/news>

Check out the latest Pragmatic developments, new titles, and other offerings.

Buy the Book

If you liked this ebook, perhaps you'd like to have a paper copy of the book. Paperbacks are available from your local independent bookstore and wherever fine books are sold.

Contact Us

Online Orders: <https://pragprog.com/catalog>

Customer Service: support@pragprog.com

International Rights: translations@pragprog.com

Academic Use: academic@pragprog.com

Write for Us: <http://write-for-us.pragprog.com>